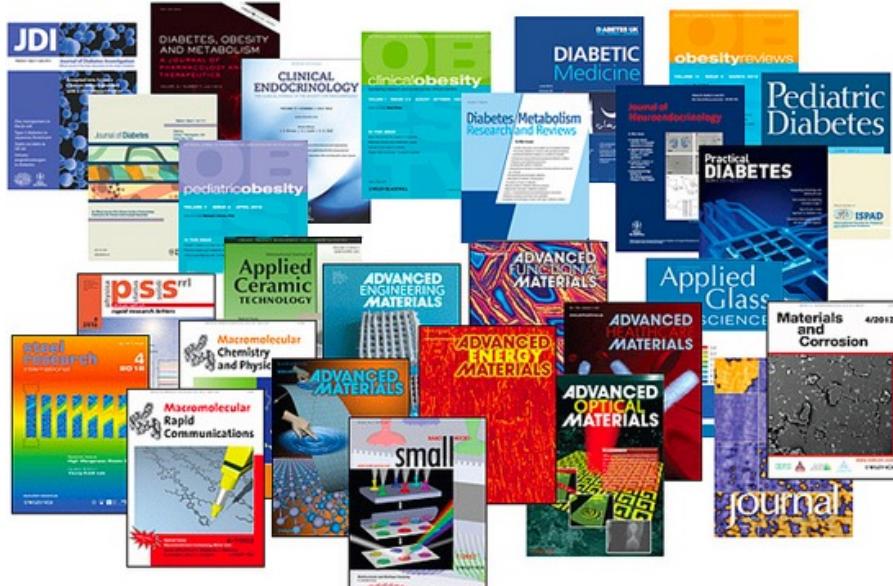


## Part 2

**Efficient** Knowledge Graph Construction  
Ningyu Zhang

# Knowledge Graph Construction

Information extraction for KG construction



Text2Knowledge

Language Pretrain  
BERT、GPT.....

属性：字

张之洞，晚清名臣、清代洋务派代表人物，祖籍直隶南皮，出生于贵州，字孝达。

实体

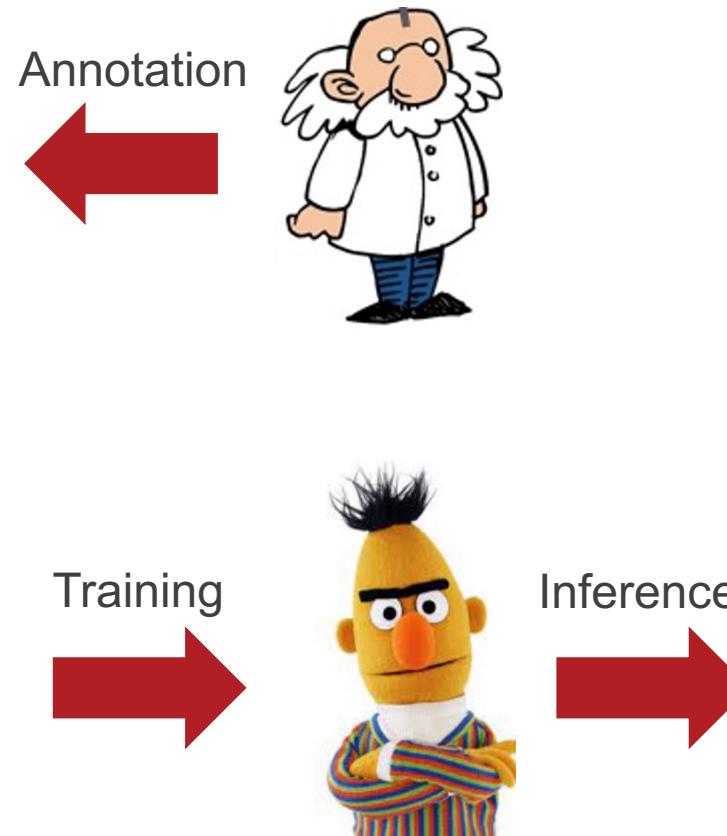
属性值

isSpouseOf

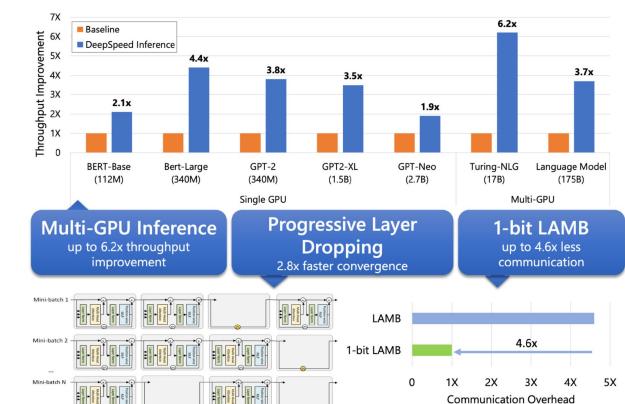
Jack is married to the microbiologist known as Dr. Germ in the USA.

Entity Pair





Expert Annotation Cost!  
GPU Computation Cost!  
Inference Time Cost!



□ What have been done?

**Cutting-edge Methods**

□ What can we do?

Remaining **Issues** and **Future Works**

□ How can we use?

Open-sourced **Toolkits**

## Cutting-edge Methods

# Scope: Efficient Knowledge Graph Construction



**Data Efficiency:** few-shot learning with different types of data

- Prompt learning, Retrieval-augmentation



**Model Efficiency:** parameter-efficient learning, unified architecture

- Parameter-efficient learning



**Inference Efficiency:** fast Inference

- FastRE, sequence-to-sequence

# Why Data-Efficient Knowledge Graph Construction?



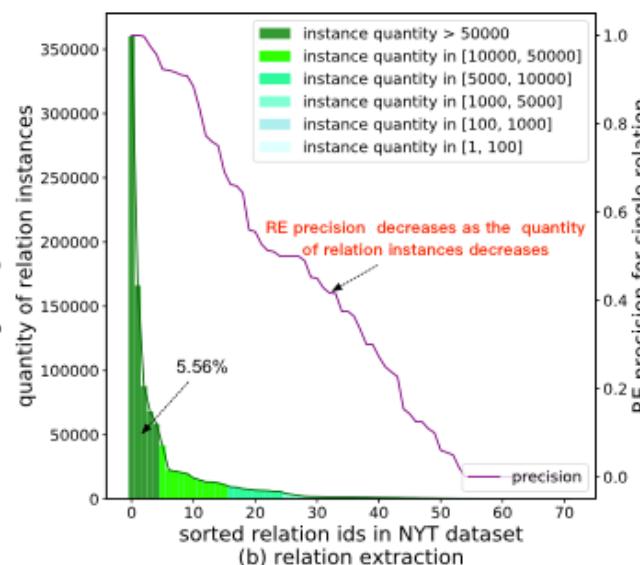
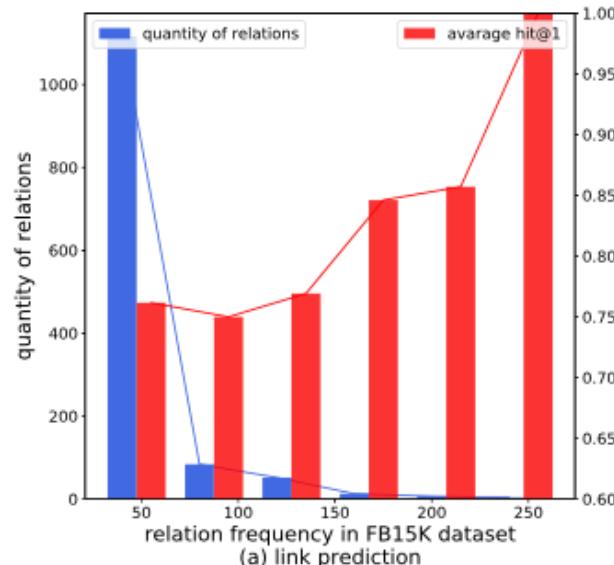
AACL IJCNLP 2022

Data Sparsity

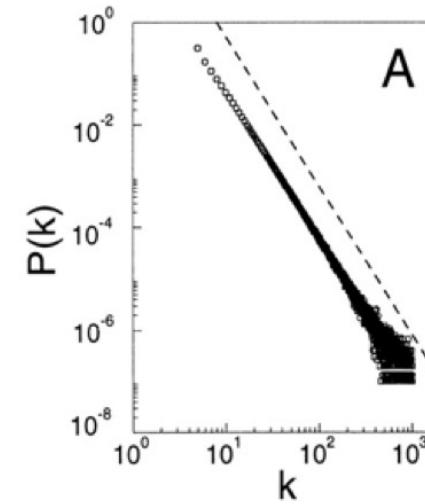


Annotation Cost

Knowledge: Long-tailed distribution



Graph: Power-law distributions

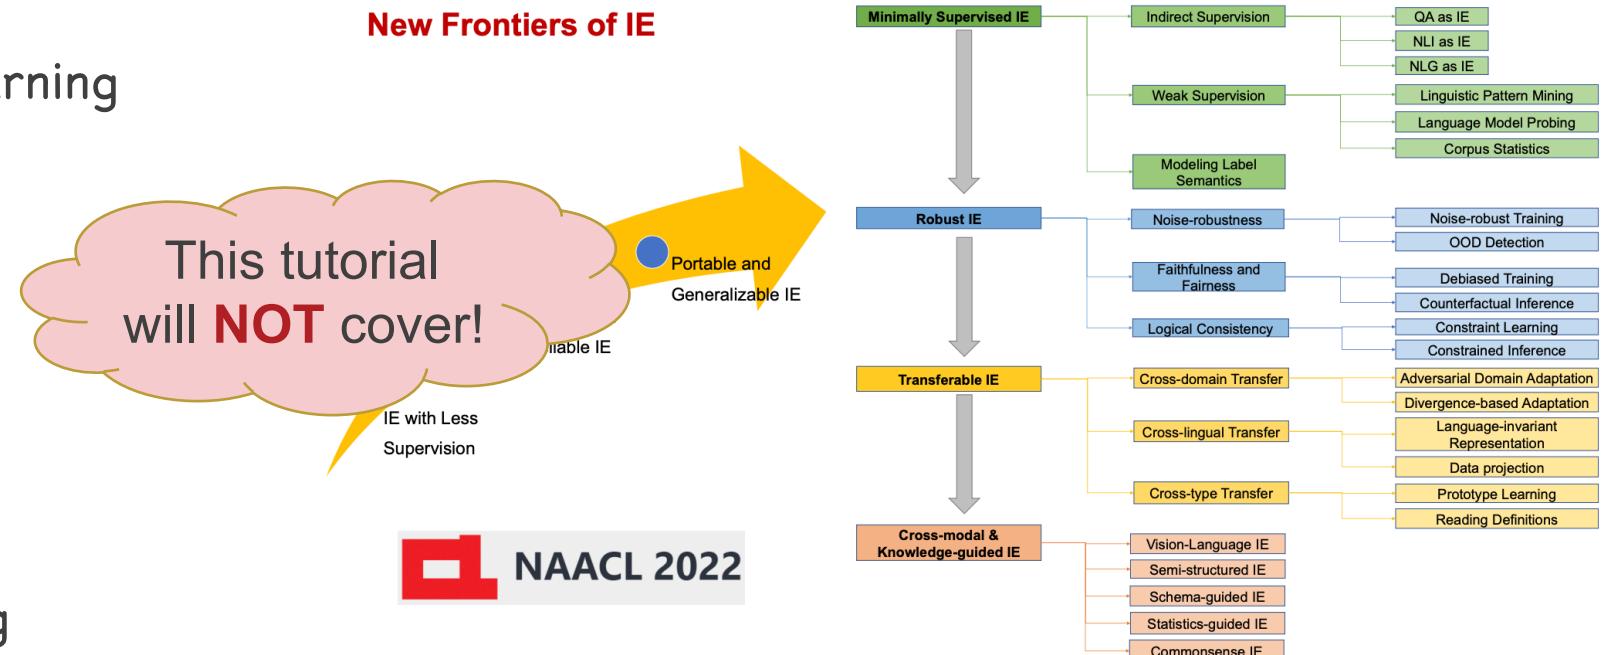


# Outline

- Transfer Learning



- Minimally Supervised Learning

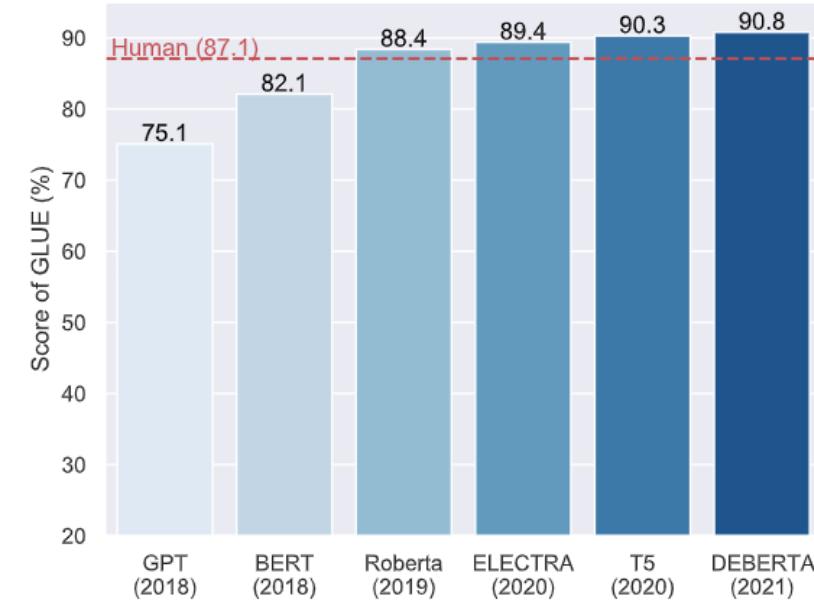
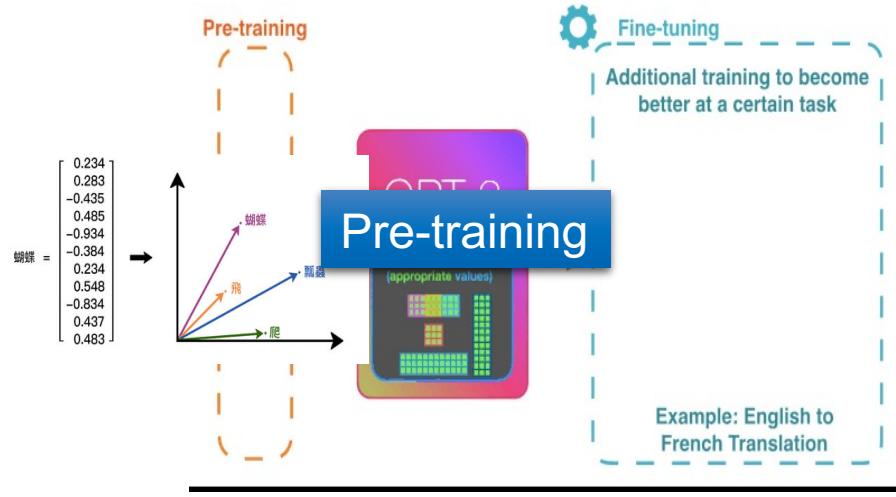


- Few/Zero-shot Learning

- Data Augmentation

- Semi-supervised Learning

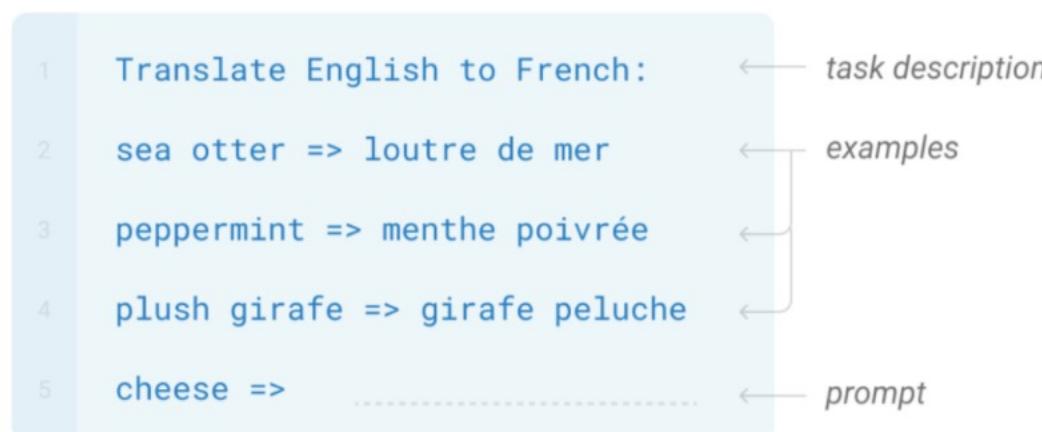
# Pre-trained Language Models



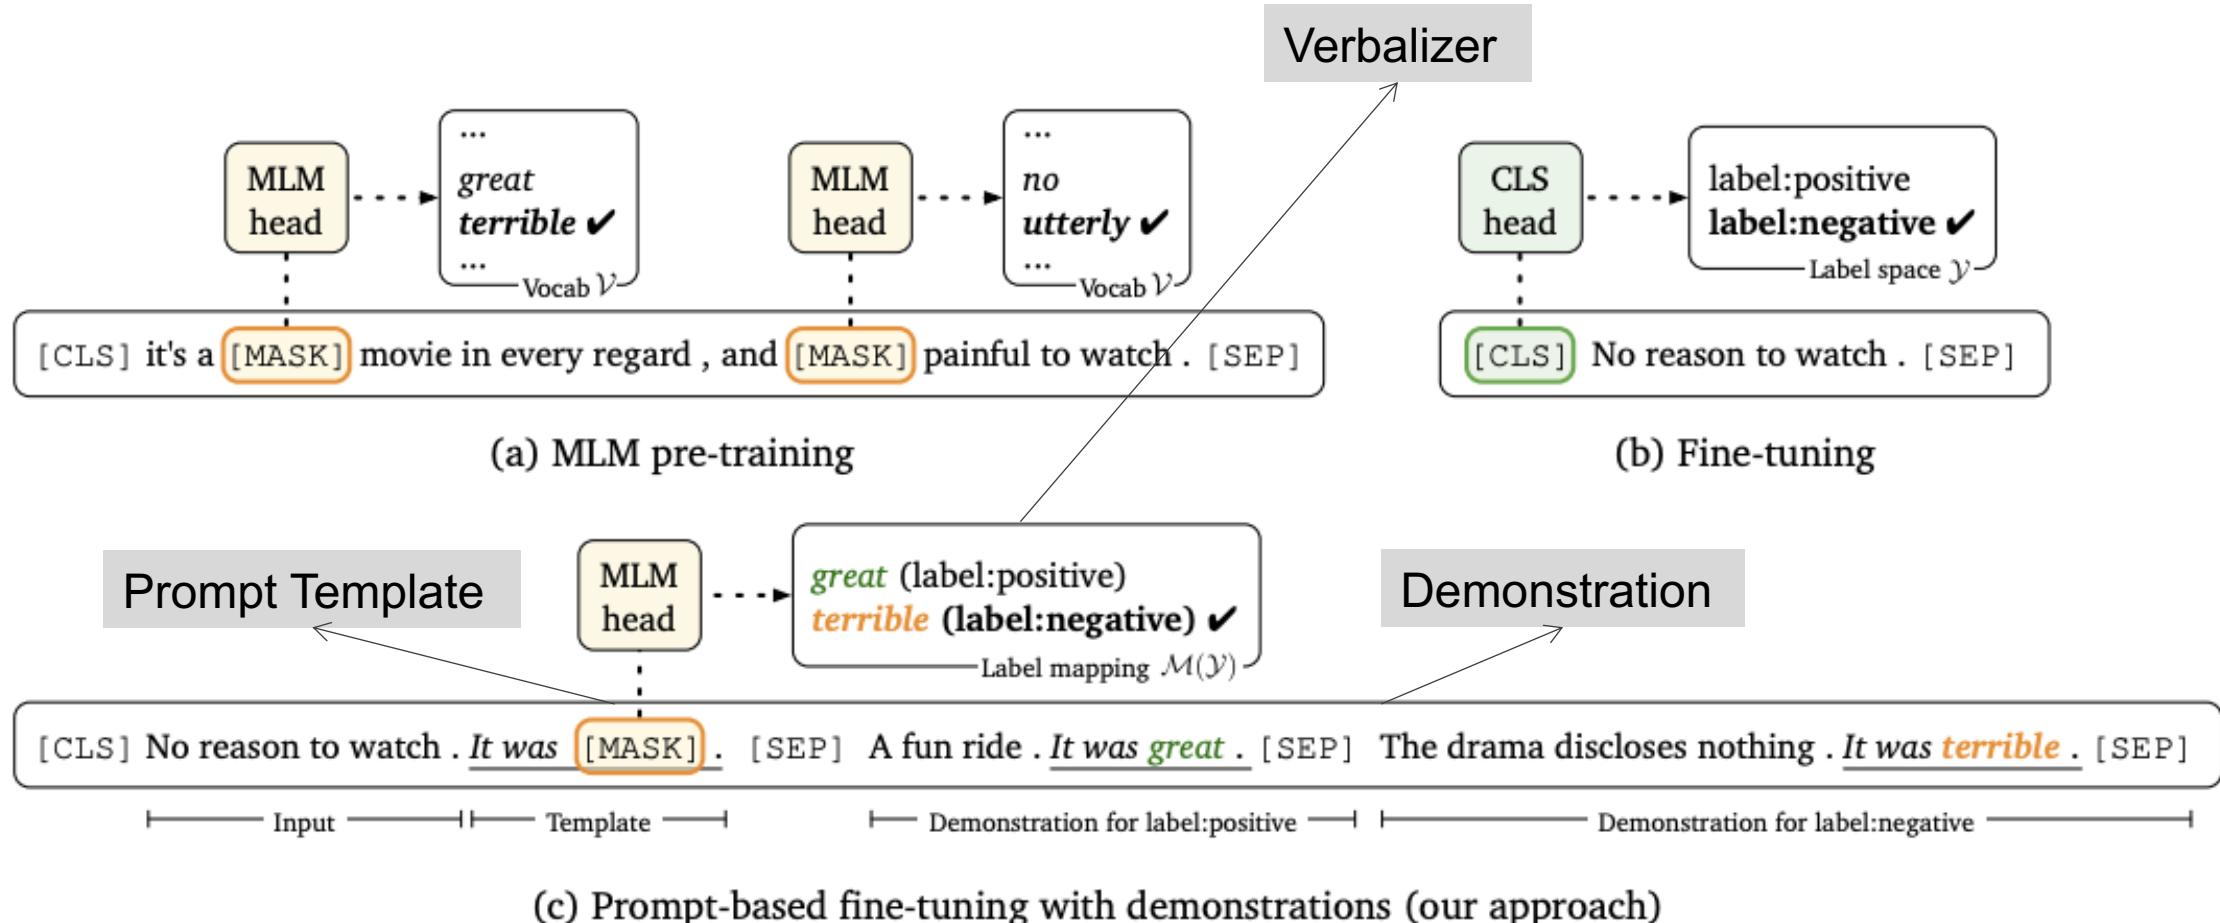
Prompt learning from GPT3 : tremendous promise in few-shot/zero-shot NLP

## Few-shot

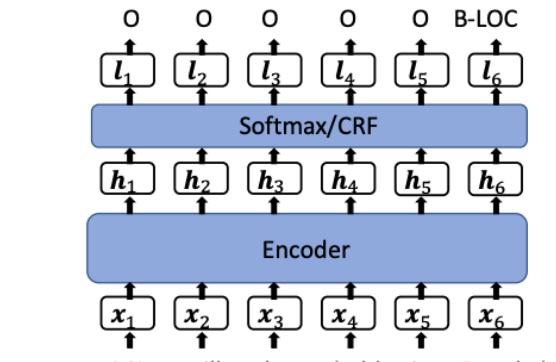
In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.



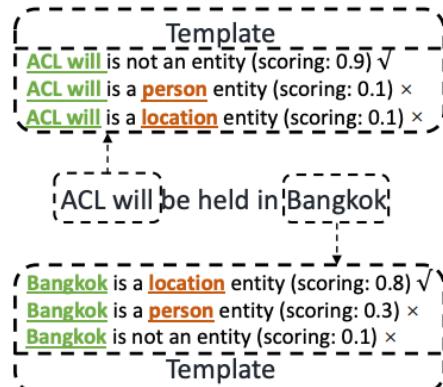
## Prompt learning for data-efficient NLP



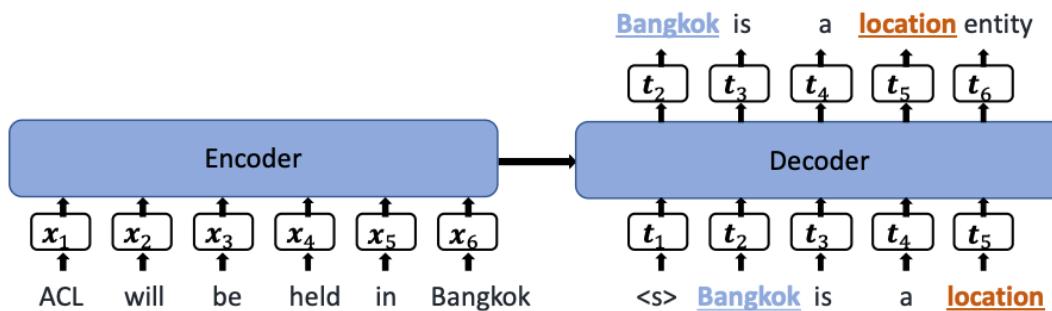
## Prompt learning for data-efficient NER



(a) Traditional sequence labeling method.



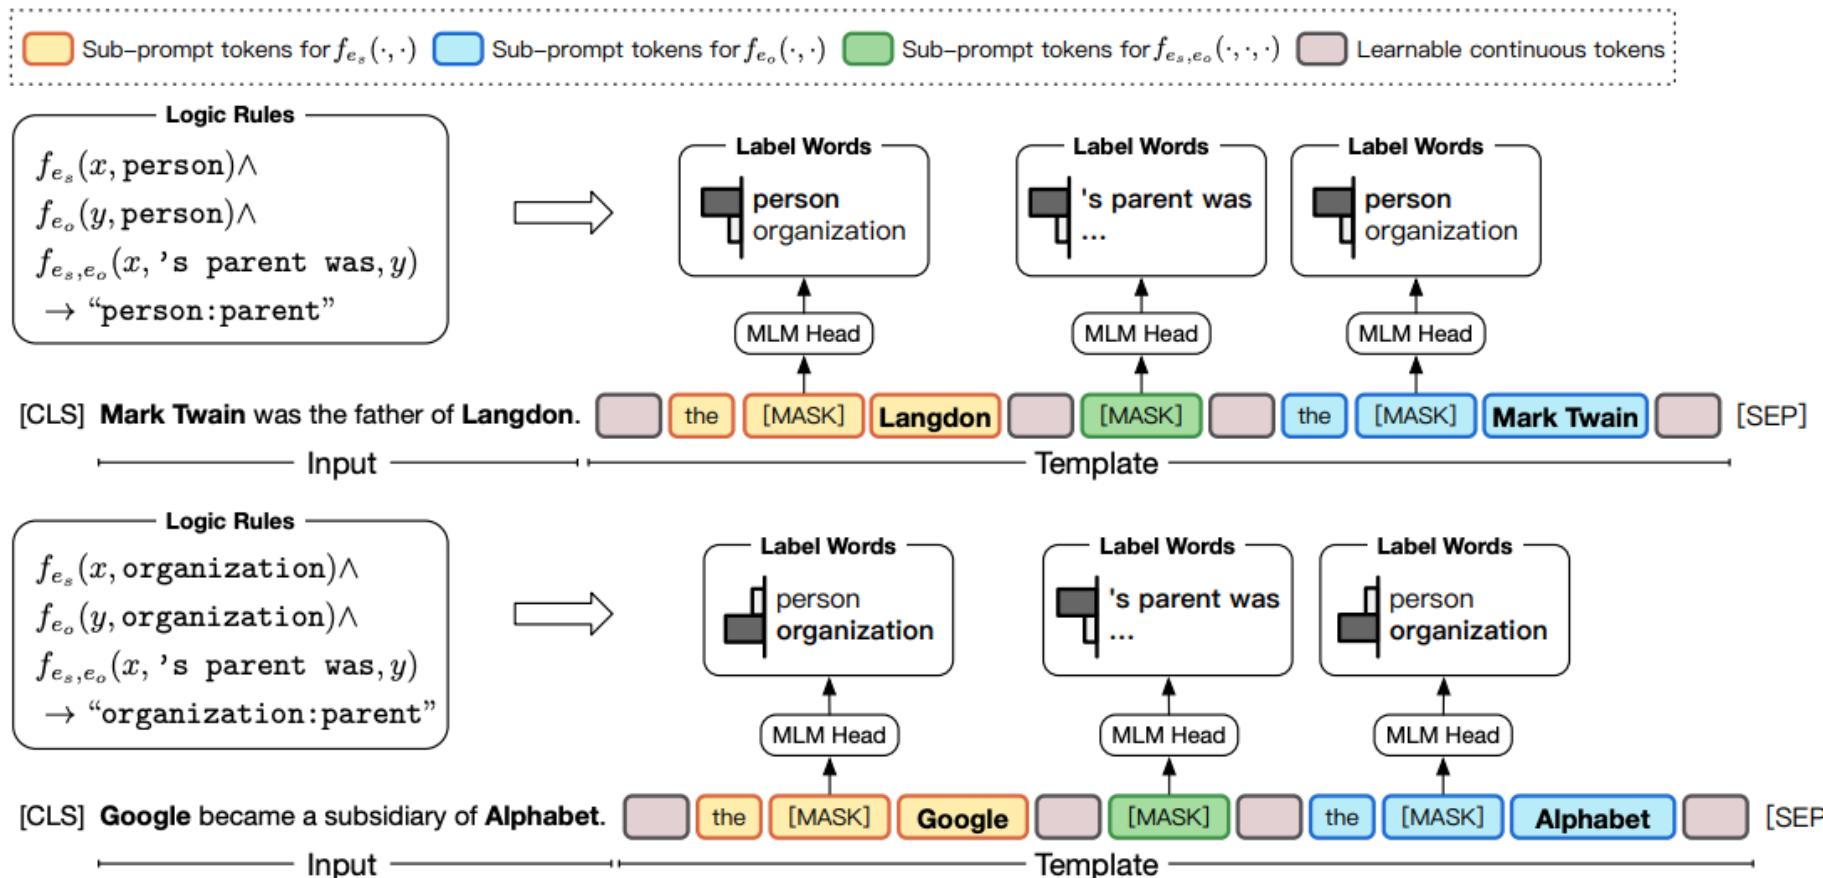
(b) Inference of template-based method.



(c) Training of template-based method. The template we use here is " $\langle x_{i:j} \rangle$  is a  $\langle y_k \rangle$  entity".

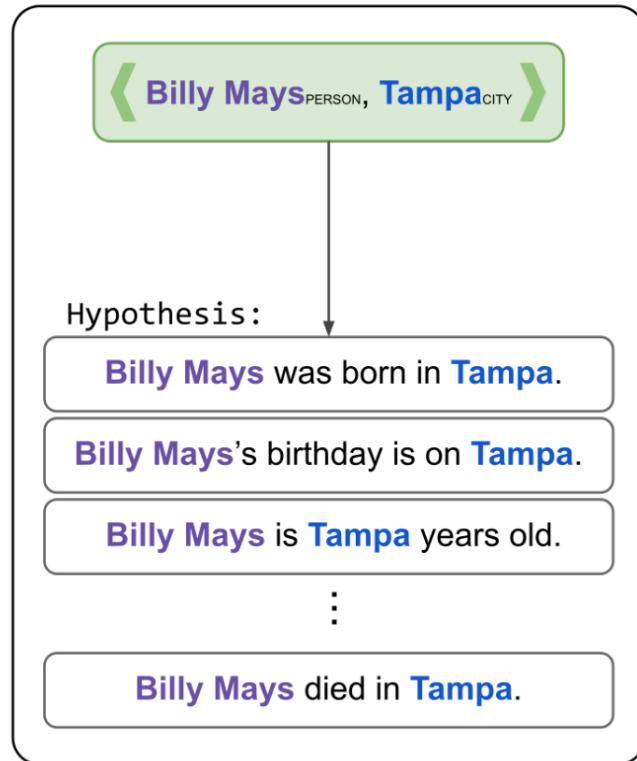
MIT Movie							
Source	Methods	10	20	50	100	200	500
None	Sequence Labeling BERT Template-based BART	25.2 37.3	42.2 48.5	49.64 52.2	50.7 56.3	59.3 62.0	74.4 74.9
CoNLL03	Wiseman and Stratos (2019)	3.1	4.5	4.1	5.3	5.4	8.6
	Ziyadi et al. (2020)	40.1	39.5	40.2	40.0	40.0	39.5
	Huang et al. (2020)*	36.4	36.8	38.0	38.2	35.4	38.3
	Sequence Labeling BERT	28.3	45.2	50.0	52.4	60.7	76.8
	Sequence Labeling BART	13.6	30.4	47.8	49.1	55.8	66.9
	Template-based BART	42.4	54.2	59.6	65.3	69.6	80.3
MIT Restaurant							
None	Sequence Labeling BERT Template-based BART	21.8 46.0	39.4 57.1	52.7 58.7	53.5 60.1	57.4 62.8	61.3 65.0
CoNLL03	Wiseman and Stratos (2019)	4.1	3.6	4.0	4.6	5.5	8.1
	Ziyadi et al. (2020)	27.6	29.5	31.2	33.7	34.5	34.6
	Huang et al. (2020)	46.1	48.2	49.6	50.0	50.1	
	Sequence Labeling BERT	27.2	40.9	56.3	57.4	58.6	75.3
	Sequence Labeling BART	8.8	11.1	42.7	45.3	47.8	58.2
	Template-based BART	53.1	60.3	64.1	67.3	72.2	75.7
ATIS							
None	Sequence Labeling BERT Template-based BART	44.1 71.7	76.7 79.4	90.7 92.6	-	-	-
CoNLL03	Wiseman and Stratos (2019)	6.7	8.8	11.1	-	-	-
	Ziyadi et al. (2020)	17.4	19.8	22.2	-	-	-
	Huang et al. (2020)	71.2	74.8	76.0	-	-	-
	Sequence Labeling BERT	53.9	78.5	92.2	-	-	-
	Sequence Labeling BART	51.3	74.4	89.9	-	-	-
	Template-based BART	77.3	88.9	93.5	-	-	-

## Prompt learning for data-efficient RE

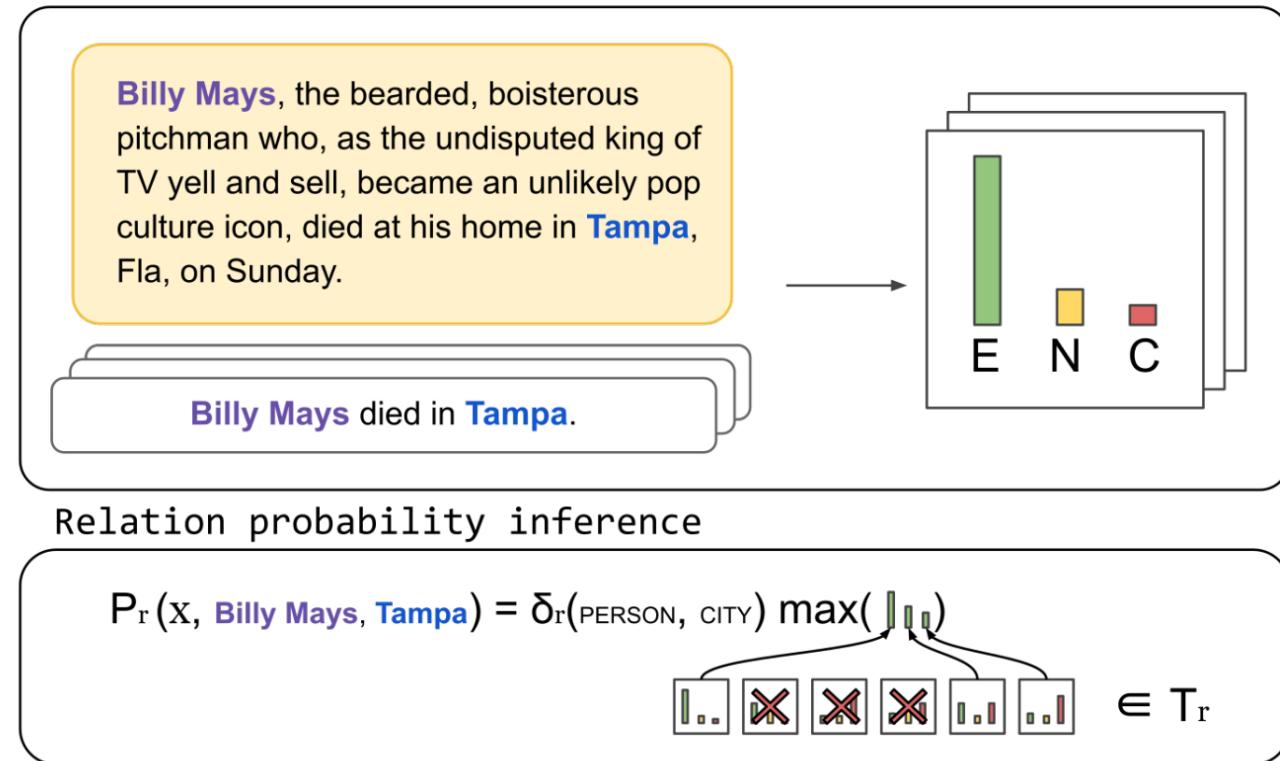


## Label Verbalization and Entailment

Verbalizer



NLI Model



## Prompting language models to generate structured texts

Relation	Sentence
Sibling	She was the mother of <b>Michael</b> and <b>Joel Douglas</b> .
Manufacturer	In late 2012, <b>Samsung</b> announced its <b>NX300</b> camera.
Architect	His <b>house</b> was designed by <b>Henry Hob Richardson</b> .

(a) Annotation samples of seen relations for training.

Relation	Sentence
Military Rank	Their grandson was <b>Group Captain Nicolas Tindal</b> .
Position Played	Made <b>Chad Brown</b> the highest paid <b>linebacker</b> in NFL.
Record Label	Deadsy signed onto <b>Immortal Records</b> to release " <b>Phantasmagore</b> ".

(b) Annotation samples of unseen relations for evaluation.

Relation	Sentence
Military Rank	She is the wife of <b>Lieutenant Colonel George Hocham</b> .
Position Played	However, it was <b>Dario Argentino</b> who defended the <b>midfield</b> .
Record Label	" <b>The Sun</b> " was first recorded by <b>Pavement</b> in 1982.

(c) Generated synthetic samples of unseen relations.

---

**Algorithm 1 RelationPrompt:** Prompting language models to generate synthetic data for ZeroRTE.

---

**Define:**

Dataset  $D = (\text{Sentences } S, \text{Triplets } T, \text{Labels } Y)$

**Require:** Train Dataset  $D_s$ , Test Dataset  $D_u$ , Relation Generator  $M_g$ , Relation Extractor  $M_e$ .

**Ensure:**  $Y_s \cap Y_u = \emptyset$ .

- 1:  $M_{g,finetune} \leftarrow Train(M_g, D_s)$
  - 2:  $M_{e,finetune} \leftarrow Train(M_e, D_s)$
  - 3:  $D_{synthetic} \leftarrow Generate(M_{g,finetune}, Y_u)$
  - 4:  $M_{e,final} \leftarrow Train(M_{e,finetune}, D_{synthetic})$
  - 5:  $\hat{T}_u \leftarrow Predict(M_{e,final}, S_u)$
  - 6: **return** Extracted Triplets  $\hat{T}_u$
-

## Challenges

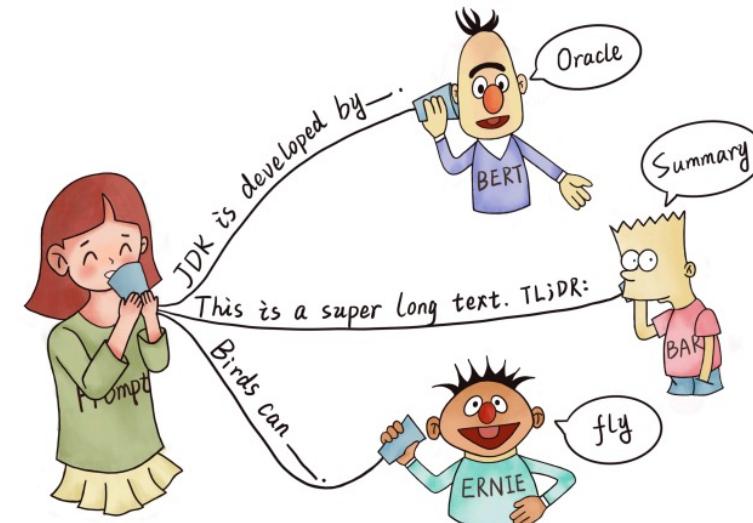
### □ Prompt Design with Complex Semantics

The relation labels of *isSpouseOf* and *org* : *city\_of\_headquarters* cannot specify a single suitable label word in the vocabulary.

### □ Inductive Bias For KG

If a pair of entities contains the semantics of “person” and “country”, the prediction probability of the [MASK] on the relation “org:city\_of\_headquarters” will be lower.

**isSpouseOf**  
Jack is married to the microbiologist known as Dr. Germ in the USA.  
Entity Pair



Prior knowledge (Schema) in KG as a prompt to stimulate the potential knowledge in PLMs for RE

## Challenges

### □ Prompt Design with Complex Semantics

The relation labels of *isSpouseOf* and *org : city\_of\_headquarters* cannot specify a single suitable label word in the vocabulary.

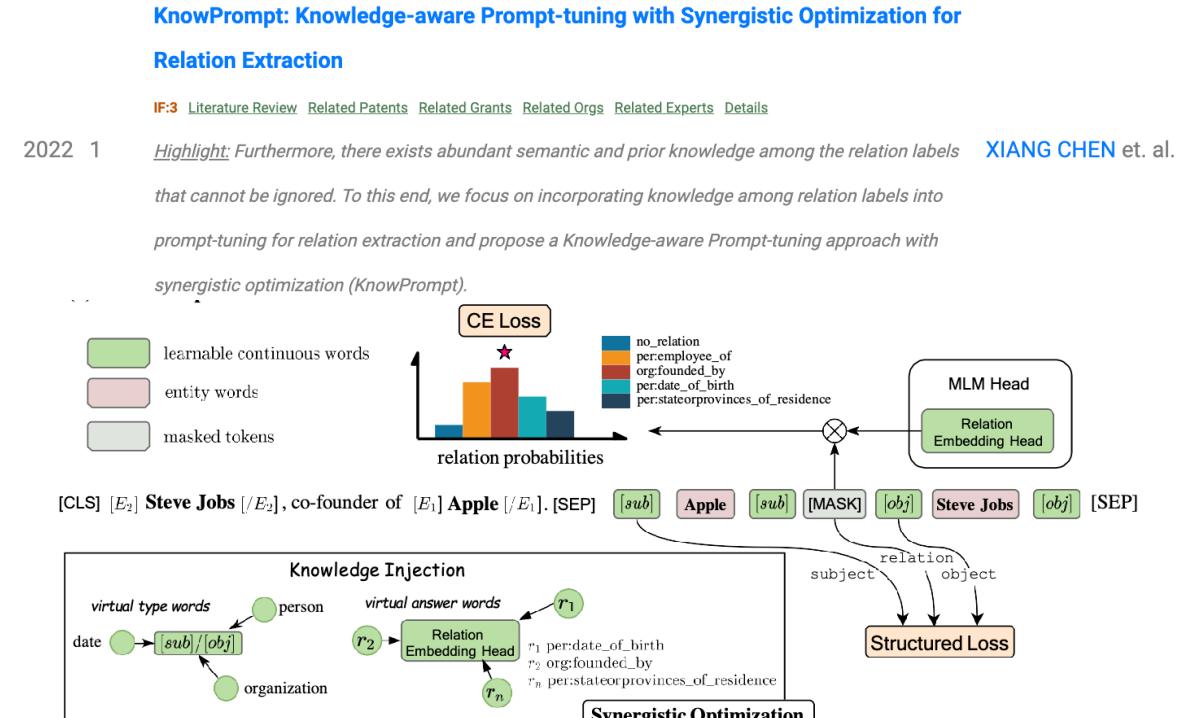
### □ Inductive Bias For KG

If a pair of entities contains the semantics of “person” and “country”, the prediction probability of the [MASK] on the relation “org:city\_of\_headquarters” will be lower

TABLE 1: Most Influential WWW Papers

YEAR RANK PAPER

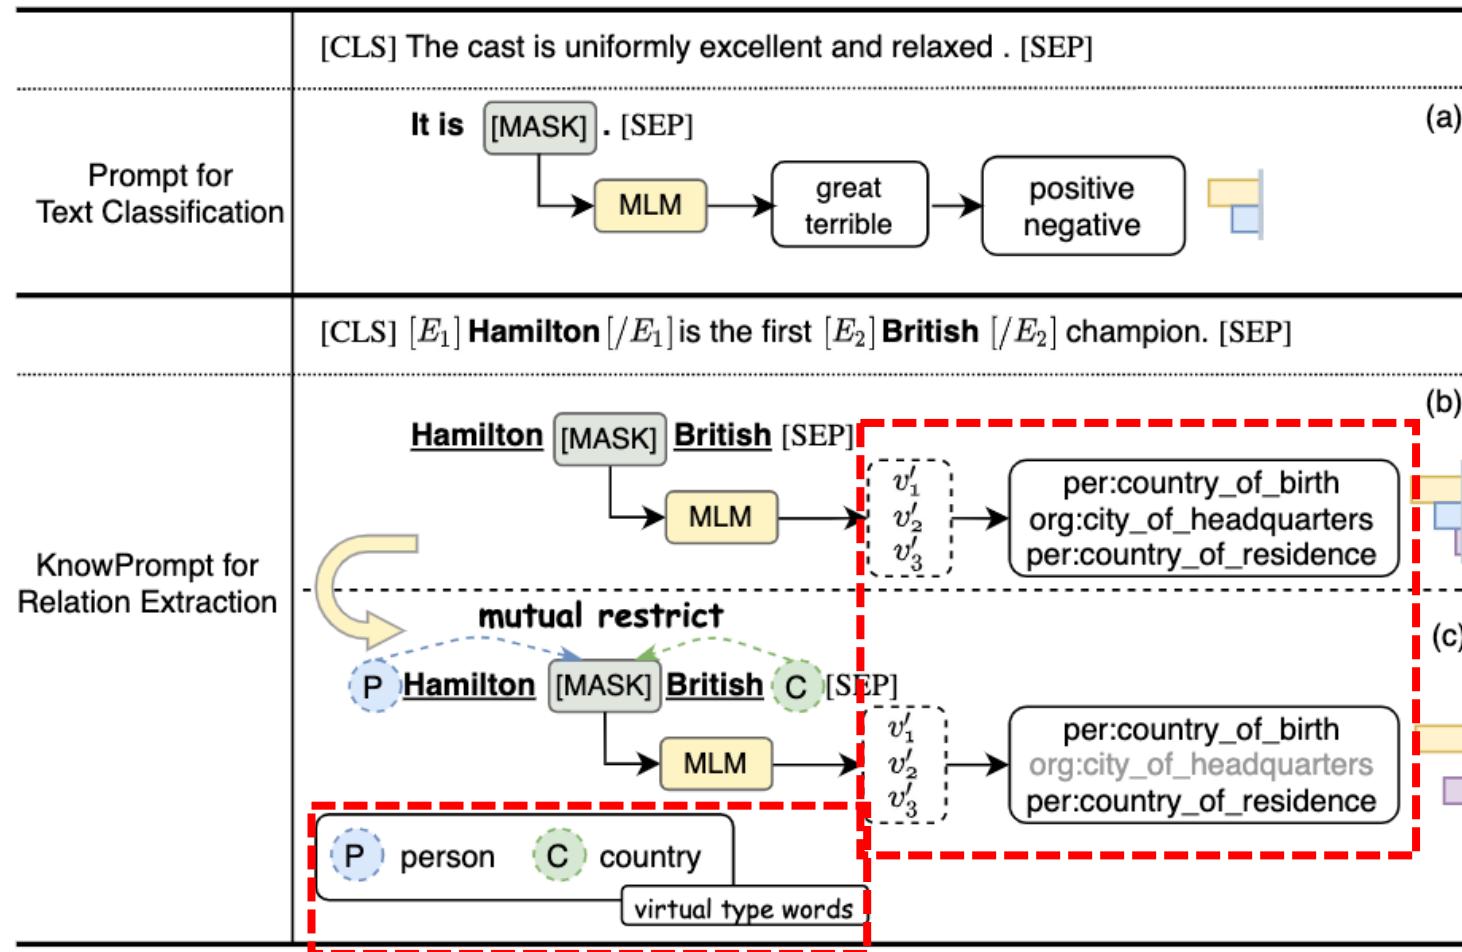
AUTHOR(S)



KnowPrompt'WWW22

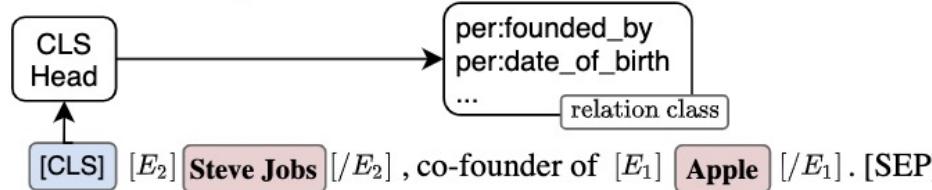
Prior knowledge (Schema) in KG as a prompt to stimulate the potential knowledge in PLMs for RE

**From discrete to continuous:** virtual embeddings as an alternative to entity and relation types

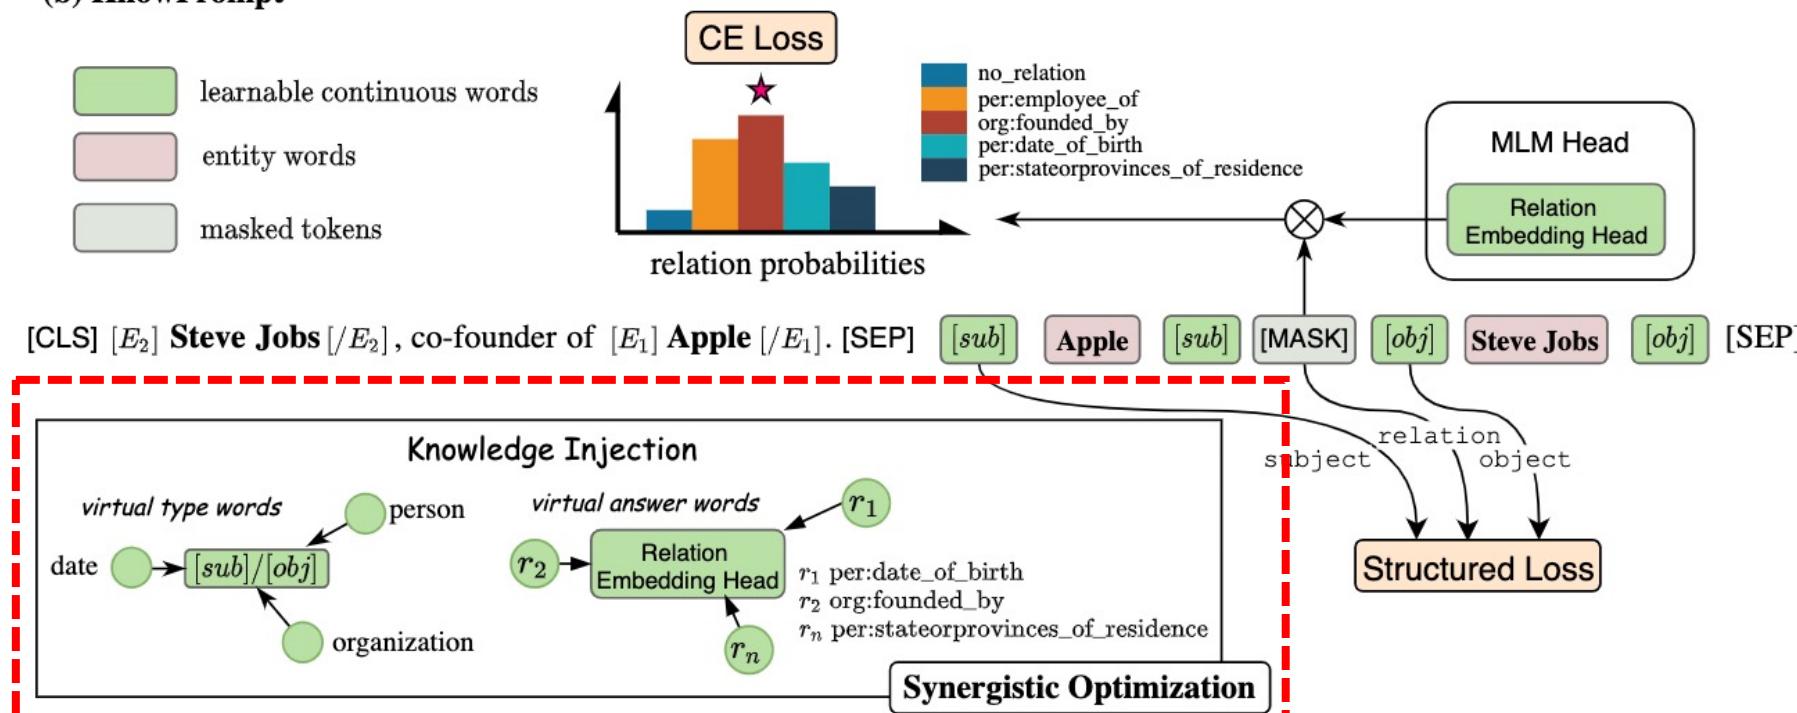


**Knowledge injection:** constructing prior structural constraints on entity and relation

(a) Fine-Tuning for RE



(b) KnowPrompt



KnowPrompt can obtain better performance on **full-supervised** setting

Code is available in <https://github.com/zjunlp/KnowPrompt>

Methods	<i>Standard Supervised Setting</i>						
	Extra Data	SemEval	DialogRE†	TACRED	TACRED-Revisit	Re-TACRED	
Fine-tuning pre-trained models							
FINE-TUNING-[ROBERTA]	w/o	87.6	57.3	68.7	76.0	84.9	
SPANBERT [30]	w/	-	-	70.8	78.0	85.3	
KNOWBERT [38]	w/	89.1	-	71.5	79.3	89.1	
LUKE [52]	w/	-	-	72.7	80.6	-	
MTB [3]	w/	89.5	-	70.1	-	-	
GDPNET [51]	w/o	-	64.9	71.5	79.3	-	
DUAL [2]	w/o	-	67.3	-	-	-	
Prompt-tuning pre-trained models							
PTR-[ROBERTA] [22]	w/o	89.9	63.2	72.4	81.4	90.9	
<b>KNOWPROMPT-[ROBERTA]</b>	w/o	<b>90.2 (+0.3)</b>	<b>68.6 (+5.4)</b>	<b>72.4 (-0.3)</b>	<b>82.4 (+1.0)</b>	<b>91.3 (+0.4)</b>	

# KnowPrompt: Experiments

KnowPrompt can achieve large improvement on **low-resource** setting

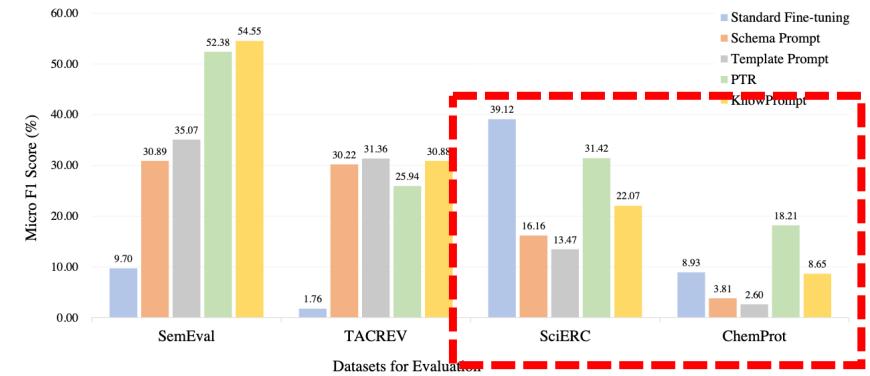
Code is available in <https://github.com/zjunlp/KnowPrompt>

<i>Low-Resource Setting</i>							
Split	Methods	SemEval	DialogRE†	TACRED	TACRED-Revisit	Re-TACRED	Average
K=8	FINE-TUNING	41.3	29.8	12.2	13.5	28.5	25.1
	GDPNET	42.0	28.6	11.8	12.3	29.0	24.7
	PTR	70.5	35.5	28.1	28.7	51.5	42.9
	<b>KNOWPROMPT</b>	<b>74.3 (+33.0)</b>	<b>43.8 (+14.0)</b>	<b>32.0 (+19.8)</b>	<b>32.1 (+18.6)</b>	<b>55.3 (+26.8)</b>	<b>47.5 (+22.4)</b>
K=16	FINE-TUNING	65.2	40.8	21.5	22.3	49.5	39.9
	GDPNET	67.5	42.5	22.5	23.8	50.0	41.3
	PTR	81.3	43.5	30.7	31.4	56.2	48.6
	<b>KNOWPROMPT</b>	<b>82.9 (+17.7)</b>	<b>50.8 (+10.0)</b>	<b>35.4 (+13.9)</b>	<b>33.1 (+10.8)</b>	<b>63.3 (+13.8)</b>	<b>53.1 (+13.2)</b>
K=32	FINE-TUNING	80.1	49.7	28.0	28.2	56.0	48.4
	GDPNET	81.2	50.2	28.8	29.1	56.5	49.2
	PTR	84.2	49.5	32.1	32.4	62.1	52.1
	<b>KNOWPROMPT</b>	<b>84.8 (+4.7)</b>	<b>55.3 (+3.6)</b>	<b>36.5 (+8.5)</b>	<b>34.7 (+6.5)</b>	<b>65.0 (+9.0)</b>	<b>55.3 (+6.9)</b>

# KnowPrompt: More results on LREBench

KnowPrompt obtain better performance on **most** datasets

Code is available in <https://github.com/zjunlp/LREBench>



## □ Finding 1

Prompt-based tuning largely outperforms standard fine-tuning

## □ Finding 3

Balancing methods may fail on those challenging datasets

Dataset	Metric	Fine-Tune						Prompt					
		Normal			Balance		DA	ST	Normal			Balance	
		8-shot	10%	100%	10%	100%	10%	10%	100%	10%	100%	10%	10%
SemEval	MaF1	2.69	34.63	81.88	41.84	82.44	69.84	60.10	58.89	44.71	83.40	54.54	<b>83.20</b>
	MiF1	9.70	54.61	89.10	58.26	89.44	78.98	74.12	62.52	69.90	90.01	76.53	92.31
TACREV	MaF1	1.02	47.32	63.41	48.64	<b>63.38</b>	50.68	48.84	29.46	61.40	67.08	63.09	69.63
	MiF1	1.76	65.43	71.68	67.19	73.86	65.99	66.89	30.88	77.00	78.30	<b>76.25</b>	81.41
Wiki80	MaF1	37.89	37.82	71.31	44.27	72.26	52.79	52.79	63.99	83.72	63.40	60.86	-
	MiF1	44.85	46.50	72.8	-	-	-	-	67.86	83.86	66.96	65.04	-
SciERC	MaF1	10.41	10.31	-	-	-	-	-	55	84.83	65.98	56.94	-
	MiF1	39.12	54.66	-	-	-	-	-	76.90	90.04	79.92	76.32	-
ChemProt	MaF1	2.18	27.96	47.35	33.56	77.35	56.17	56.17	36.43	<b>47.16</b>	38.99	<b>47.07</b>	37.44
	MiF1	8.93	49.20	68.81	54.98	<b>68.77</b>	56.58	54.17	8.65	56.96	69.14	57.28	<b>69.12</b>
DialogRE	MaF1	1.13	2.17	25.31	5.84	27.28	-	<b>0.00</b>	44.96	45.42	69.99	49.68	<b>69.73</b>
	MiF1	3.92	23.37	41.52	24.53	<b>41.24</b>	-	<b>0.00</b>	45.70	58.66	76.20	<b>58.16</b>	<b>75.88</b>
DuIE2.0	MaF1	36.62	90.46	95.01	92.91	96.00	-	<b>89.27</b>	80.31	93.48	95.73	93.70	96.01
	MiF1	39.00	94.42	96.22	94.46	<b>96.13</b>	-	<b>93.81</b>	82.14	95.09	96.43	95.23	96.44
CMeIE	MaF1	13.68	62.30	84.37	67.22	86.31	-	<b>58.46</b>	36.54	67.59	86.42	67.84	86.68
	MiF1	17.05	79.82	90.48	80.43	90.56	-	<b>78.92</b>	38.02	83.38	92.08	83.40	92.14

Eight datasets of single-sentence & document with various languages

## □ Finding 2

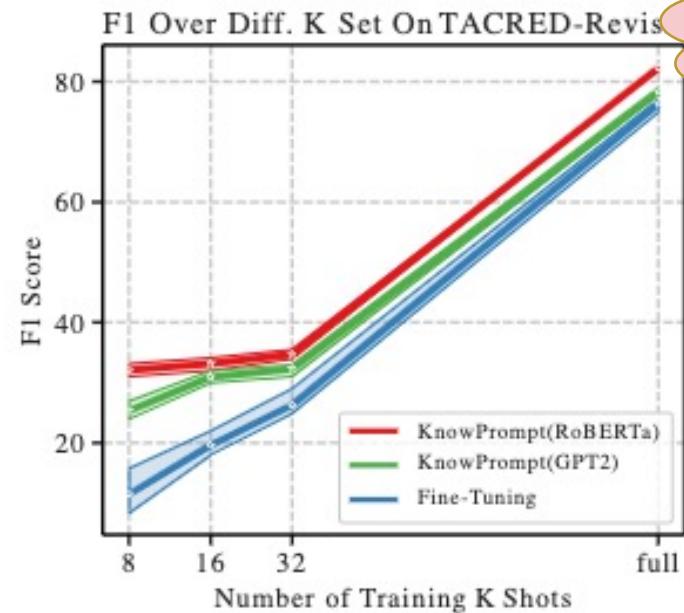
Data augmentation achieves much gain on RE

## □ Finding 4

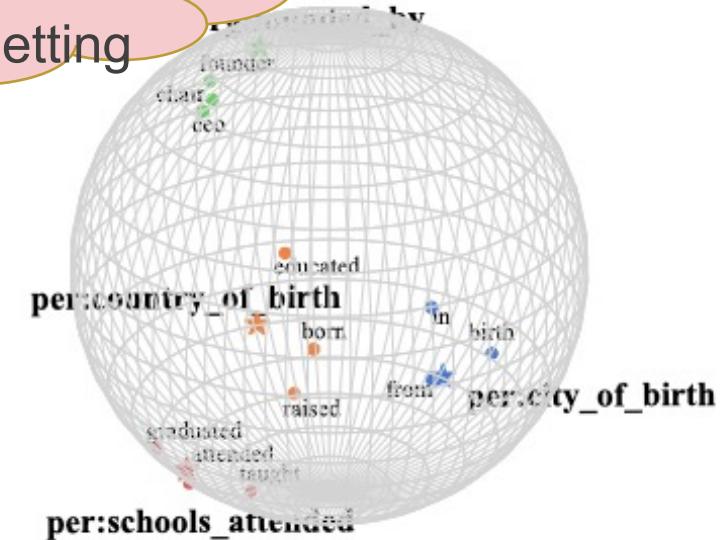
RE systems may fail on cross-sentence contexts and among multiple triples

# KnowPrompt: Analysis

Input Example of our KnowPrompt	Top 3 words around [sub]	Top 3 words around [obj]
x:[CLS] It sold [E <sub>1</sub> ] <b>ALICO</b> [/E <sub>1</sub> ] to [E <sub>2</sub> ] <b>MetLife Inc</b> [E <sub>2</sub> ] for \$ 162 billion. [SEP] [sub] ALICO [sub] [MASK] [obj] MetLife Inc [obj]. [SEP] y: "org : member_of"	organization group corporation	company plc organization
x: [CLS] [E <sub>1</sub> ] <b>Ismael Rukwago</b> [/E <sub>1</sub> ], a senior [E <sub>2</sub> ] ADF [E <sub>2</sub> ] commander, denied any involvement. [SEP] [sub] Ismael Rukwago [sub] [MASK] [obj] ADF [obj]. [SEP] y: "per : employee_of"	person commander colonel	intelligence organization command

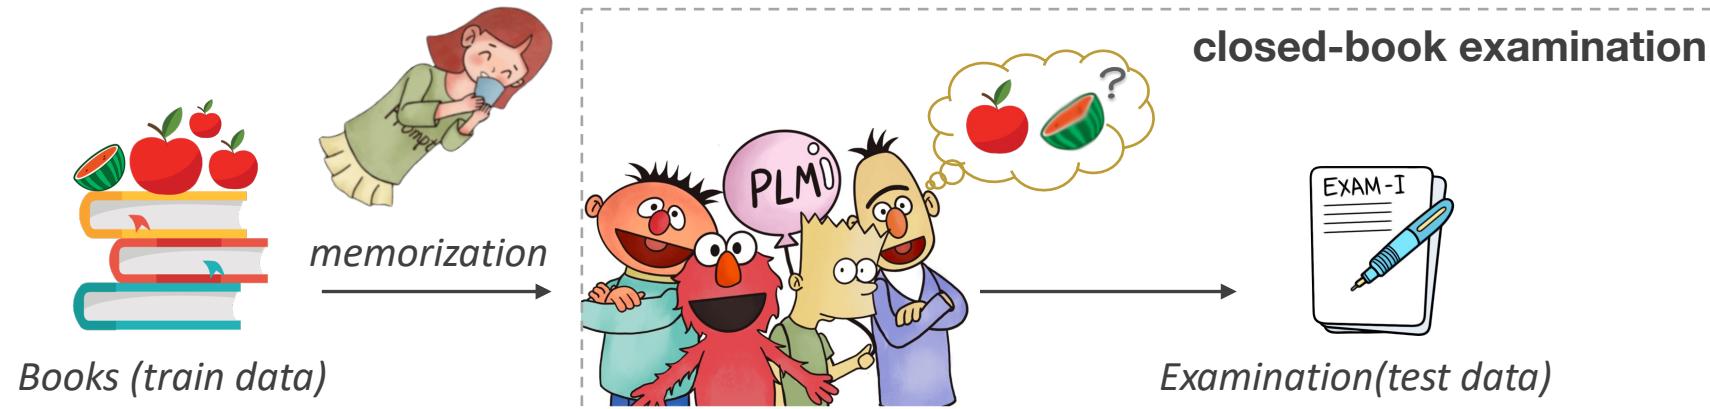


Knowprompt fail  
on **extreme** low-  
resource setting

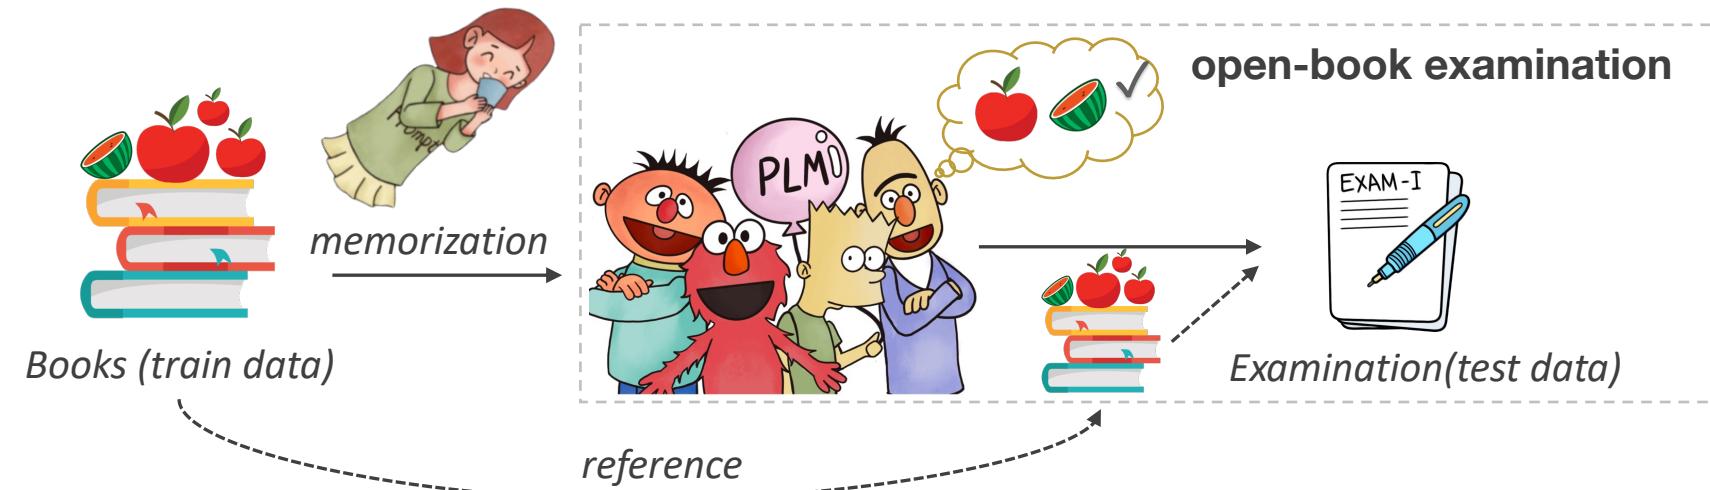


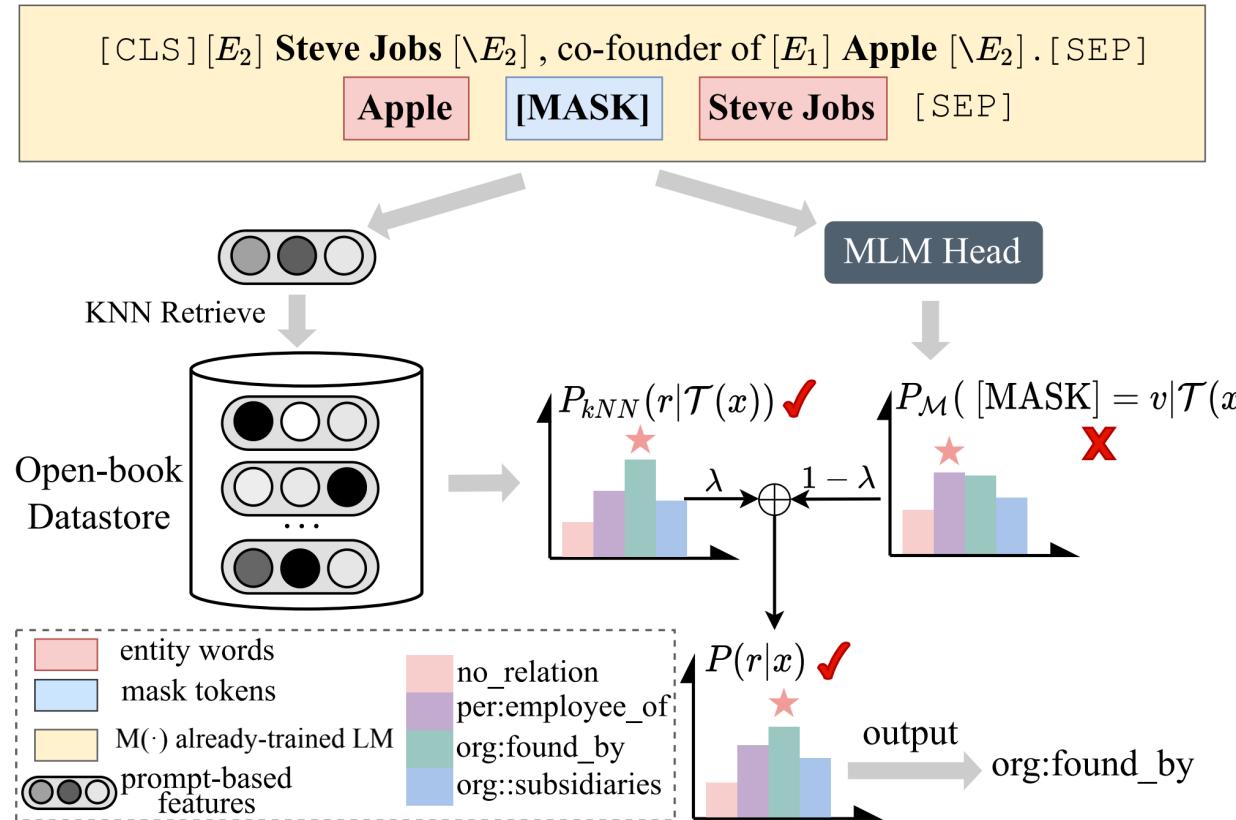
# What's Missing? More knowledge in Training Data

Prompt-tuning



**Retrieval-enhanced**  
prompt tuning

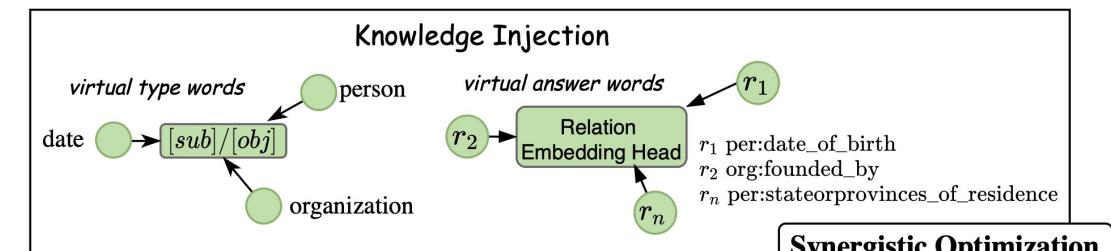




$$P_M([\text{MASK}] = v | \mathcal{T}(x))$$

$$P_{kNN}(r | \mathcal{T}(x)) \propto \sum_{(c_i, r_i) \in \mathcal{N}} 1_{r=r_i} \exp(-d(h_x, h_{c_i})).$$

$$P(r | x) = \lambda P_{kNN}(r | \mathcal{T}(x)) + (1 - \lambda) P_M([\text{MASK}] = v | \mathcal{T}(x)).$$



Soft label words & Knowledge Injection (from *KnowPrompt*<sup>[1]</sup>)

# Experiments

Code is available in <https://github.com/zjunlp/PromptKG/tree/main/research/RetrievalRE>

**Table 1: Model performance of RE models in the low-resource setting. We report the mean and standard deviation performance of micro  $F_1$  scores (%) over 5 different splits. The best numbers are highlighted in each column.**

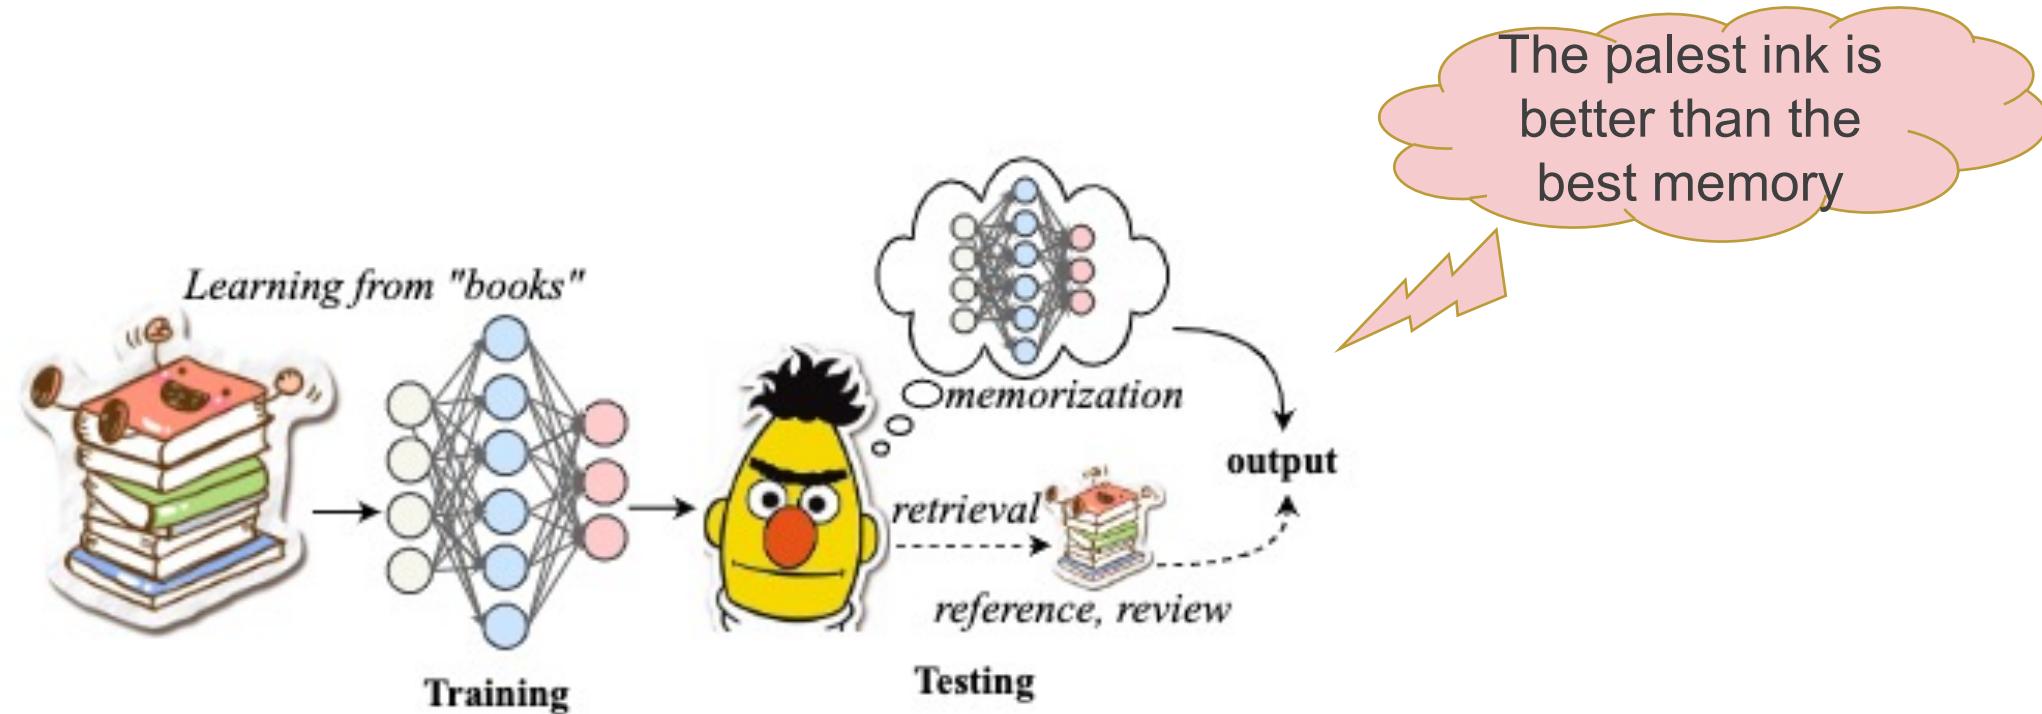
Source	Model	SemEval			TACRED			TACREV		
		K=1	K=5	K=16	K=1	K=5	K=16	K=1	K=5	K=16
None	FINE-TUNING	18.5 ( $\pm$ 1.4)	41.5 ( $\pm$ 2.3)	66.1 ( $\pm$ 0.4)	7.6 ( $\pm$ 3.0)	16.6 ( $\pm$ 2.1)	26.8 ( $\pm$ 1.8)	7.2 ( $\pm$ 1.4)	16.3 ( $\pm$ 2.1)	25.8 ( $\pm$ 1.2)
	GDPNET	10.3 ( $\pm$ 2.5)	42.7 ( $\pm$ 2.0)	67.5 ( $\pm$ 0.8)	4.2 ( $\pm$ 3.8)	15.5 ( $\pm$ 2.3)	28.0 ( $\pm$ 1.8)	5.1 ( $\pm$ 2.4)	17.8 ( $\pm$ 2.4)	26.4 ( $\pm$ 1.2)
	PTR	14.7 ( $\pm$ 1.1)	53.9 ( $\pm$ 1.9)	80.6 ( $\pm$ 1.2)	8.6 ( $\pm$ 2.5)	24.9 ( $\pm$ 3.1)	30.7 ( $\pm$ 2.0)	9.4 ( $\pm$ 0.7)	26.9 ( $\pm$ 1.5)	31.4 ( $\pm$ 0.3)
	KnowPrompt	28.6 ( $\pm$ 6.2)	66.1 ( $\pm$ 8.6)	80.9 ( $\pm$ 1.6)	17.6 ( $\pm$ 1.8)	28.8 ( $\pm$ 2.0)	34.7 ( $\pm$ 1.8)	17.8 ( $\pm$ 2.2)	30.4 ( $\pm$ 0.5)	33.2 ( $\pm$ 1.4)
<b>RetrievalRE</b>		<b>33.3</b> ( $\pm$ 1.6)	<b>69.7</b> ( $\pm$ 1.7)	<b>81.8</b> ( $\pm$ 1.0)	<b>19.5</b> ( $\pm$ 1.5)	<b>30.7</b> ( $\pm$ 1.7)	<b>36.1</b> ( $\pm$ 1.2)	<b>18.7</b> ( $\pm$ 1.8)	<b>30.6</b> ( $\pm$ 0.2)	<b>35.3</b> ( $\pm$ 0.3)

**Table 3: Standard RE performance of  $F_1$  scores (%) on different test sets. Best results are bold.**

Methods	Standard Supervised Setting				
	SemEval	DialogRE	TACRED	TACREV	Re-TACRED
Fine-tuning pre-trained models					
FINE-TUNING	87.6	57.3	68.7	76.0	84.9
SPANBERT [22]	-	-	70.8	78.0	85.3
KNOWBERT [30]	89.1	-	71.5	79.3	89.1
LUKE [38]	-	-	<b>72.7</b>	80.6	-
MTB [4]	89.5	-	70.1	-	-
GDPNET [37]	-	64.9	71.5	79.3	-
DUAL [3]	-	67.3	-	-	-
Prompt-tuning pre-trained models					
PTR [13]	89.9	63.2	72.4	81.4	90.9
KNOWPROMPT	90.2	68.6	72.4	82.4	91.3
<b>RETRIEVALRE</b>	<b>90.4</b>	<b>69.4</b>	<b>72.7</b>	<b>82.7</b>	<b>91.5</b>

- Our RetrievalRE achieves improvements over all baselines in the **low-resource setting**.
- RetrievalRE can also achieve improvements of F1 scores when the training data is **exceptionally scarce**.
- Slight improvement in the **standard supervised setting**.

## Memorisation versus Generalisation



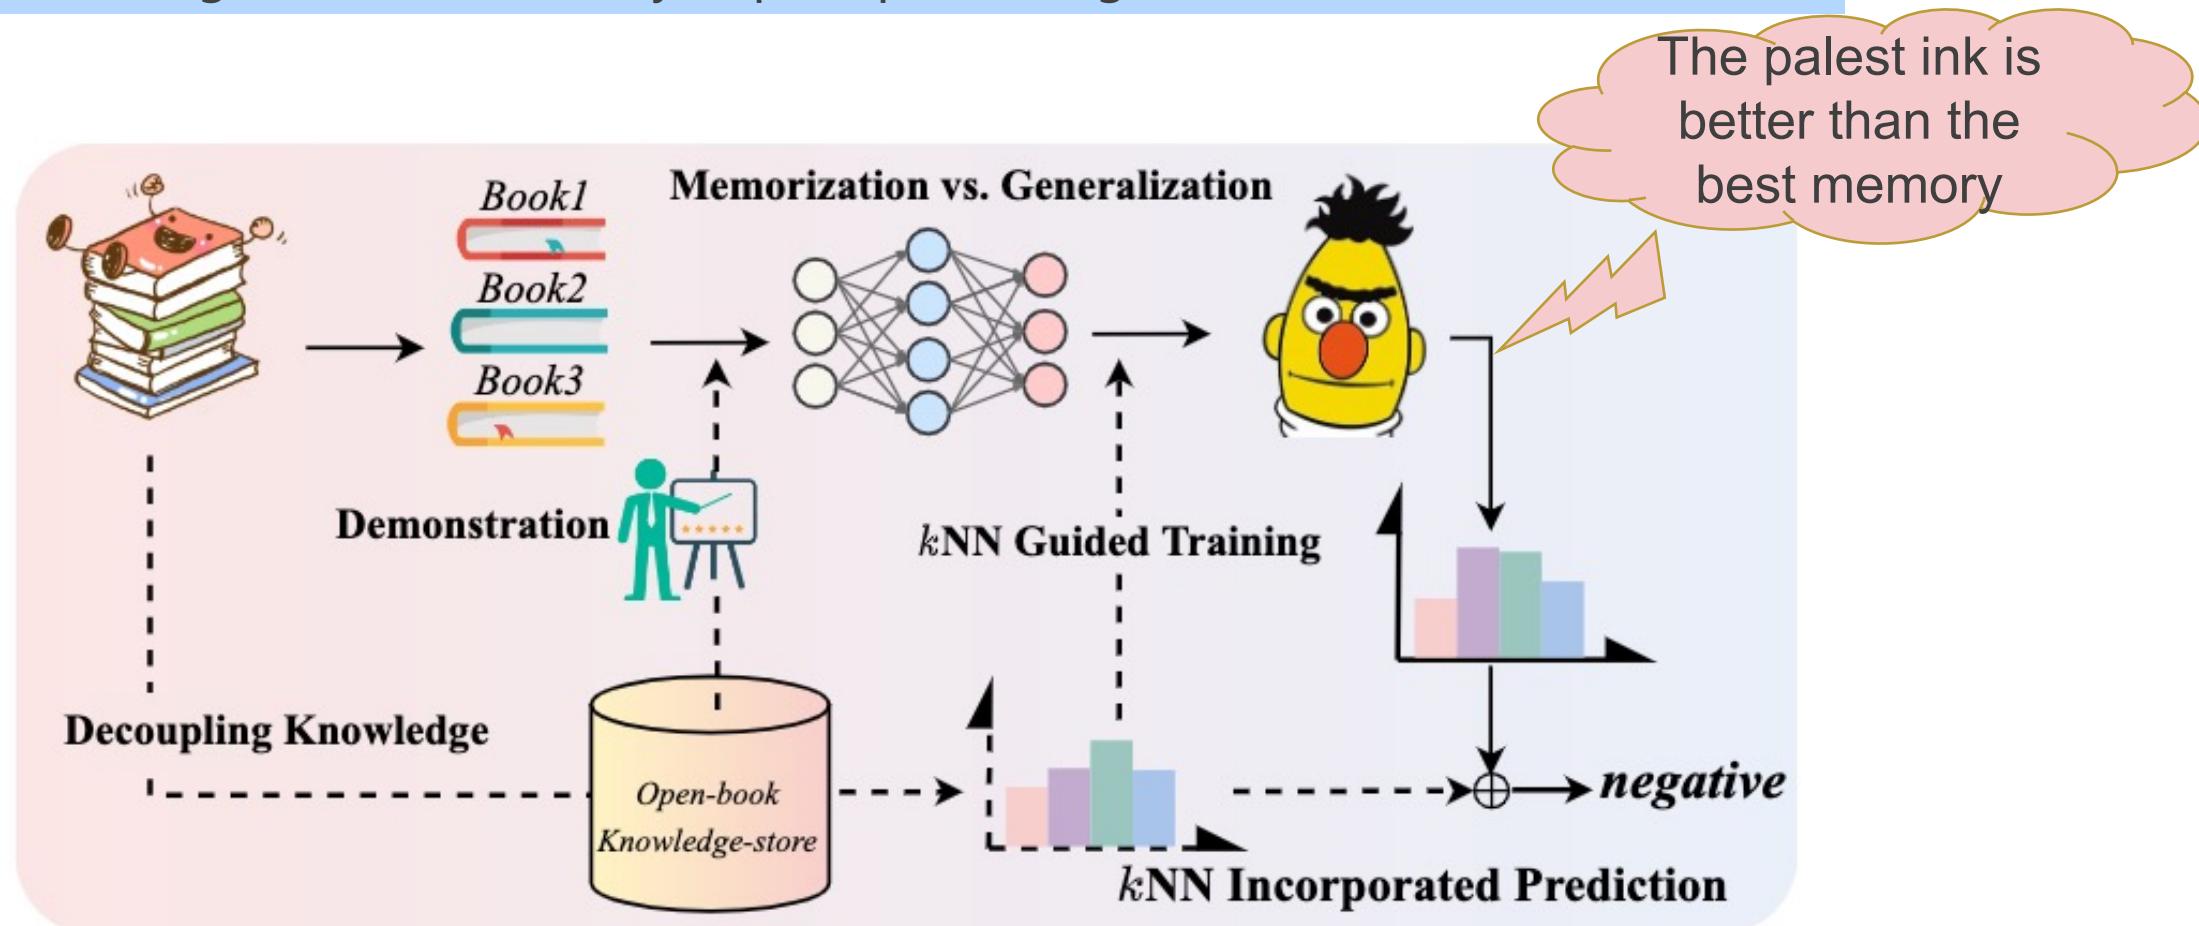
Vitaly Feldman. 2020. Does learning require memorization? a short tale about a long tail  
An Empirical Study of Memorization in NLP (ACL2022)  
Memorisation versus Generalisation in Pre-trained Language Models (ACL2022)

# Motivation: DO NOT LEARN BY ROT !



AACL IJCNLP 2022

Improve the generalization ability of prompt learning with retrieval and association.



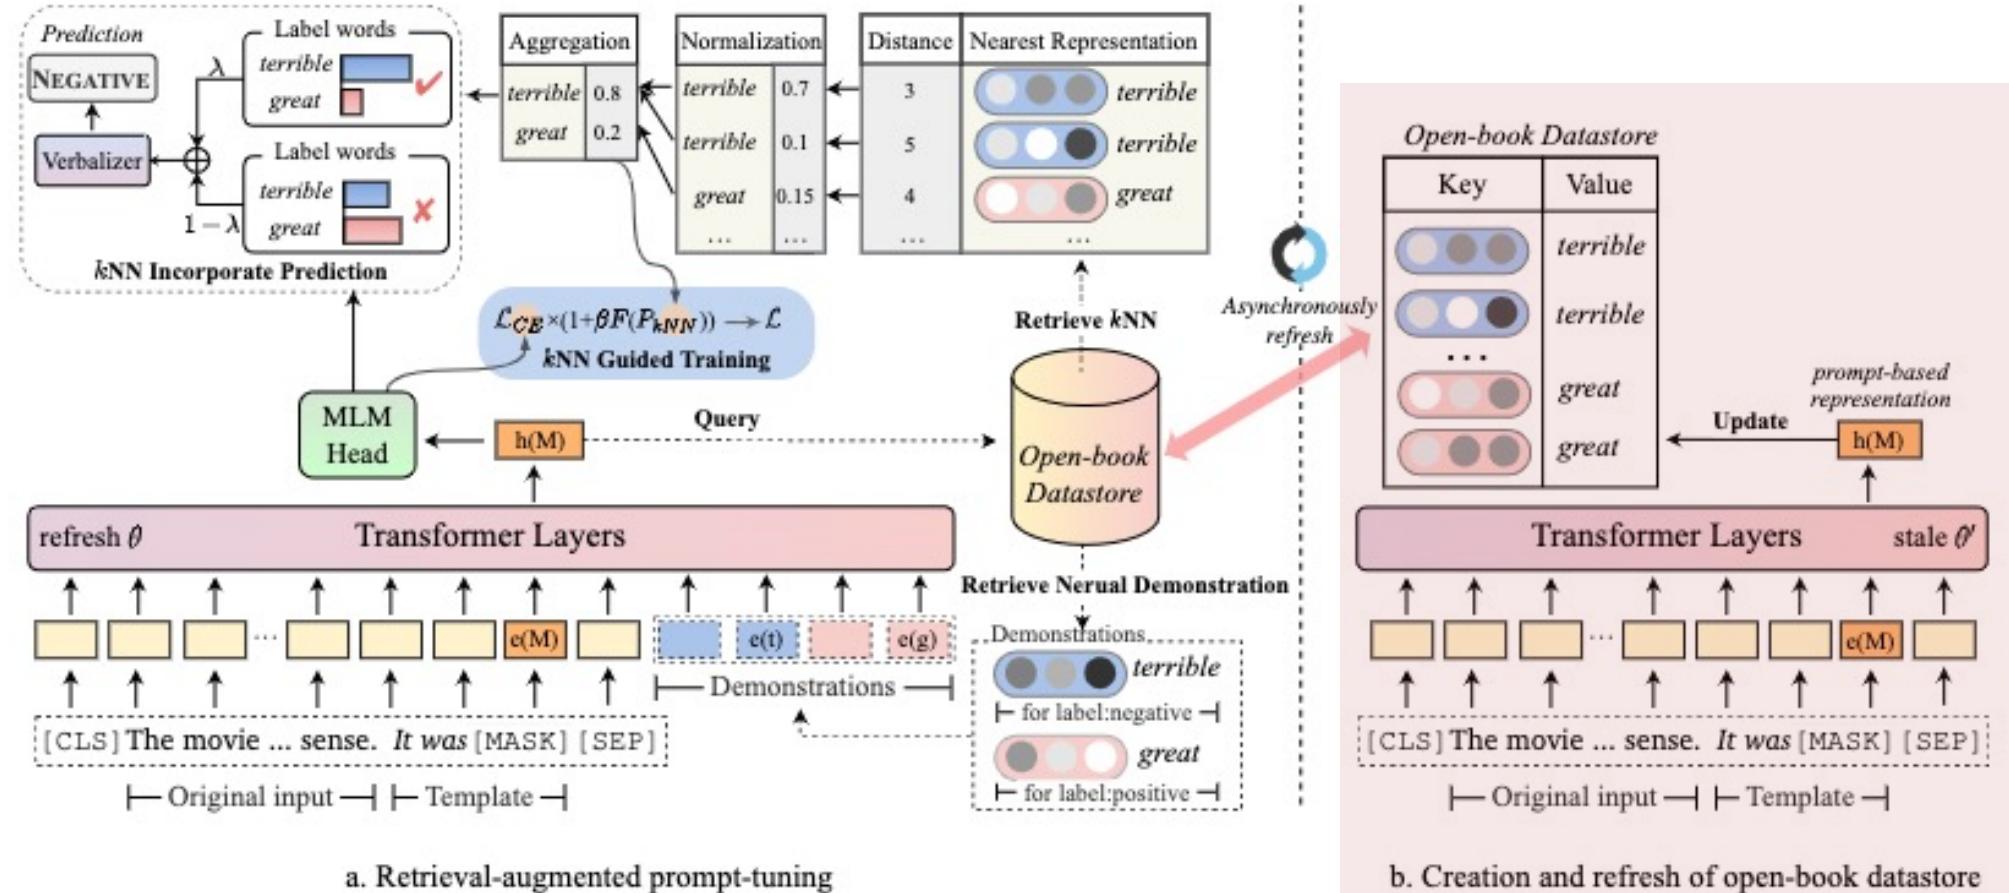


Figure 2: Overview of our RETRO\_PROMPT. Note that  $e(\cdot)$  donates the word embedding in the PLM  $\mathcal{M}$ , while “M”, “t” and “g” in the  $e(\cdot)$  specifically refers to the “[MASK]”, “terrible” and “great”.

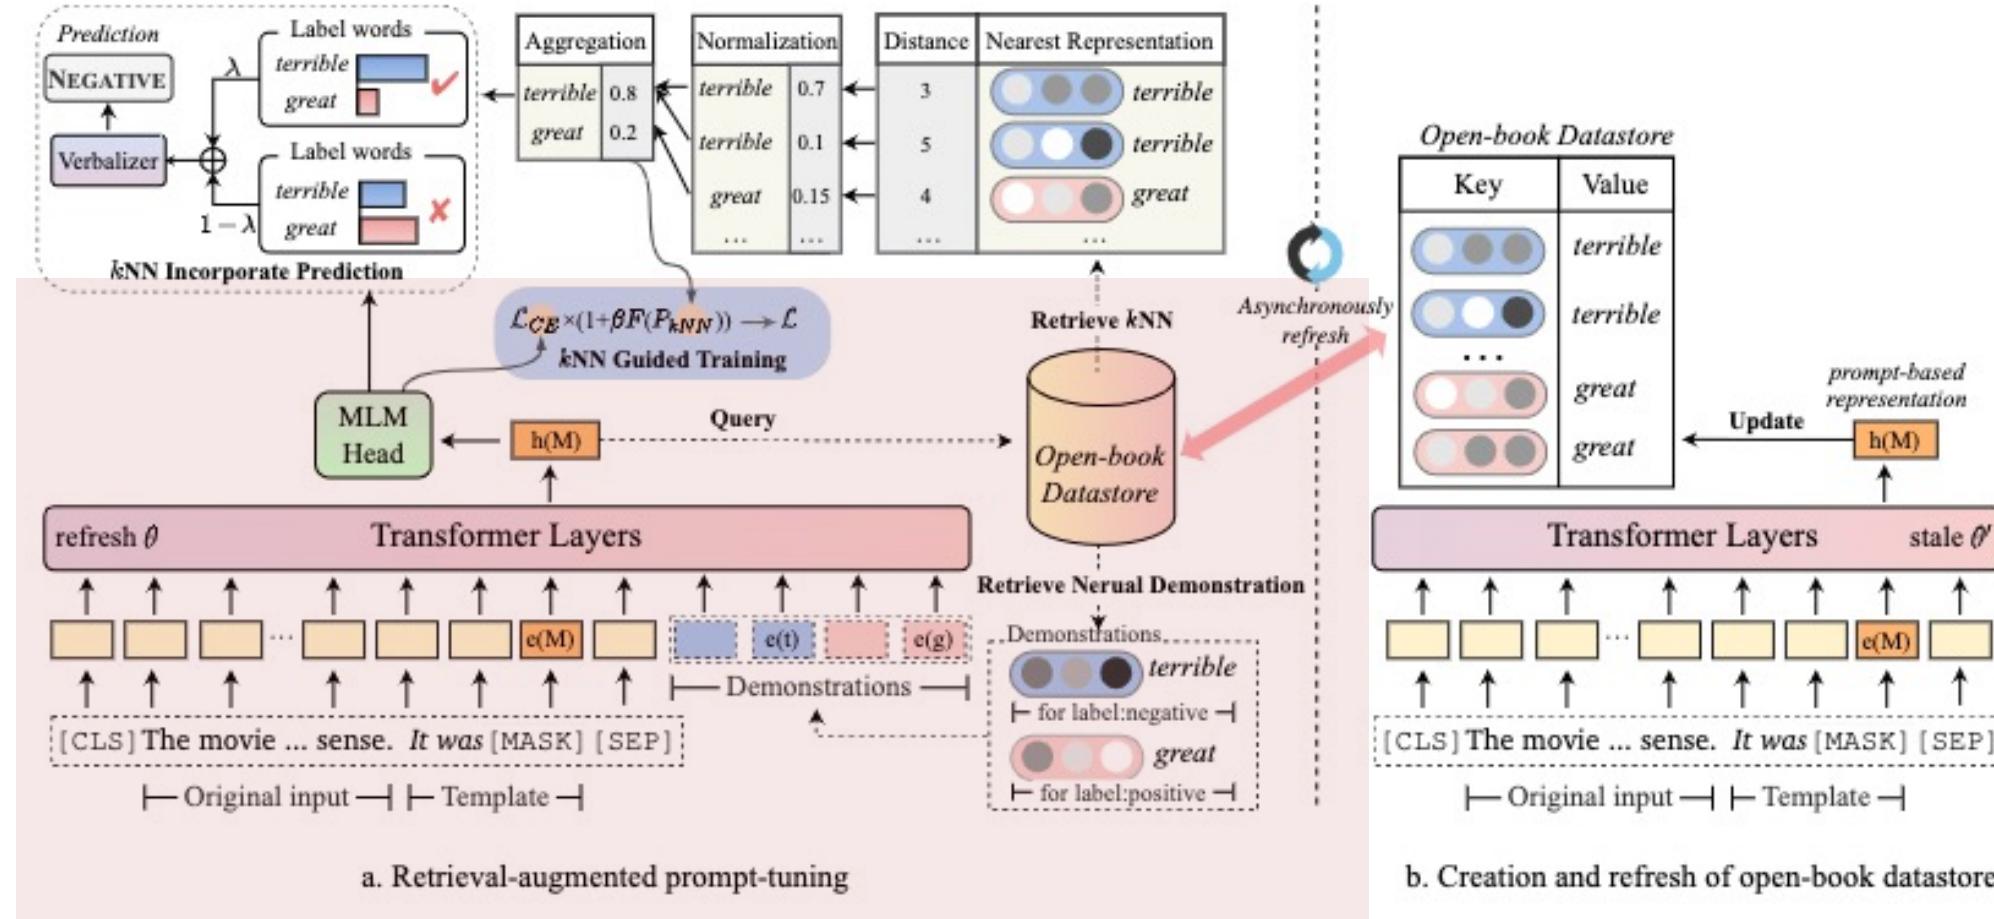


Figure 2: Overview of our RETRO PROMPT. Note that  $e(\cdot)$  donates the word embedding in the PLM  $\mathcal{M}$ , while “M”, “t” and “g” in the  $e(\cdot)$  specifically refers to the “[MASK]”, “terrible” and “great”.

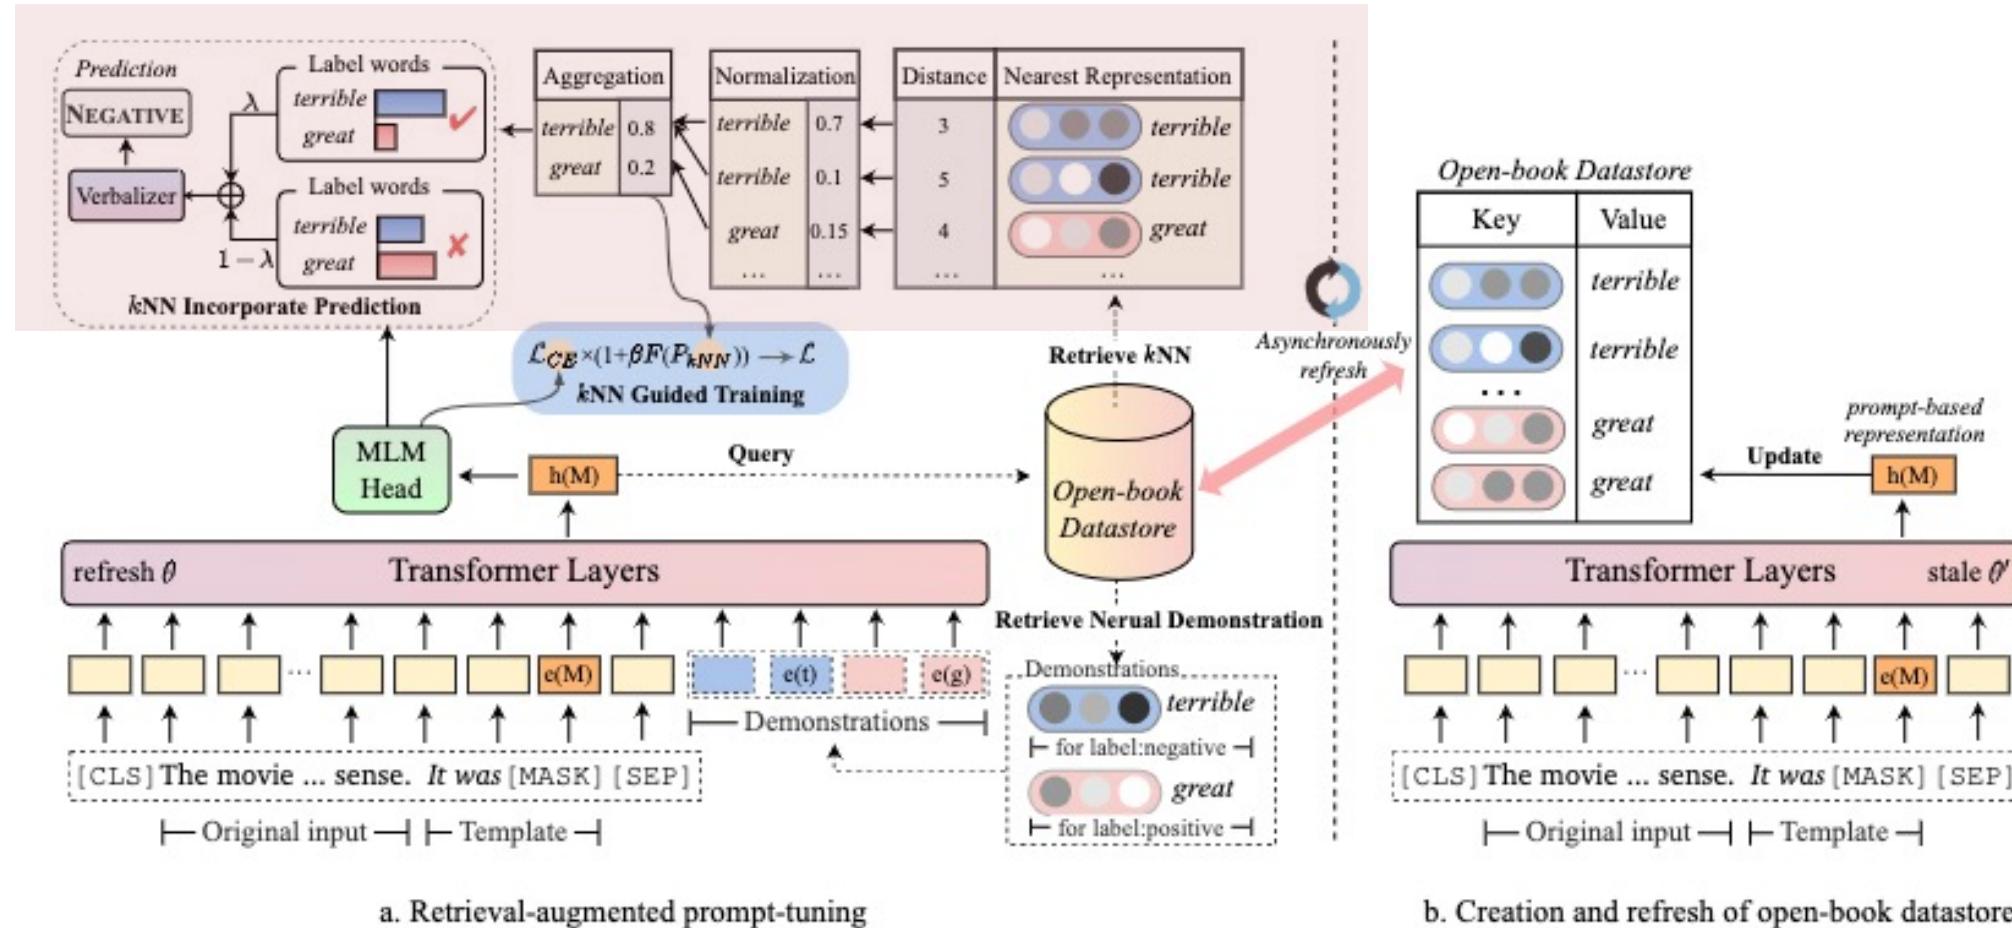


Figure 2: Overview of our RETRO\_PROMPT. Note that  $e(\cdot)$  donates the word embedding in the PLM  $\mathcal{M}$ , while “M”, “t” and “g” in the  $e(\cdot)$  specifically refers to the “[MASK]”, “terrible” and “great”.



Code is available in <https://github.com/zjunlp/PromptKG/tree/main/research/RetroPrompt>

St.	Model	Single Sentence			Sentence Pair			Model	Information Extraction			Avg.
		SST-2 (acc)	MR (acc)	CR (acc)	MNLI (acc)	QNLI (acc)	QQP (F1)		FewN (acc)	SemEval (acc)	TACRED (F1)	
16	FT	81.4 (3.8)	76.9 (5.9)	75.8 (3.2)	45.8 (6.4)	60.2 (6.5)	60.7 (4.3)	FT	52.7 (2.2)	66.1 (1.2)	25.8 (2.8)	60.6
	LM-BFF (man)	91.6 (1.2)	87.0 (2.0)	90.3 (1.6)	64.3 (2.5)	64.6 (5.4)	65.4 (5.3)	KnPr	65.3 (1.1)	80.9 (2.5)	33.2 (2.0)	71.4
	LM-BFF (D-demo)	91.8 (1.2)	86.6 (1.8)	90.2 (1.4)	64.8 (2.3)	69.2 (5.4)	68.2 (3.2)	KnPr (D-demo)	—	—	—	72.2*
	KPT †	90.3 (1.6)	86.8 (1.8)	88.8 (3.7)	61.4 (2.1)	61.5 (2.8)	71.6 (2.7)	KPT †	65.9 (1.5)	78.8 (2.1)	32.8 (1.7)	70.9
<b>Ours</b>		<b>93.9</b> (0.4)	<b>88.0</b> (0.8)	<b>91.9</b> (0.7)	<b>71.1</b> (1.8)	<b>71.6</b> (1.8)	<b>74.0</b> (2.0)	<b>Ours</b>	<b>67.3</b> (0.9)	<b>81.5</b> (1.3)	<b>40.7</b> (0.7)	<b>75.6</b>
4	FT	60.2 (2.8)	57.6 (1.4)	66.4 (5.5)	35.0 (0.3)	54.2 (3.9)	52.8 (4.7)	FT	32.7 (2.9)	38.8 (2.0)	14.7 (2.8)	45.8
	LM-BFF (man)	90.7 (0.8)	85.2 (2.8)	89.9 (1.8)	51.0 (2.5)	61.1 (6.1)	48.0 (4.9)	KnPr	52.5 (1.5)	58.4 (3.7)	28.8 (2.5)	62.8
	LM-BFF (D-demo)	90.2 (1.5)	85.5 (2.1)	89.7 (0.6)	56.1 (1.0)	61.7 (7.6)	63.2 (5.6)	KnPr (D-demo)	—	—	—	65.1*
	KPT †	88.2 (5.7)	83.4 (1.5)	87.2 (2.5)	53.7 (2.7)	59.2 (2.8)	54.9 (7.9)	KPT †	58.8 (2.2)	57.2 (3.2)	27.5 (2.2)	63.3
<b>Ours</b>		<b>91.5</b> (1.8)	<b>87.4</b> (0.5)	<b>91.4</b> (0.6)	<b>57.6</b> (5.5)	<b>62.2</b> (6.0)	<b>66.1</b> (4.1)	<b>Ours</b>	<b>60.9</b> (1.9)	<b>59.2</b> (3.0)	<b>32.1</b> (2.0)	<b>67.6</b>
0	LOTClass♦	71.8	81.7	50.1	50.4	36.5	55.9	LOTClass♦	11.5	9.8	2.5	41.1
	FT	49.1	50.0	49.8	34.4	49.5	31.6	FT	10.0	6.2	0.5	31.2
	LM-BFF (man)	83.5	80.3	78.4	49.7	50.5	49.7	KnPr	15.9	10.3	2.3	46.7
	LM-BFF (D-demo)	82.9	80.7	<b>81.4</b>	52.2	53.5	44.0	KnPr (D-demo)	—	—	—	47.0*
	KPT †	78.4	81.9	71.4	37.1	55.3	47.5	KPT †	24.6	11.6	0.8	45.7
	<b>Ours</b>	<b>86.8</b>	<b>83.5</b>	79.7	<b>53.7</b>	<b>56.2</b>	<b>56.7</b>	<b>Ours</b>	<b>41.3</b>	<b>12.2</b>	<b>2.8</b>	<b>52.5</b>

“

- **Prompt learning** with pre-trained language model can obtain promising performance and yield more gains with **retrieval augmentation**
- **Data augmentation (e.g., via prompting)** is powerful for few-shot learning
- Task-specific **inductive bias (Schema)** can help data-efficient knowledge graph construction



This is very cool



# Scope: Efficient Knowledge Graph Construction



**Data Efficiency:** few-shot learning with different types of data

- Prompt learning, Retrieval-augmentation



**Model Efficiency:** parameter-efficient learning, unified architecture

- Parameter-efficient learning



**Inference Efficiency:** fast Inference

- FastRE, sequence-to-sequence

# Why Model-Efficient Knowledge Graph Construction?

Large-scale pre-trained models

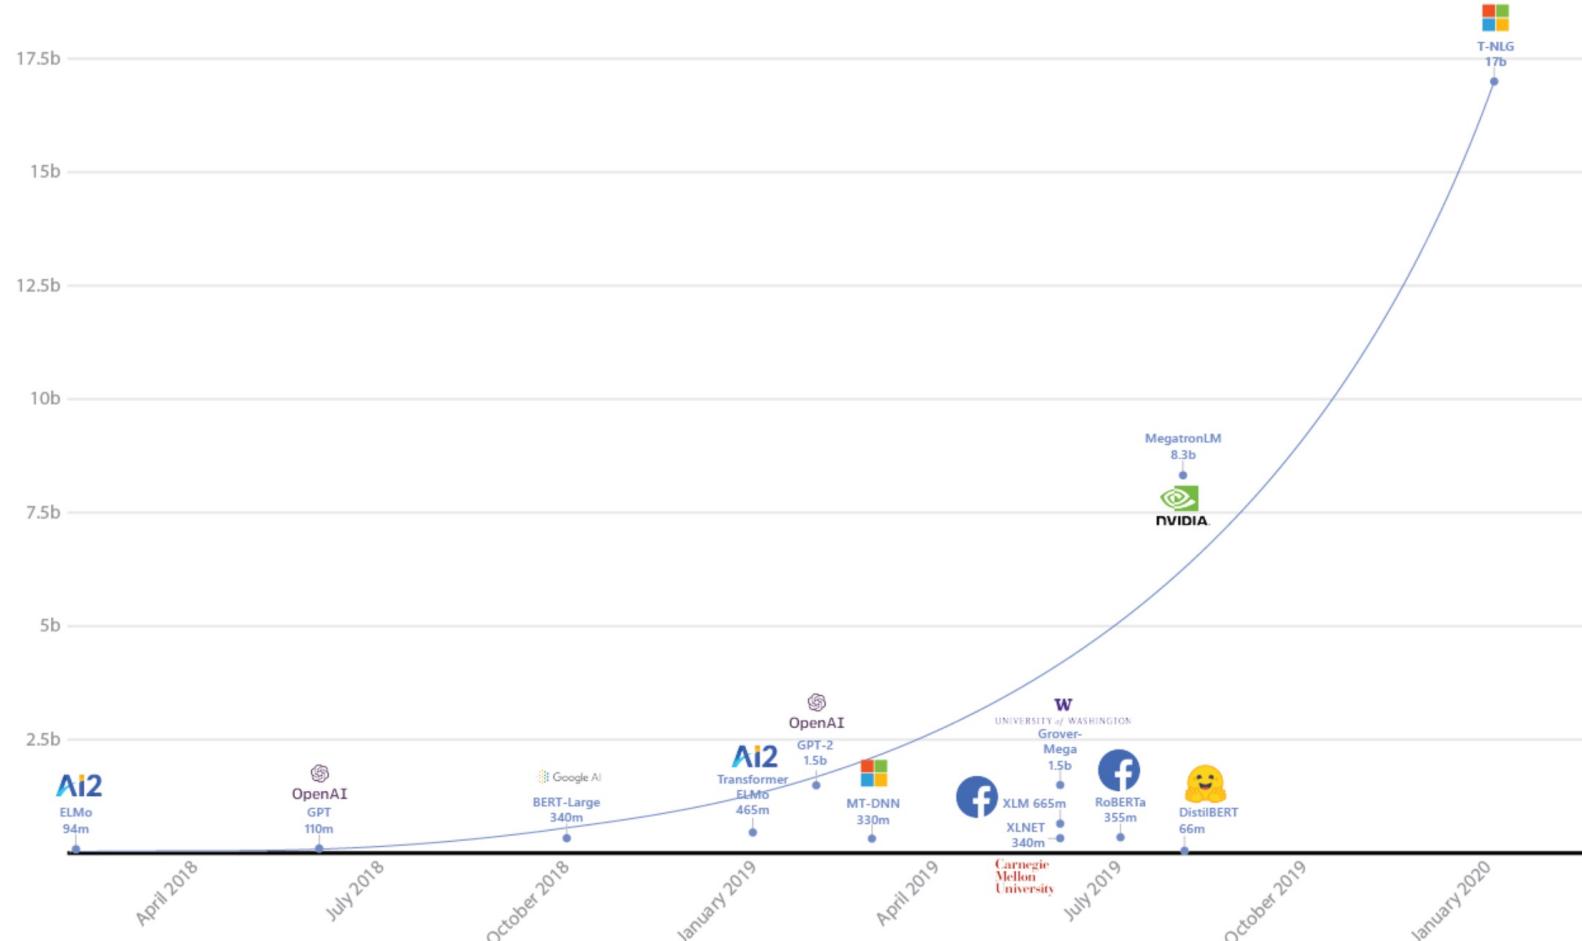


Figure from Turing-NLG

# Why Model-Efficient Knowledge Graph Construction?



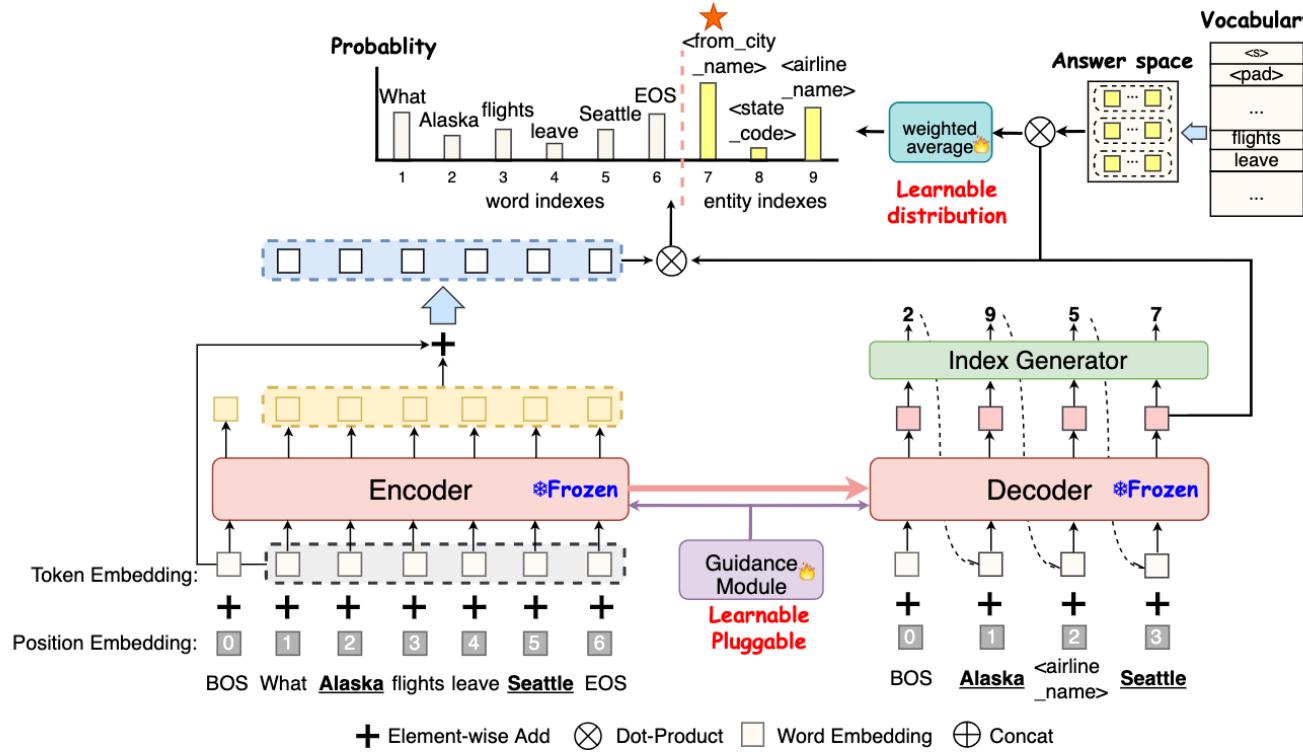
AACL IJCNLP 2022

## Efficient transfer learning

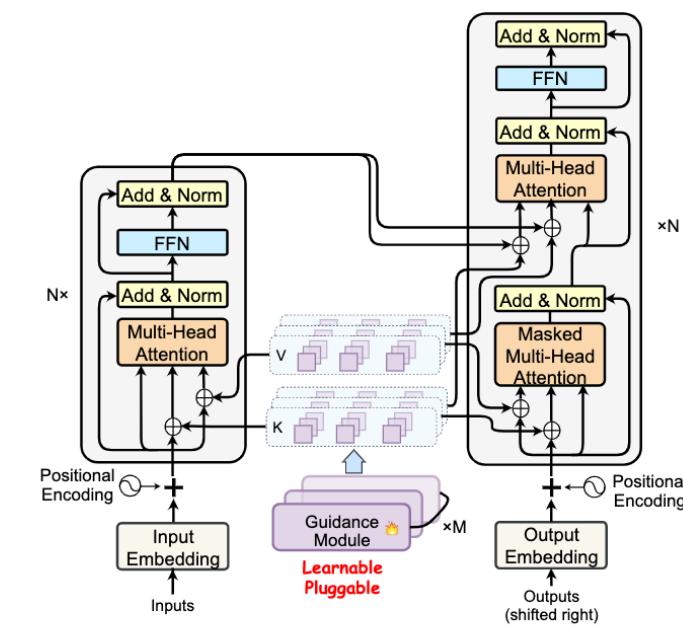




## Continuous prompts for NER



(a) Overview of LightNER.



(b) Attention with guidance.

Figure 2: Overview of our LightNER framework.

Code is available in <https://github.com/zjunlp/DeepKE/tree/main/example/ner/few-shot>

Source	Methods	MIT Movie						MIT Restaurant						ATIS		
		10	20	50	100	200	500	10	20	50	100	200	500	10	20	50
None	LC-BERT	25.2	42.2	49.6	50.7	59.3	74.4	21.8	39.4	52.7	53.5	57.4	61.3	44.1	76.7	90.7
	LC-BART	10.2	27.5	44.2	47.5	54.2	64.1	6.3	8.5	51.3	52.2	56.3	60.2	42.0	72.7	87.5
	Template	37.3	48.5	52.2	56.3	62.0	74.9	46.0	57.1	58.7	60.1	62.8	65.0	71.7	79.4	92.6
	BERT-MRC†	18.7	48.3	55.5	62.5	80.2	82.1	12.3	37.1	53.5	63.9	65.5	70.4	35.3	63.2	90.2
	LightNER	<b>41.7</b>	<b>57.8</b>	<b>73.1</b>	<b>78.0</b>	<b>80.6</b>	<b>84.8</b>	<b>48.5</b>	<b>58.0</b>	<b>62.0</b>	<b>70.8</b>	<b>75.5</b>	<b>80.2</b>	<b>76.3</b>	<b>85.3</b>	<b>92.8</b>
CoNLL03	Neigh.Tag.	0.9	1.4	1.7	2.4	3.0	4.8	4.1	3.6	4.0	4.6	5.5	8.1	2.4	3.4	5.1
	Example.	29.2	29.6	30.4	30.2	30.0	29.6	25.2	26.1	26.8	26.2	25.7	25.1	22.9	16.5	22.2
	MP-NSP	36.4	36.8	38.0	38.2	35.4	38.3	46.1	48.2	49.6	49.6	50.0	50.1	71.2	74.8	76.0
	LC-BERT	28.3	45.2	50.0	52.4	60.7	76.8	27.2	40.9	56.3	57.4	58.6	75.3	53.9	78.5	92.2
	LC-BART	13.6	30.4	47.8	49.1	55.8	66.9	8.8	11.1	42.7	45.3	47.8	58.2	51.3	74.4	89.9
	Template	42.4	54.2	59.6	65.3	69.6	80.3	53.1	60.3	64.1	67.3	72.2	75.7	77.3	88.9	93.5
	BERT-MRC†	20.2	50.8	56.3	62.9	81.5	82.3	15.8	39.5	54.8	65.8	68.8	73.5	40.5	66.7	91.8
	LightNER	<b>62.9</b>	<b>75.6</b>	<b>78.8</b>	<b>82.2</b>	<b>84.5</b>	<b>85.7</b>	<b>58.1</b>	<b>67.4</b>	<b>69.5</b>	<b>73.7</b>	<b>78.4</b>	<b>81.1</b>	<b>86.9</b>	<b>89.4</b>	<b>93.9</b>

Table 1: Model performance (F1 score) in the cross-domain low-resource setting. “†” indicates that we rerun their public code in this setting. All of our experiments and baselines adopt large version of LMs.

# Low-high Layer vs. High-low Layer

Guidance module applied to **higher layers of LMs can better stimulate knowledge** from PLMs for downstream tasks more efficiently

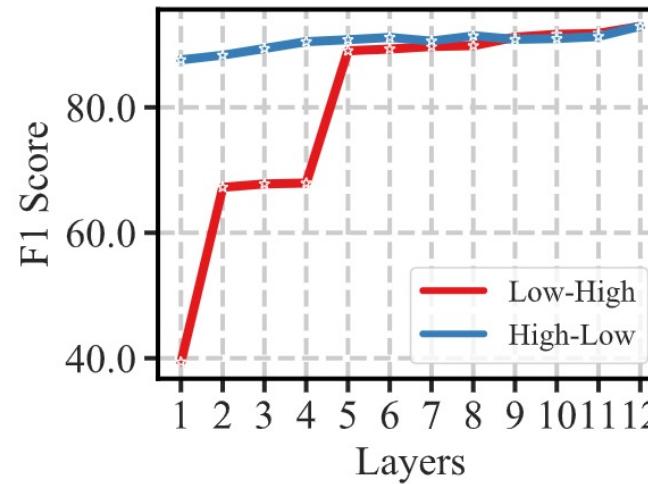


Figure 3: Performances on CoNLL03 as the layers of guidance module varies.

# Parameter-efficient Knowledge Graph Completion

Reformulating KG completion as a “fill-in-the-blank” with parameter-lite encoder

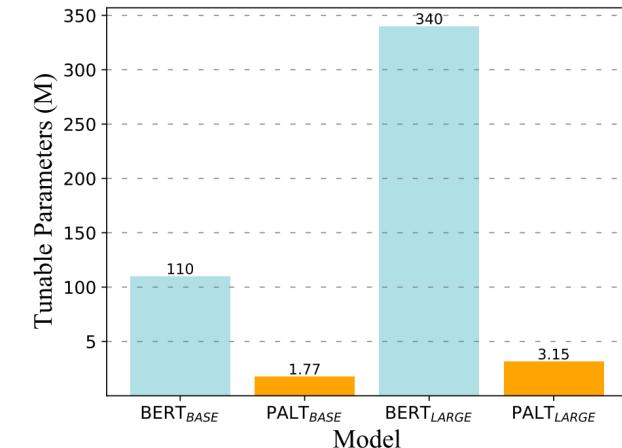
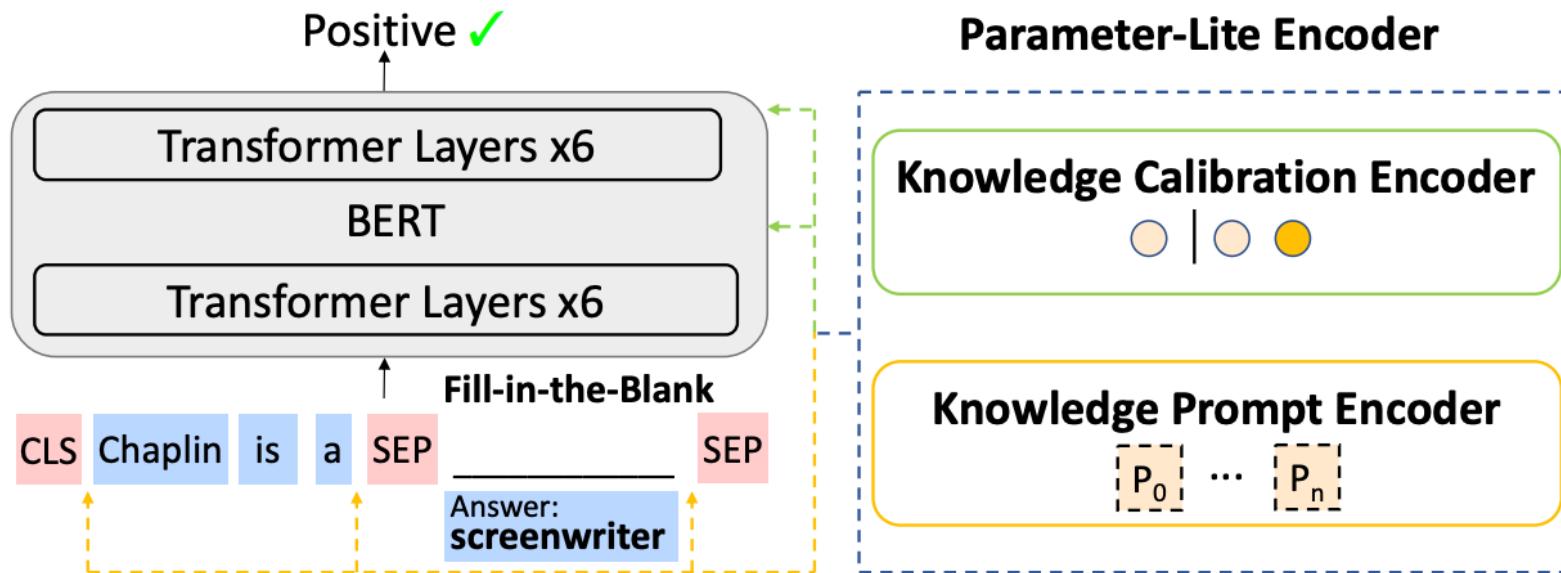


Figure 2: Compare the number of tunable parameters of PALT and BERT.

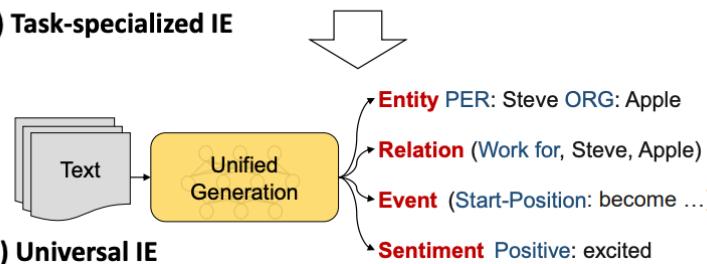
Any other solution?

# Model-efficient Knowledge Graph Construction

Unified one model for **ALL**

Task	Schema	Instance
Entity	PER: _ ORG:_	<b>PER</b> <b>ORG</b> In 1997, Steve was excited to become the CEO of Apple.
Relation	(_, Work for, _)	<b>Work For</b> In 1997, Steve was excited to become the CEO of Apple.
Event	Type Start Position employee employer ...	<b>Start-Position</b> <b>Person</b> <b>Entity</b> In 1997, Steve was excited to become the CEO of Apple.
Sentiment	Positive { Opinion: _; Target: _ }	<b>Opinion</b> <b>Positive</b> <b>Target</b> In 1997, Steve was excited to become the CEO of Apple.

(a) Task-specialized IE



(b) Universal IE

Figure 1: From (a) Task-specialized IE: different tasks, different structures, different schemas to (b) Universal IE: unified modeling via structure generation.

```

( Spot Name: Info Span
  (Asso Name: Info Span)
  (Asso Name: Info Span)
)
)
)
  
```

(a) Structured extraction language (SEL) for Universal IE.

```

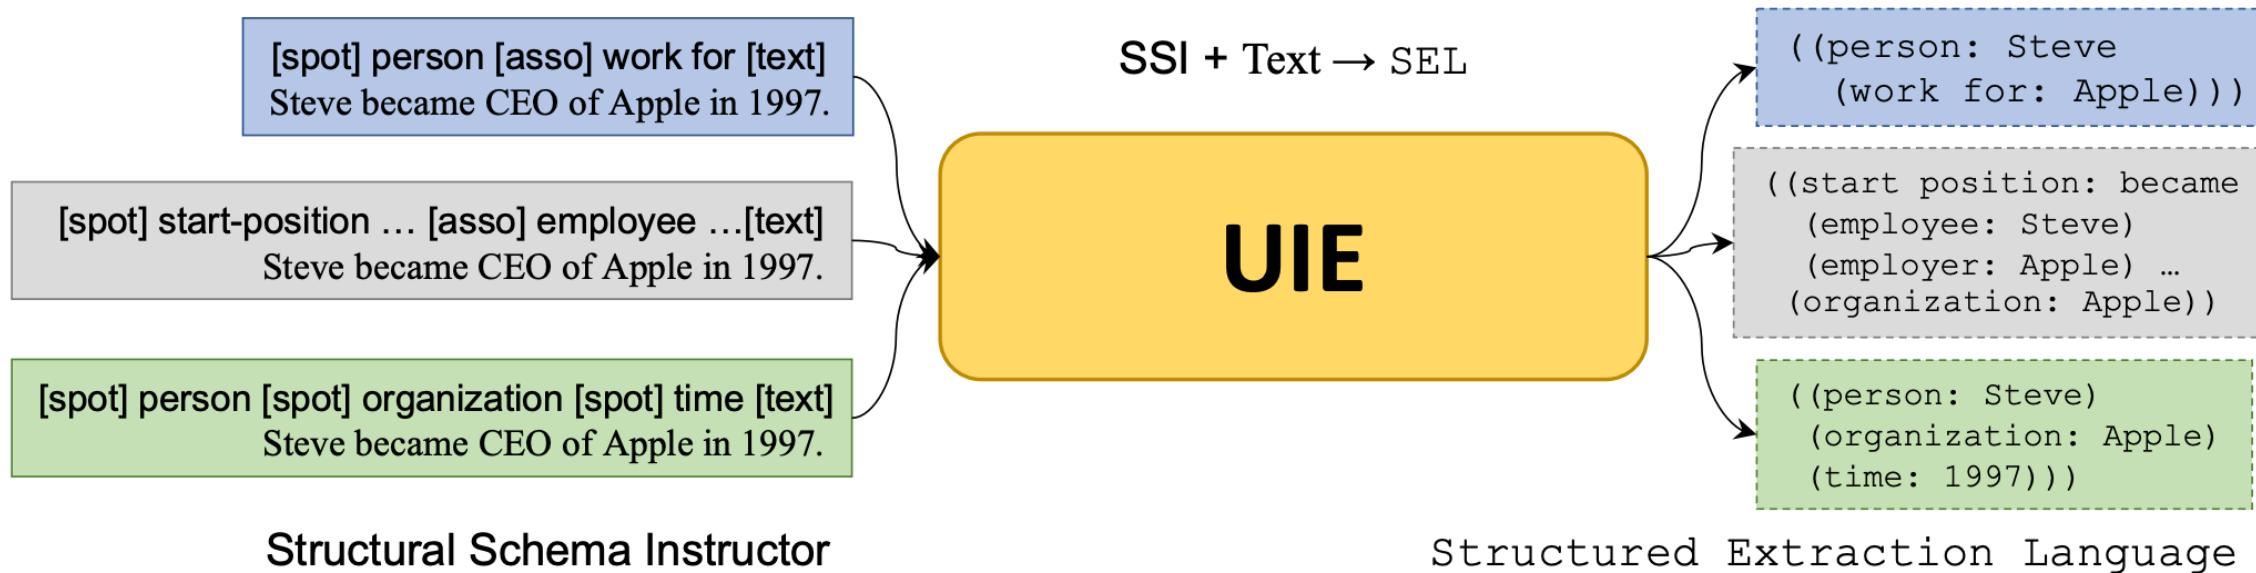
(
  (person: Steve
    (work for: Apple)
  )
  (start-position: became
    (employee: Steve)
    (employer: Apple)
    (time: 1997)
  )
  (organization: Apple)
  (time: 1997)
)
  
```

(b) The SEL representation of the extraction structure of “Steve became CEO of Apple in 1997.”, where the relation structure is marked blue, the event structure is marked red, and the rest are entities.

Figure 2: Illustrations of structured extraction language.

# Model-efficient Knowledge Graph Construction

Unified one model for **ALL**



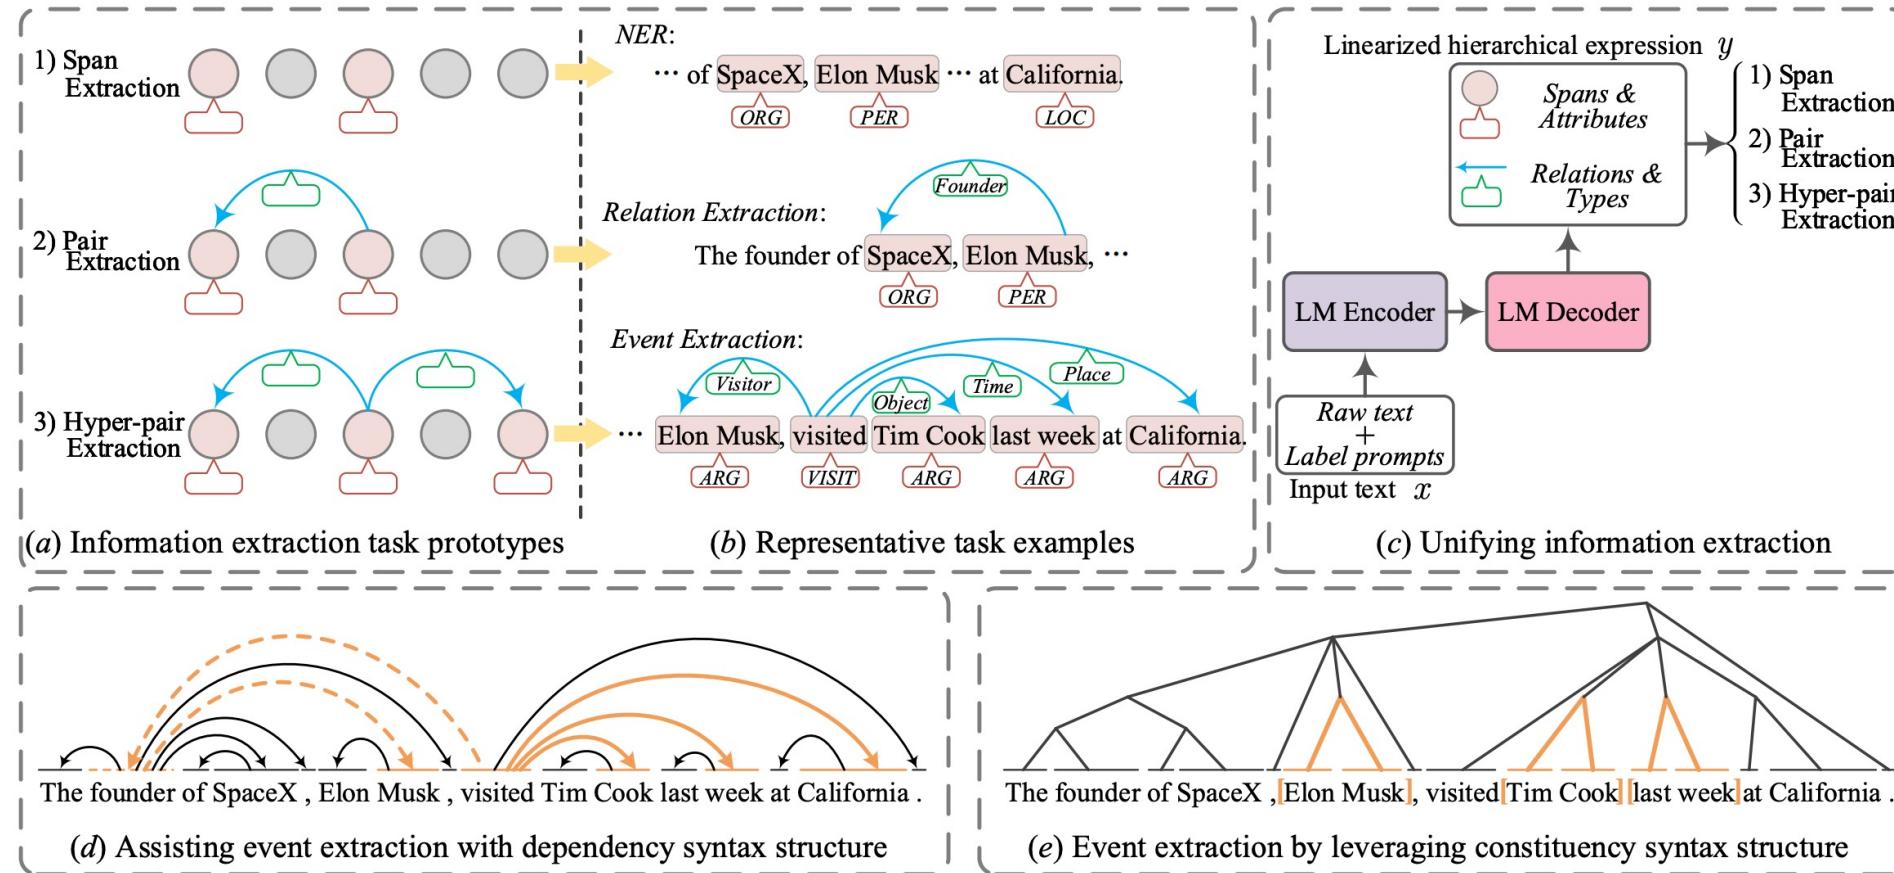
# Model-efficient Knowledge Graph Construction

Unified one model for **ALL**

Model		1-Shot	5-Shot	10-Shot	AVE-S	1%	5%	10%	AVE-R
<b>Entity</b> <small>(CoNLL03)</small>	T5-v1.1-base	12.73	30.17	58.89	33.93	75.74	85.71	87.70	83.05
	Fine-tuned T5-base	24.93	54.85	65.31	48.36	78.51	87.67	88.91	85.03
	UIE-base w/o SSI	43.52	64.76	72.47	60.25	81.91	<b>88.41</b>	<b>89.84</b>	86.72
	UIE-base	<b>46.43</b>	<b>67.09</b>	<b>73.90</b>	<b>62.47</b>	<b>82.84</b>	88.34	89.63	<b>86.94</b>
<b>Relation</b> <small>(CoNLL04)</small>	T5-v1.1-base	2.35	7.99	25.98	12.11	6.08	32.38	41.87	26.78
	Fine-tuned T5-base	4.24	28.16	41.44	24.61	12.89	37.75	49.95	33.53
	UIE-base w/o SSI	13.21	40.35	49.47	34.34	24.21	48.70	56.59	43.17
	UIE-base	<b>22.05</b>	<b>45.41</b>	<b>52.39</b>	<b>39.95</b>	<b>30.77</b>	<b>51.72</b>	<b>59.18</b>	<b>47.22</b>
<b>Event Trigger</b> <small>(ACE05-Evt)</small>	T5-v1.1-base	19.40	43.35	50.57	37.77	25.59	49.47	57.18	44.08
	Fine-tuned T5-base	30.18	48.31	51.27	43.25	31.08	51.16	57.76	46.67
	UIE-base w/o SSI	32.07	48.11	51.00	43.73	32.71	53.20	59.26	48.39
	UIE-base	<b>38.14</b>	<b>51.21</b>	<b>53.23</b>	<b>47.53</b>	<b>41.53</b>	<b>55.70</b>	<b>60.29</b>	<b>52.51</b>
<b>Event Argument</b> <small>(ACE05-Evt)</small>	T5-v1.1-base	2.75	20.21	27.53	16.83	3.59	21.53	30.90	18.67
	Fine-tuned T5-base	<b>6.96</b>	25.07	30.96	21.00	7.39	24.97	33.90	22.09
	UIE-base w/o SSI	9.31	23.99	30.31	21.20	9.57	27.25	34.18	23.67
	UIE-base	<b>11.88</b>	<b>27.44</b>	<b>33.64</b>	<b>24.32</b>	<b>12.80</b>	<b>30.43</b>	<b>36.28</b>	<b>26.50</b>
<b>Sentiment</b> <small>(16res)</small>	T5-v1.1-base	0.04	2.11	12.66	4.94	3.50	27.08	45.97	25.52
	Fine-tuned T5-base	6.55	21.06	29.92	19.18	18.72	39.63	51.65	36.67
	UIE-base w/o SSI	7.79	17.77	32.07	19.21	19.14	42.76	53.44	38.45
	UIE-base	<b>10.50</b>	<b>26.24</b>	<b>39.11</b>	<b>25.28</b>	<b>24.24</b>	<b>49.31</b>	<b>57.61</b>	<b>43.72</b>

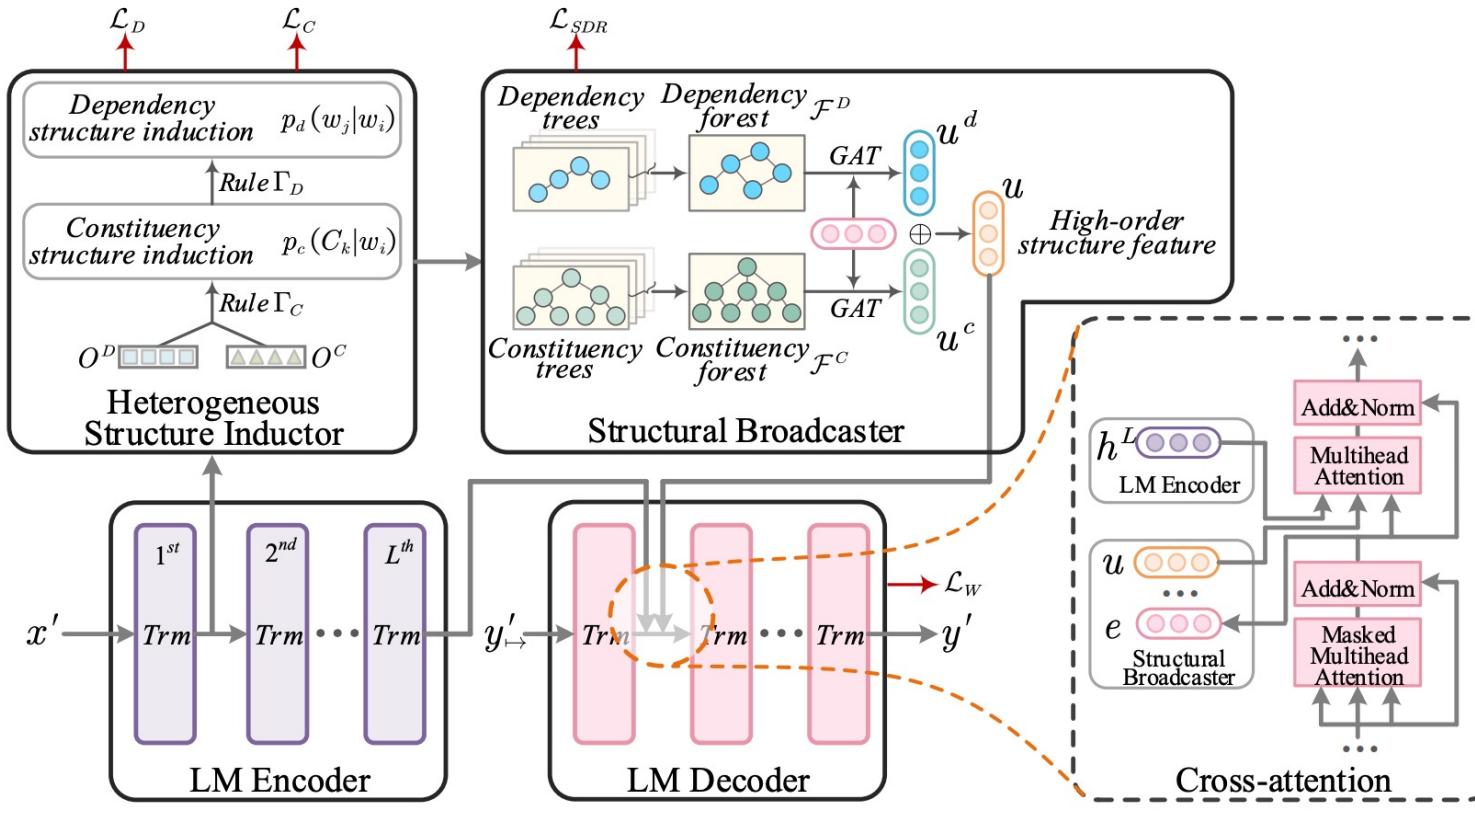
# Model-efficient Knowledge Graph Construction

Unified one model for ALL

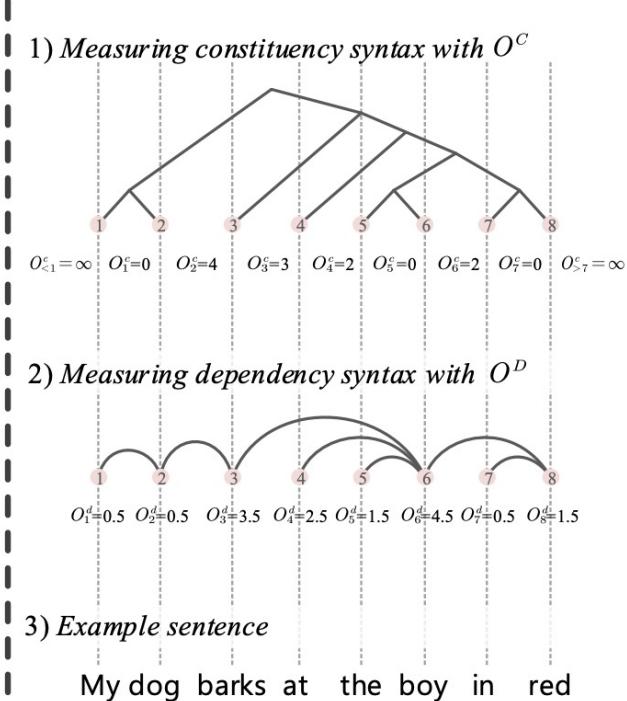


# Model-efficient Knowledge Graph Construction

## Unified one model for **ALL**



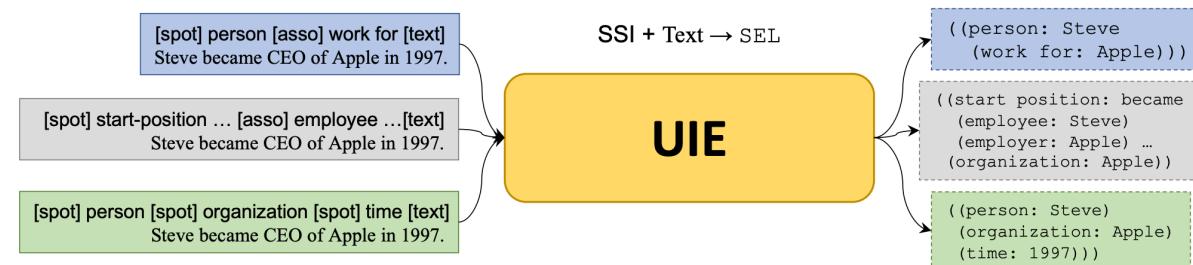
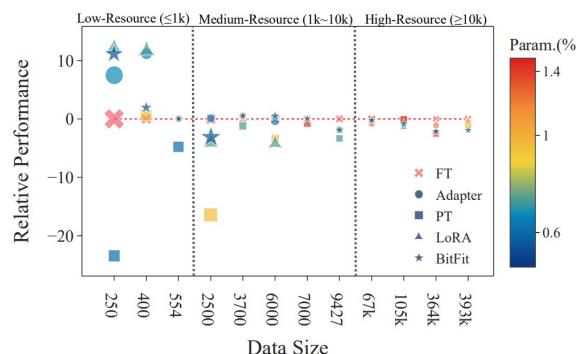
(a) Unsupervised structure-aware post-training



### (b) Heterogeneous syntax measurements

“

- By tuning just a fraction amount of parameters comparing to full model finetuning, **parameter-efficient tuning methods** can achieve performance on par with vanilla finetuning
- Instead of task-specialized modeling, it is beneficial to advocate for a **unifying view** of knowledge graph construction



# Scope: Efficient Knowledge Graph Construction



**Data Efficiency:** few-shot learning with different types of data

- Prompt learning, Retrieval-augmentation



**Model Efficiency:** parameter-efficient learning, unified architecture

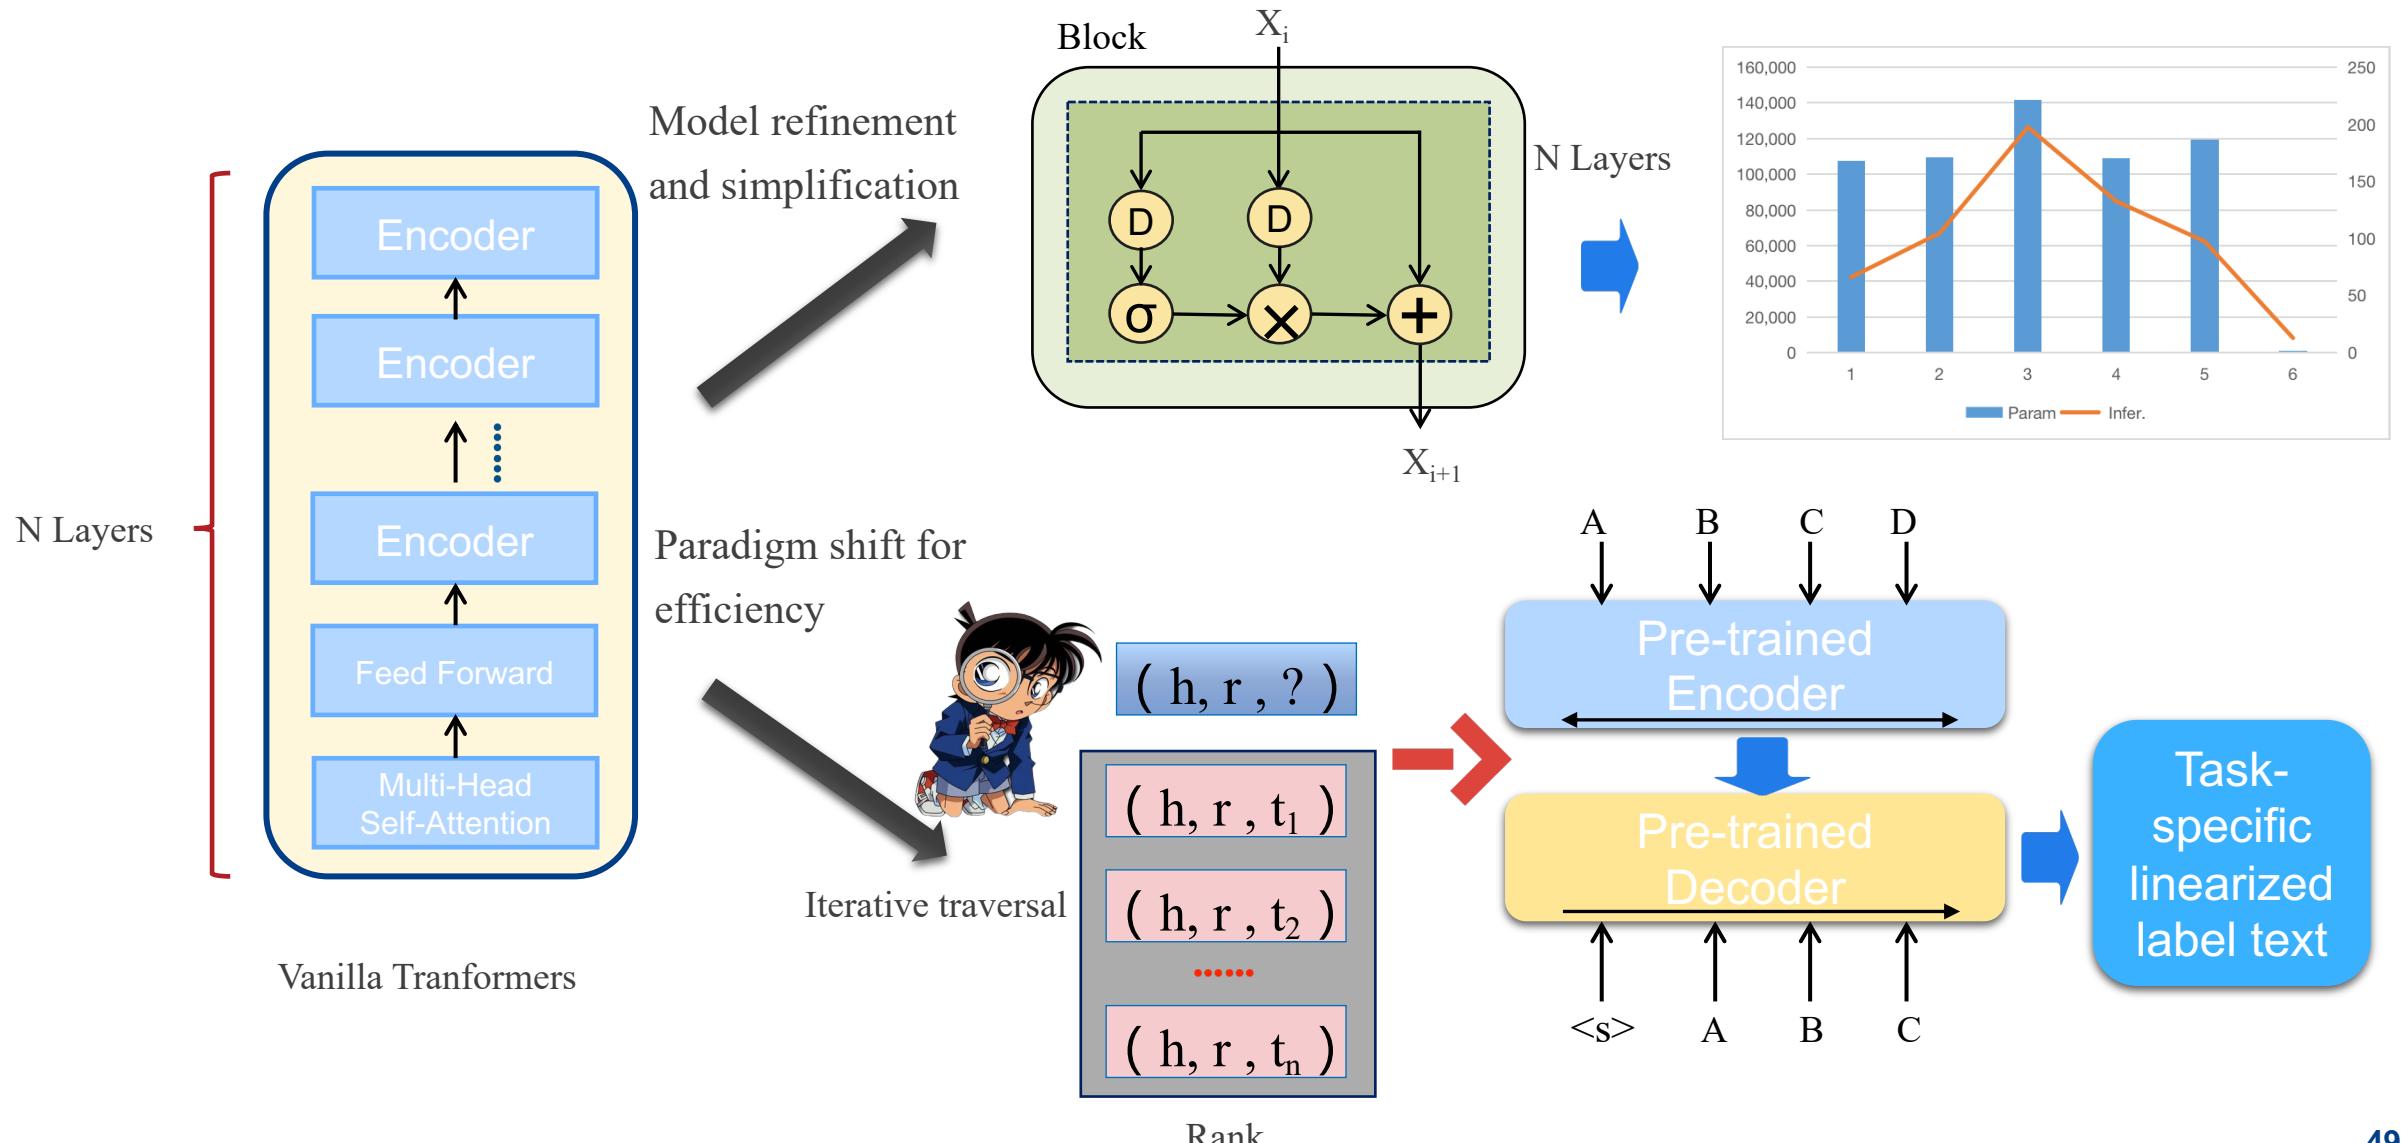
- Parameter-efficient learning



**Inference Efficiency:** fast Inference

- FastRE, sequence-to-sequence

# Very Slow Inference



# Inference-efficient Knowledge Graph Construction



AAACL IJCNLP 2022

Convolutional encoder, NOT BERT!

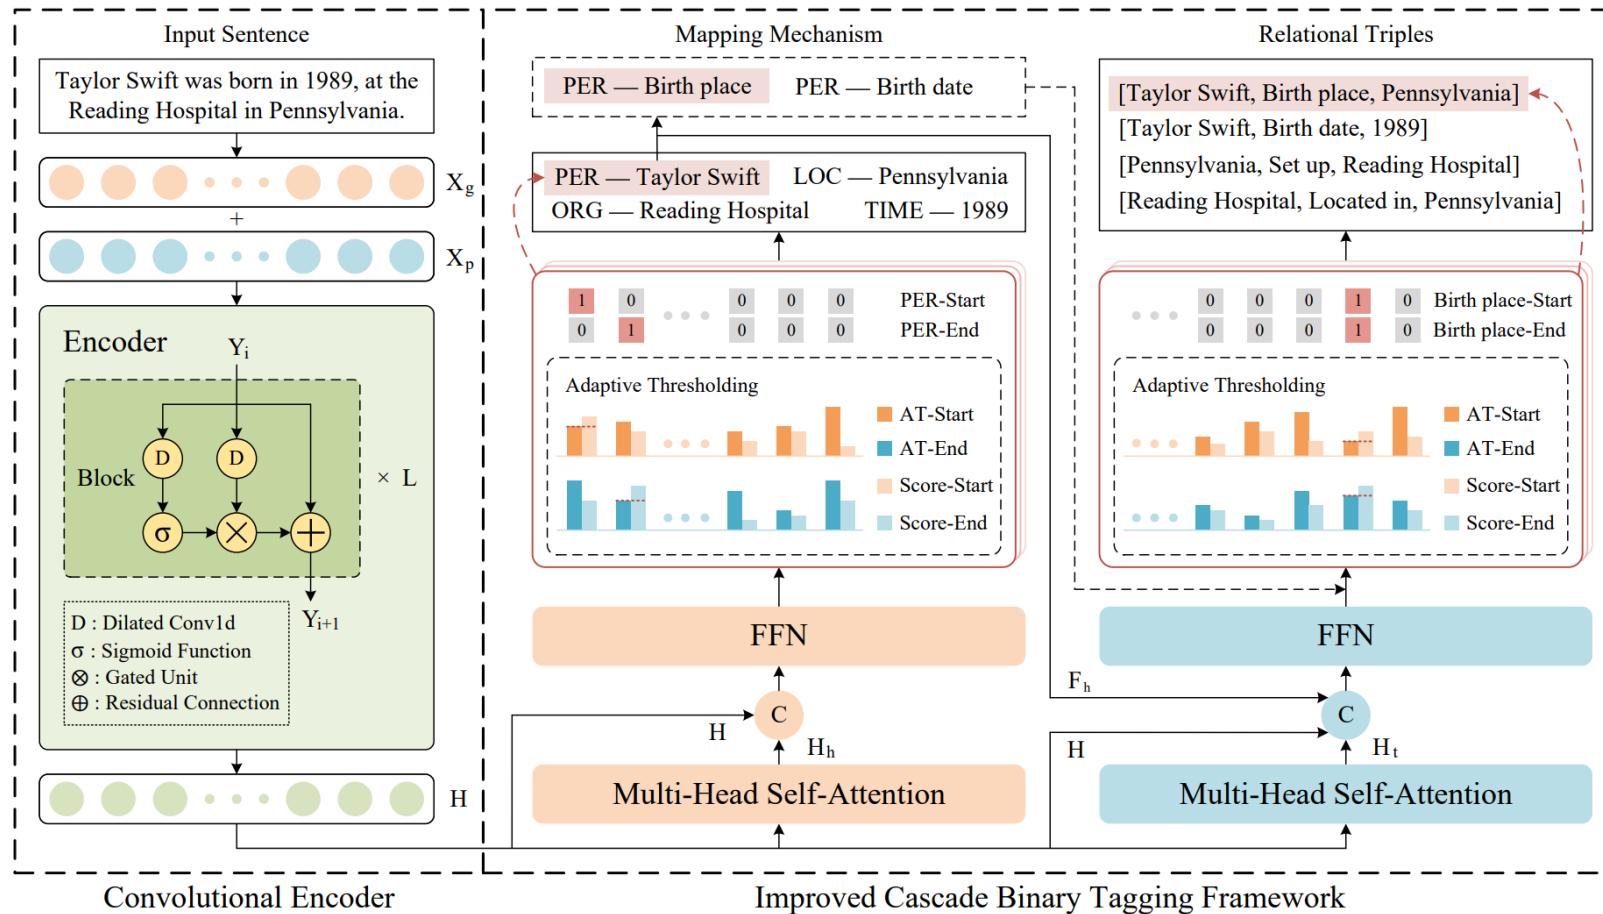


Figure 1: The overall structure of FastRE.

FastRE: Towards Fast Relation Extraction with Convolutional Encoder  
and Improved Cascade Binary Tagging Framework (IJCAI2022)

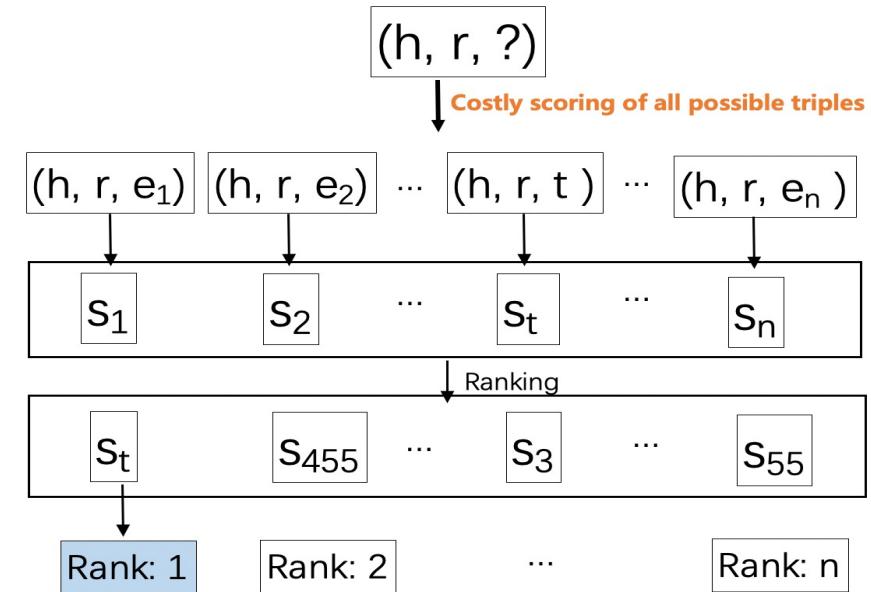
Any other solution?

# Efficient Knowledge Graph Completion



AAACL IJCNLP 2022

1. Traditional KGE models embed the entities and relations into a vector space and obtain the predicted triples by leveraging **pre-defined scoring functions** to those vectors.
2. A discrimination strategy has to **costly scoring of all possible triples** in inference and suffer from the instability of negative sampling.

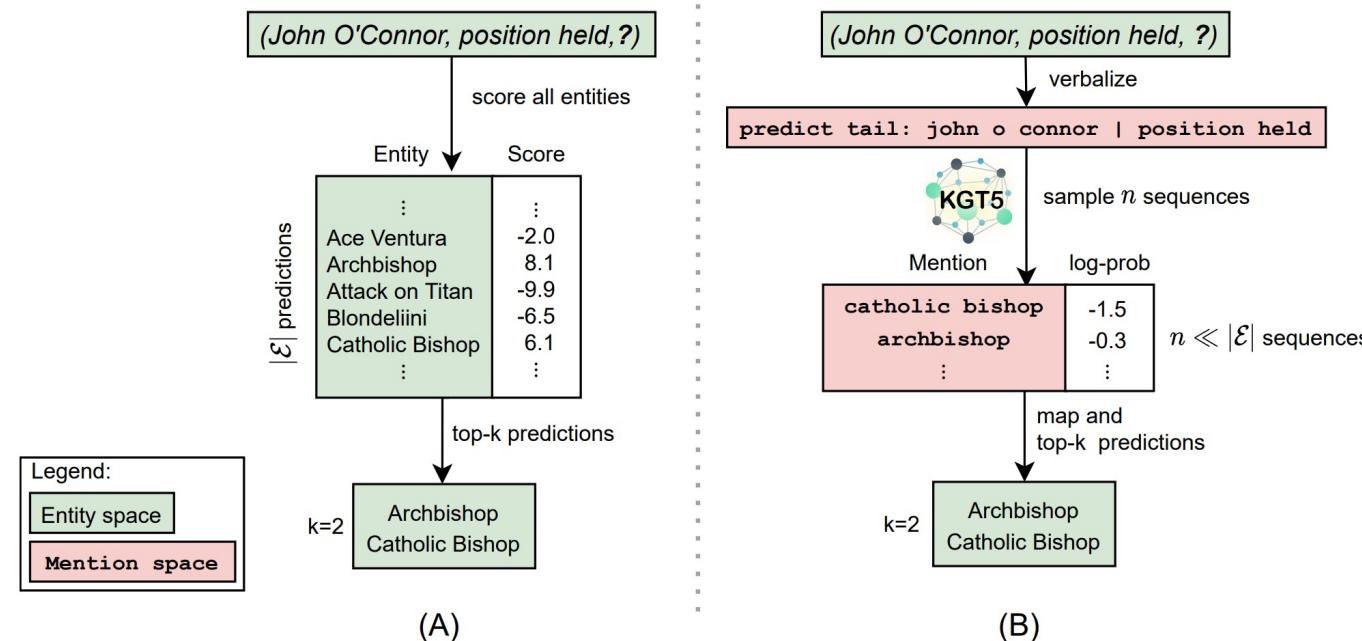
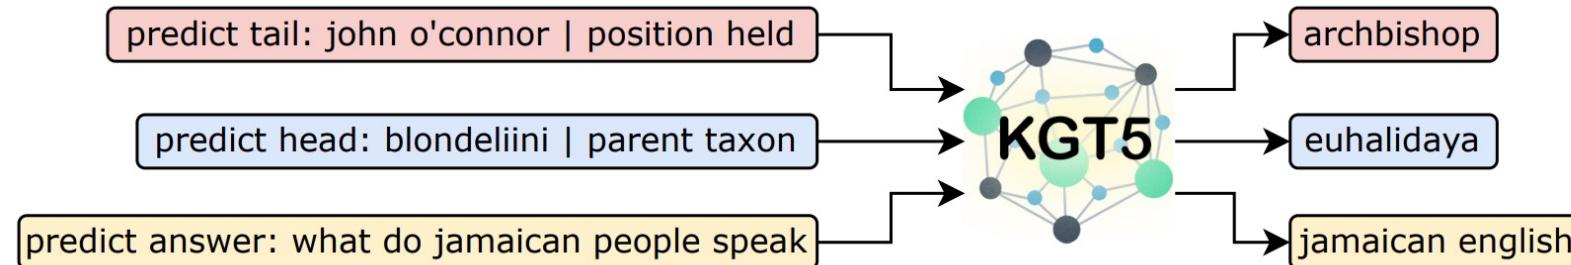


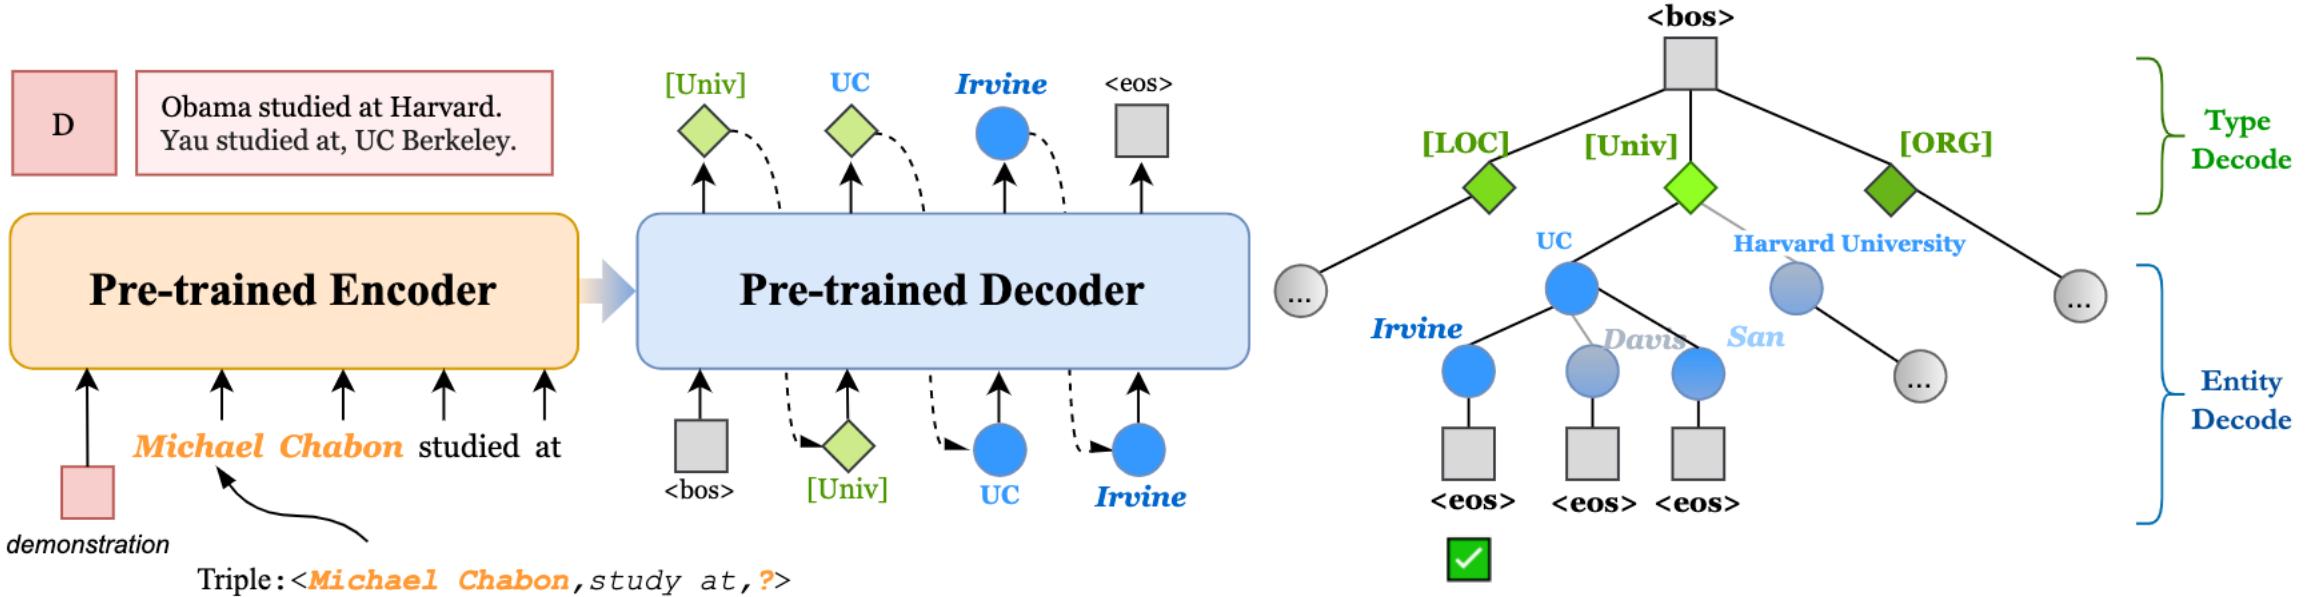
For One Triple	Method	Complexity	Time under RTX 3090
Training	TransE	$O(k + 1)$	0.08ms
	KG-BERT	$O( d ^2 \times (k + 1))$	72ms
	<b>GenKGC</b>	$O( d ^2)$	2.35ms
Inference	TransE	$O( \mathcal{E} )$	0.02s
	KG-BERT	$O( d ^2 \times  \mathcal{E} )$	10100s
	<b>GenKGC</b>	$O( d ^2 \times  d ^k)$	0.96s

How to complete Knowledge Graph with PLMs **efficiently?**

One solution: *Fast Inference with Generation*

**Encoder-decoder Transformer** model can serve as a scalable and versatile model





$$p_{\theta}(y | x) = \prod_{i=1}^{|c|} p_{\theta}(z_i | z_{<i}, x) \prod_{i=|c|+1}^N p_{\theta}(y_i | y_{<i}, x), \quad (1)$$

$$\mathcal{L} = -\log p_{\theta}(y | x)$$

## Benchmark and real-world datasets

### Datasets:

- WN18RR
- FB15k-237
- OpenBG500 (A New KG)

### Code for reproducibility:

<https://github.com/zjunlp/PromptKG/tree/main/research/GenKG>

**Table 3: Summary statistics of benchmark datasets.**

Dataset	# Ent	# Rel	# Train	# Dev	# Test
WN18RR	40,943	11	86,835	3,034	3,134
FB15k-237	14,541	237	272,115	17,535	20,466
OpenBG500	269,658	500	1,242,550	5,000	5,000

OpenBG: Large Scale Open Business Knowledge Graph Evaluation Benchmark [Leaderboard](#)

AliOpenKG 2022-03-08 417 98

[Apply for dataset](#) [New a notebook](#)

Content Notebook Comment Leaderboard Submission

#### Description

OpenBG Benchmark is a large-scale open business knowledge graph evaluation benchmark, which includes multiple sub-datasets and sub-tasks. The dataset is built on the basis of [OpenBG](#). OpenBG is an open business knowledge graph that utilizes a unified Schema covering multi-modal datasets in large scale, which contains millions of products and consumer demand provided by Alibaba Knowledge Engine Group and Zhejiang University. The goal of OpenBG is to use open business knowledge to discover the value of the social economy, promote cross-disciplinary research in the fields of digital commerce and digital economy, and serve the national strategic needs of the healthy development of digital economy.

Website: <https://tianchi.aliyun.com/OpenBG>

Github: <https://github.com/OpenGBBenchmark>

<https://github.com/OpenGBBenchmark/OpenBG>

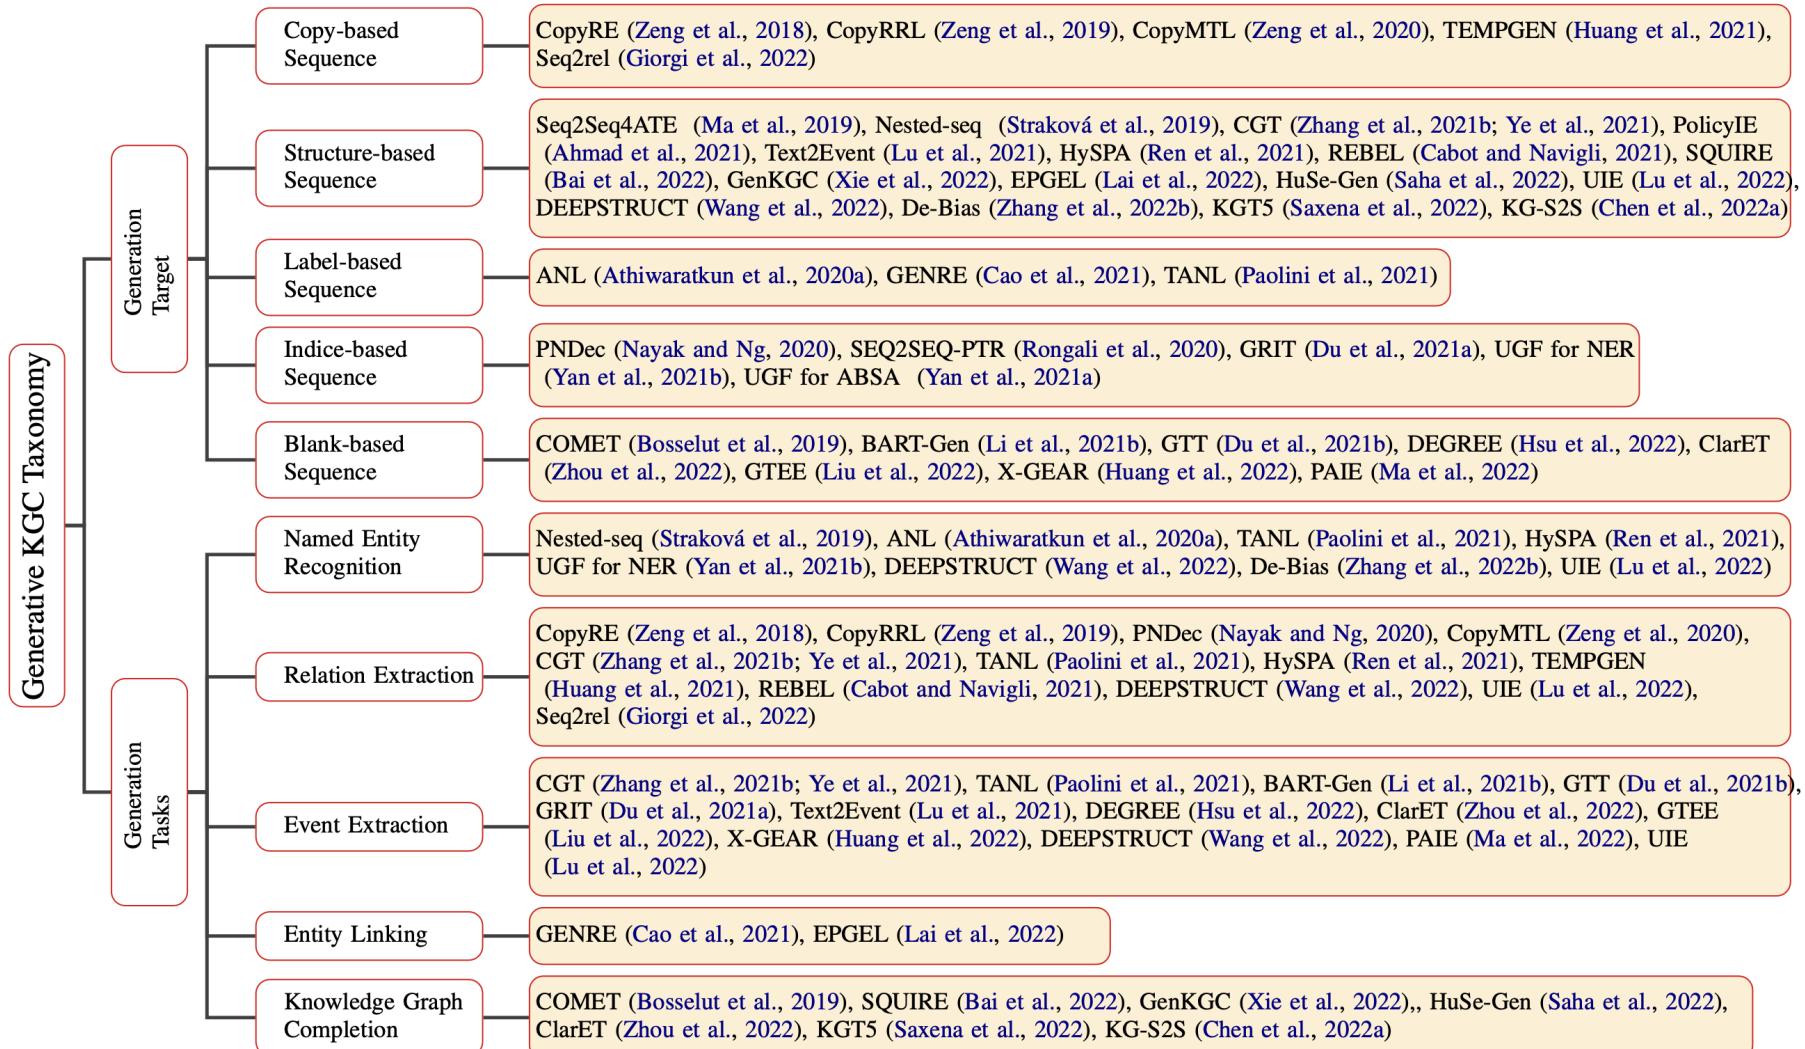
Method	WN18RR			FB15k-237			OpenBG500			
	Hits@1	Hits@3	Hits@10	Hits@1	Hits@3	Hits@10	Hits@1	Hits@3	Hits@10	
<i>Graph embedding approach</i>										
TransE [2] ◊	260M Para.			0.532	0.198	0.376	0.441	0.207	0.340	0.513
DistMult [13] ◊	0.412	0.470	0.504	0.199	0.301	0.446	0.049	0.088	0.216	
ComplEx [11] ◊	0.409	0.469	0.530	0.194	0.297	0.450	0.053	0.120	0.266	
RotatE [9]	0.428	0.492	0.571	0.241	0.375	0.533	-	-	-	
TuckER [1]	0.443	0.482	0.526	0.226	0.394	0.544	-	-	-	
ATTH [4]	0.443	0.499	0.486	0.252	0.384	0.549	-	-	-	
<i>Textual encoder</i>										
KG-BERT [14]	0.041	0.302	0.524	100	100s	120	0.023	0.049	0.241	
StAR [12]	0.041	0.301	0.511	0.709	0.205	0.482	-	-	-	
GenKGC	110M Para.			0.535	0.199	0.96s	0.19	0.203	0.280	0.351

# Generative Knowledge Graph Construction



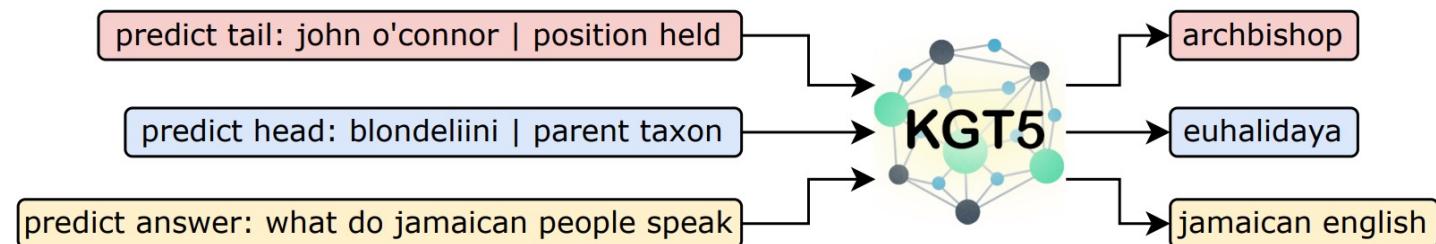
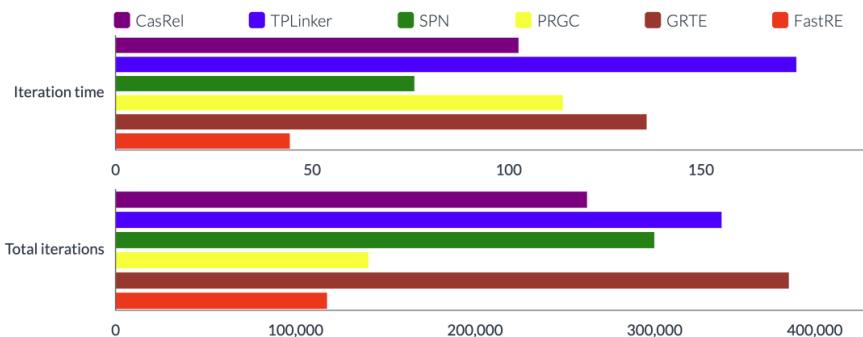
AACL IJCNLP 2022

Public Repository : [https://github.com/zjunlp/Generative\\_KG\\_Construction\\_Papers](https://github.com/zjunlp/Generative_KG_Construction_Papers)



“

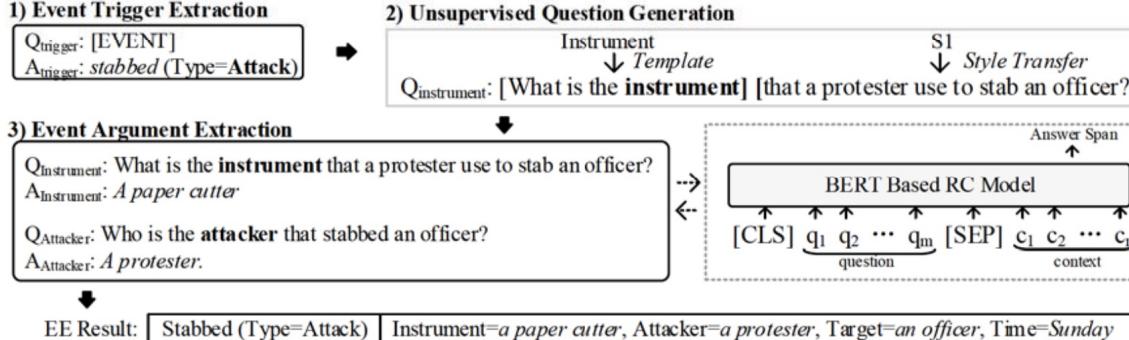
- ❑ **Vanilla encoders (such as CNN)** with several innovations can improve efficiency while also keeping promising performance
- ❑ **Paradigm shifting (such as from discrimination to generation)** can help reduce the inference time cost



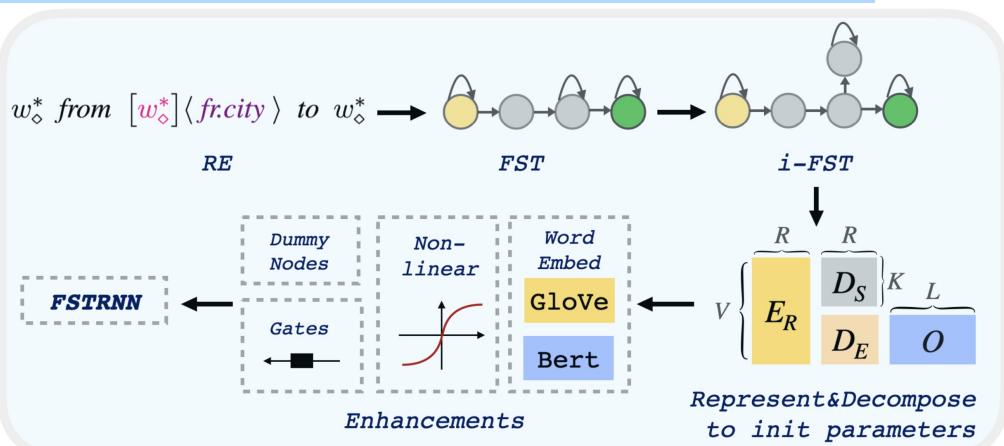
# Other Approaches

## Formulating KE as QA/MRC

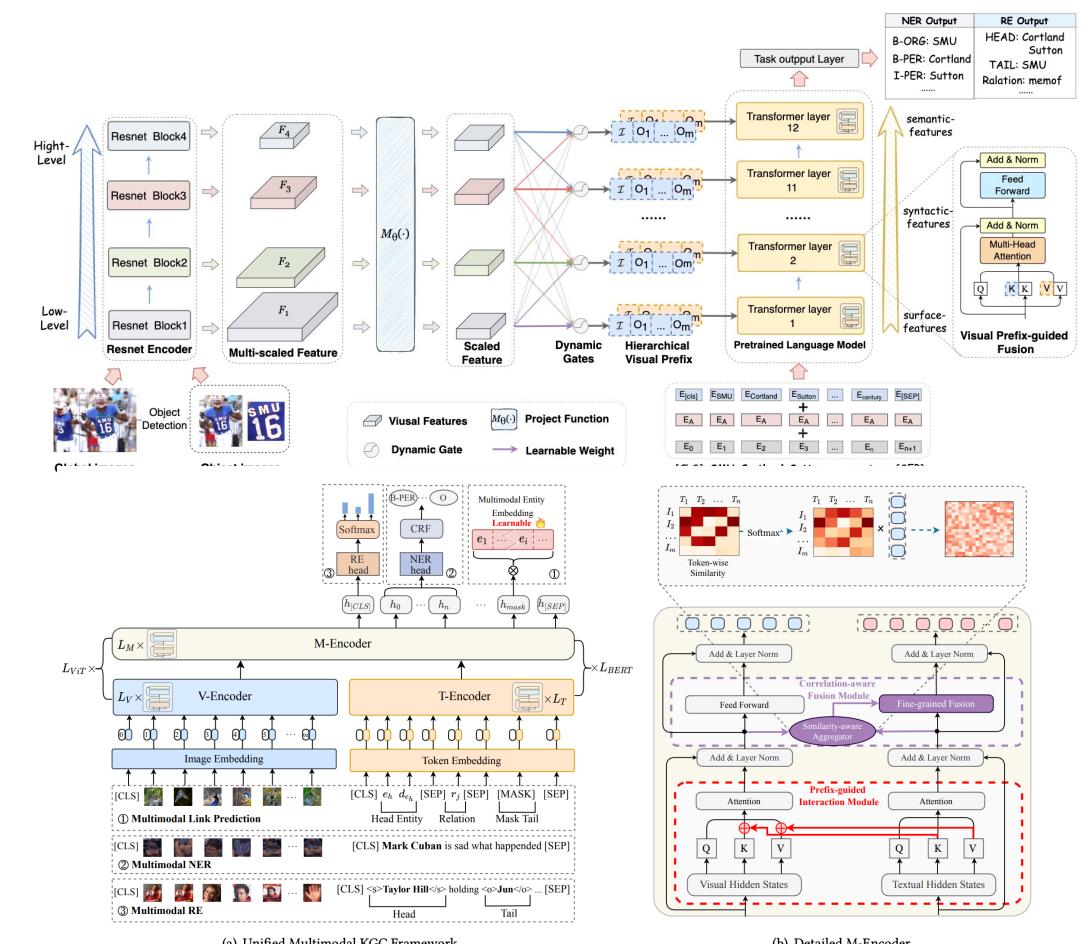
S1: On Sunday, a protester stabbed an officer with a paper cutter.



## Converting regular expressions into NN



## Multimodal enhancement



Language Model Priming for Cross-Lingual Event Extraction (AAAI2022)

Event Extraction as Machine Reading Comprehension (EMNLP2020)

Neuralizing Regular Expressions for Slot Filling (EMNLP2021)

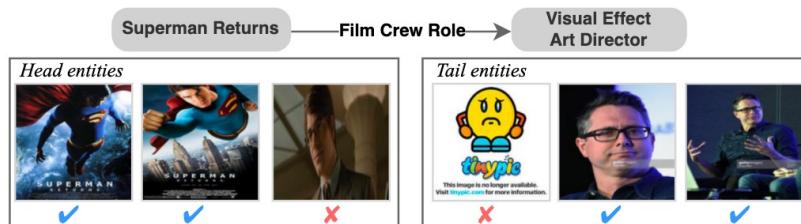
Good Visual Guidance Make A Better Extractor: Hierarchical Visual Prefix for Multimodal Entity and Relation Extraction (NAACL 2022 Findings)<sup>60</sup>

## Remaining Issues and Future Works

# Remaining Issue 1

## Efficient **multimodal** Knowledge Graph Construction

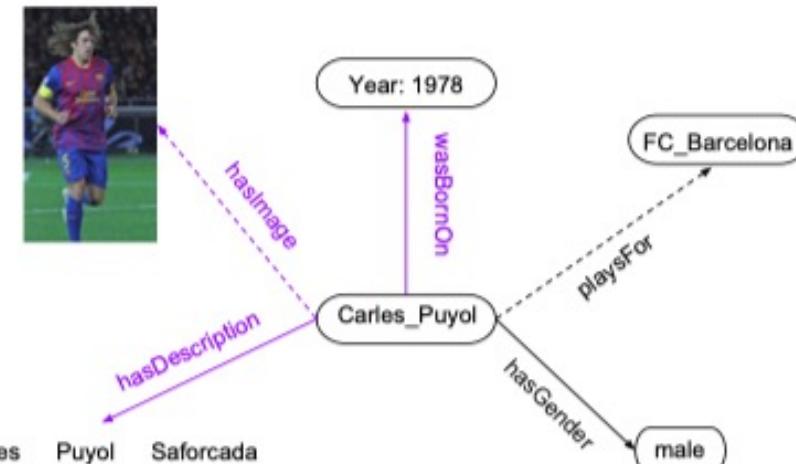
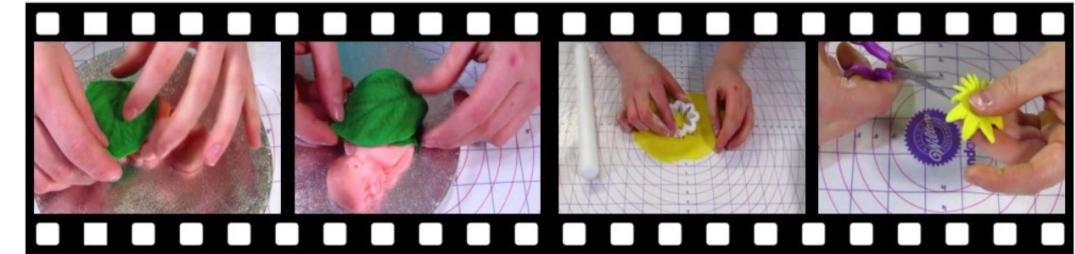
- When and why visual information is helpful?
- Rich semantic in dynamic videos



(a) Examples of Multimodal Link Prediction



(b) Examples of Multimodal Relation Extraction

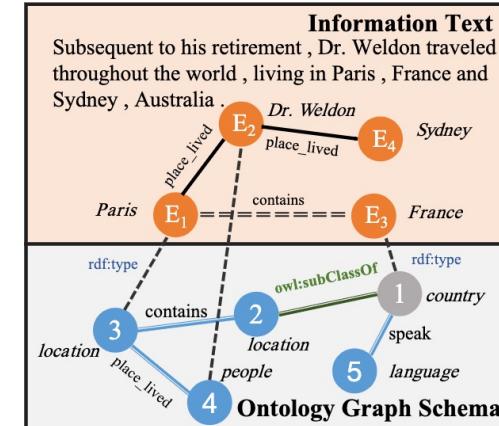


Is Visual Context Really Helpful for Knowledge Graph? A Representation Learning (MM21)  
Multi-Modal Knowledge Graph Construction and Application: A Survey 2022

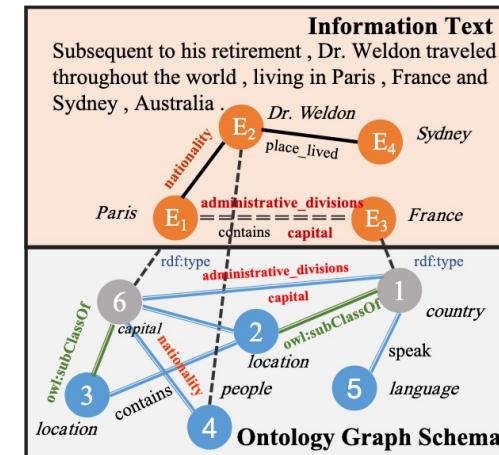
# Remaining Issue2

# Efficient **lifelong** Knowledge Graph construction

- Rich **connections** between schema
  - Schema **self-adapting**, **self-boosting** and **self-correction**

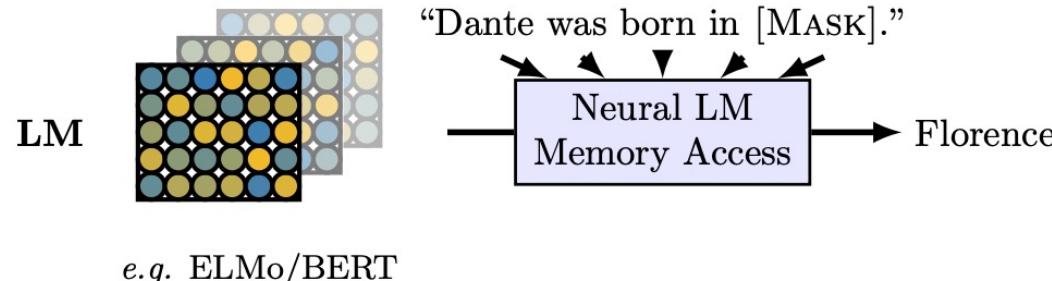


- ◆ Schema Change:
  - (1) Horizontal Node Expansion
  - (2) Vertical Node Expansion
  - (3) Analogous Node Replacement

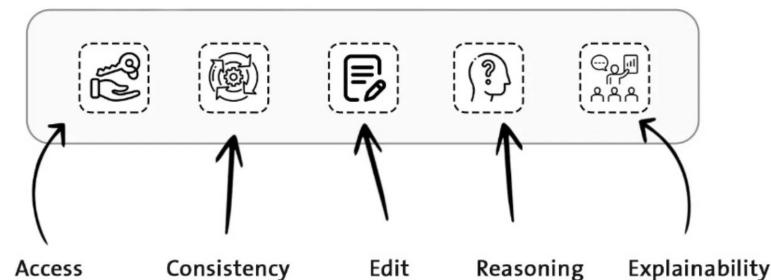


Curriculum-meta Learning for Order-robust Continual Relation Extraction (AAAI2021)  
Refining Sample Embeddings with Relation Prototypes to Enhance Continual Relation Extraction (ACL2021)

## Manipulating knowledge in PLM for efficient KGC



## LMs-as-KBs



Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: **Let's think step by step.**

(Output) There are 16 balls in total. Half of the balls are golf balls. That means that there are 8 golf balls. Half of the golf balls are blue. That means that there are 4 blue golf balls. ✓

- Use PLM to obtain fact knowledge
- Can the PLM have the reasoning abilities?

Symbolic **VS** Neural -> Symbolic + Neural

# Efficient KGC for Social Good

Efficient NER for information extraction and help **information visualization**



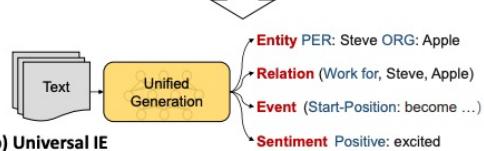
# Outlook : Changes of Interactive Objects Opportunities



AACL IJCNLP 2022

Task	Schema	Instance
Entity	PER: _ORG_	PER: ORG In 1997, Steve was excited to become the CEO of Apple.
Relation	(_, Work for, _)	Work For In 1997, Steve was excited to become the CEO of Apple.
Event	Type   Start Position employee   employer ...	Start-Position Person   Entity In 1997, Steve was excited to become the CEO of Apple.
Sentiment	Positive { Opinion: _; Target: _ }	Opinion   Positive   Target In 1997, Steve was excited to become the CEO of Apple.

(a) Task-specialized IE



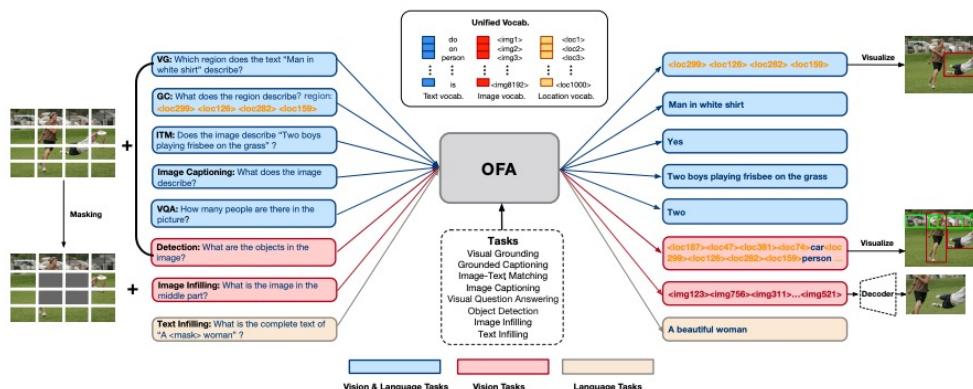
(b) Universal IE

UIE' ACL22



With massive text data, language pre-training

With multi-modal, multi-type data, multi-task/modal pre-training



With the world, environment, embodied pre-training

Q how do neural networks work

How neural networks work - A simple introduction  
Information flows through a neural network in two ways. When it's learning (being trained) or operating normally (after being trained), patterns of information are fed into the network via the input units, which trigger the layers of hidden units, and these in turn arrive at the output units. This common design is called a feedforward network.

How do Neural Networks really work? - Analytics Vidhya

people exposed to artificial intelligence generally have a good high-level idea of how a neural network works — data is passed from one layer of the neural network to the next, and this data is propagated from the topmost layer to the bottom layer until, somehow, the algorithm outputs the prediction on whether an image is that of a chihuahua or a ...

What are Neural Networks? | IBM

Neural networks rely on training data to learn and improve their accuracy over time. However, once these learning algorithms are fine-tuned for accuracy, they are powerful tools in computer science and artificial intelligence, allowing us to classify and cluster data at a high velocity. Tasks in speech recognition or image recognition can take minutes versus hours when compared to the manual ...

How Do Neural Networks Really Work? | Nick McCullum

In its most basic form, a neural network only has two layers - the input layer and the output layer. The output layer is the component of the neural net that actually makes predictions. For example, if you wanted to make predictions using a simple weighted sum (also called linear regression)

WebGPT

## Open-sourced Toolkits

## OpenPrompt

An Open-Source Framework for Prompt-learning.

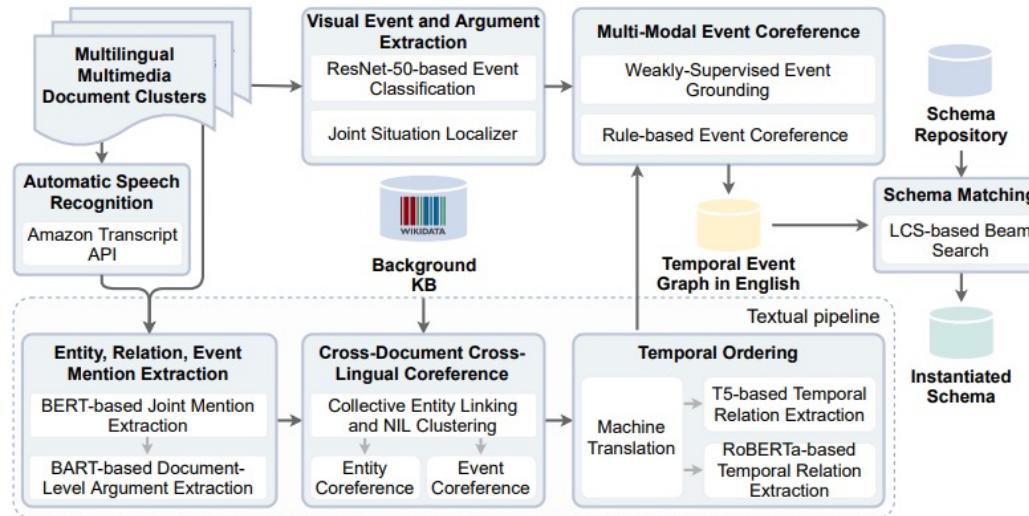


Figure 1: The architecture of RESIN schema-guided information extraction and temporal event tracking system.

RESIN

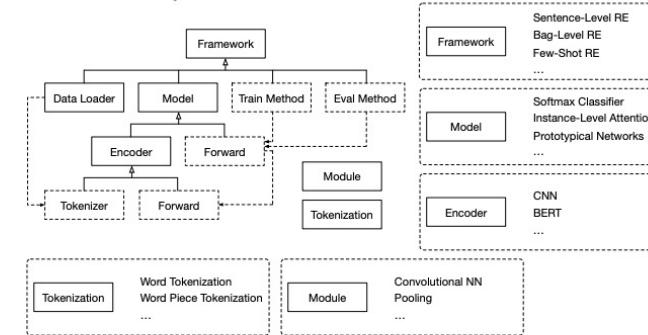
RESIN: A Dockerized Schema-Guided Cross-document Cross-lingual  
Cross-media Information Extraction and Event Tracking System (NAACL2021 Demo)  
OpenNRE: An Open and Extensible Toolkit for Neural Relation Extraction (EMNLP2019)

OpenPrompt: An Open-source Framework for Prompt-learning (ACL2022)

## OpenDelta

An Open-Source Framework for Parameter-Efficient Tuning (Delta Tuning).

### Architecture of OpenNRE



### Example Code

```

rel2id = 'relation->id dictionary'
# Define encoder
sentence_encoder = nrekit.encoder.BERTEncoder(
    max_length=80, pretrain_path='bert pretrain model path')
# Define model
model = nrekit.model.SoftmaxNN(sentence_encoder, len(rel2id), rel2id)
# Define framework
framework = nrekit.framework.SentencereRE(
    train_path='path of training data',
    val_path='path of validation data',
    test_path='path of test data',
    model=model,
    ckpt='path of checkpoint',
    batch_size=64,
    max_epoch=10,
    lr=3e-5,
    opt='bert_adam')
# Train
framework.train_model()
# Test
framework.load_state_dict(torch.load(ckpt)['state_dict'])
result = framework.eval_model(framework.test_loader)

```

OpenNRE

Public Repository : <https://github.com/BBN-E/ZS4IE>

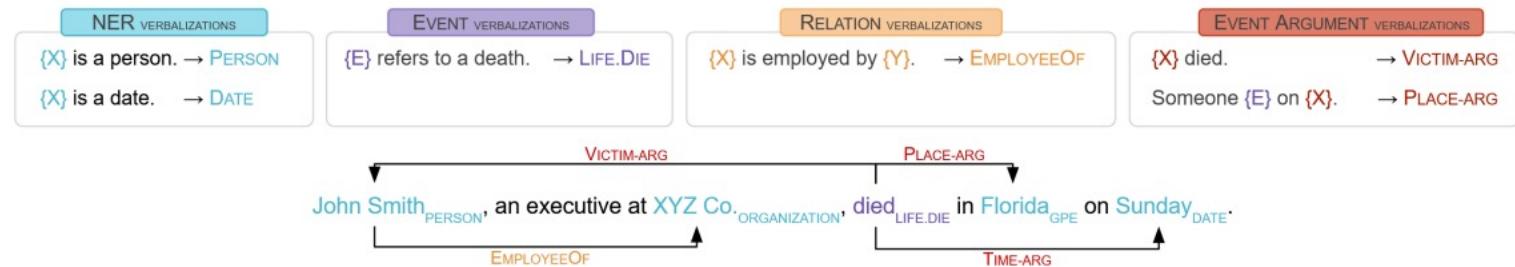


Figure 1: Verbalization templates for a sample schema involving four tasks (from left to right, NER, EE, RE, EAE), with example output (bottom). The schema contains a **EMPLOYEEOF** relation between **PERSON** and **ORGANIZATION** entities and a **LIFE.DIE** event with three argument types (**VICTIM**, **PLACE** and **TIME**) and **PERSON**, **DATE** and **GPE** entities as fillers. Due to space constraints, at most two verbalizations per task shown.



Figure 2: Three steps for entailment-based NER. The steps for the other IE tasks is analogous.

# Open-sourced Frameworks

**DeepKE** is a knowledge extraction toolkit supporting low-resource, document-level and multimodal scenarios.  
**PromptKG** is a prompt learning framework for knowledge graph representation learning and applications.



## DeepKE

*Knowledge Extraction Tool*

# PromptKG



**OpenKG.CN** 中文开放知识图谱

Homepage Demo Download Contact Us Doc OpenKG English 中文

Sentence

- Standard Named Entity Recognition
- Standard Attribution Extraction
- Standard Relation Extraction

Document

- Document Relation Extraction
  - Example
  - Input

Multimodal

- Multimodal Relation Extraction
  - Example
  - Input

**Canton Memorial Civic Center**

The {{0}|LOC|Canton Memorial Civic Center} is a multi - purpose arena located in {{1}|LOC|Canton}, {{2}|LOC|Ohio}, {{3}|LOC|United States}, and is currently the home arena for the {{4}|ORG|Canton Charge} of the {{5}|ORG|NBA G League}. Built in {{6}|TIME|1951}, previous sports teams that have played at the center include the {{7}|ORG|Canton Legends} indoor football team, {{8}|ORG|Canton Invaders} indoor soccer team, and {{9}|ORG|Ohio Aviators} of the {{10}|ORG|American Basketball Association}. The building is owned by the {{1}|LOC|City of Canton} and operated by {{11}|ORG|SMG}. Capacity is {{12}|NUM|5,200} in the arena, and up to {{13}|NUM|600} in the {{14}|LOC|McKinley Room}. The facility has over the years hosted concerts, professional wrestling cards, political rallies, family shows, and features a number of annual {{15}|MISC|Pro Football Hall of Fame} festival events.

**Paris Belongs to Us**

{{0}|MISC|Paris Belongs to Us} (sometimes translated as {{1}|MISC| Paris Is Ours}) is a {{2}|TIME|1961} {{3}|LOC|French} mystery film directed by {{4}|PER|Jacques Rivette}. Set in {{5}|LOC|Paris} in {{6}|TIME|1957} and often referencing {{7}|PER|Shakespeare}'s play {{8}|MISC|Pericles}, the title is highly ironic because the characters are immigrants or alienated and do not feel that they belong at all. The story centres on an essentially innocent young university student called {{9}|PER|Anne} who through her older brother meets a group of friends haunted by mysterious tensions and fears that lead {{10}|NUM|two} of them to commit suicide. Among them is her opposite, a femme fatale called {{11}|PER|Terry} who has had affairs with all the men. The source of the malaise affecting the group is never explained, leaving viewers to ponder how far it might be an amalgam of individual imbalances, general existentialist anxiety, or the more specific paranoia of the {{12}|MISC|Cold War} as the world faced the possibility of nuclear annihilation.

**The Great Appeal**

{{0}|MISC|The Great Appeal} ({{1}|MISC|Italian}: {{0}|MISC|Il Grande appello)) is a {{2}|TIME|1936} {{3}|LOC|Italian} war film directed by {{4}|PER|Marie Camerini} and starring {{5}|PER|Camillo Pilotto}, {{6}|PER|Roberto Villa} and {{7}|PER|Lina d'Acosta}. It is sometimes known by the alternative title {{0}|MISC|The Last Roll-Call}. {{4}|PER|Camerini} was considered to have no sympathies with the {{8}|ORG|Fascist} regime of {{9}|LOC|Italy}, but he made this propaganda film that endorsed the colonial policies of the {{3}|LOC|Italian} government. It was one of a number of {{10}|LOC|African}-set films made during the {{8}|ORG|Fascist} era including {{11}|MISC|The White Squadron} ({{2}|TIME|1936}), {{12}|MISC|Sentinels of Bronze} ({{13}|TIME|1937}) and {{14}|MISC|Luciano Serra, Pilot} ({{15}|TIME|1938}). The film portrays the rediscovery of his patriotism of an {{3}|LOC|Italian}, who eventually dies for his country.

**Click**

Visitors' number: 15650  
Copyright© 2019-2022 OpenKG & ZJUKG All Rights Reserved

<https://github.com/zjunlp>

DeepKE: A Deep Learning Based Knowledge Extraction Toolkit for Knowledge Base Population (EMNLP 2022 Demo)  
PromptKG: A Prompt Learning Framework for Knowledge Graph Representation Learning and Application 2022

# DeepKE: Diverse Scenarios

- Single Sentence

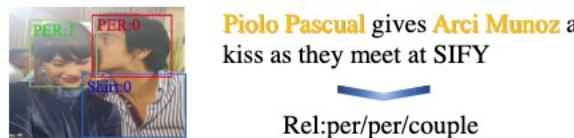
It was one o'clock when we left **Lauriston Gardens**,  
Sherlock Holmes led me meet **Gregson** from **Scotland Yard**.  
[LOC]  
[PER] [PER] [ORG]

- Document

... Elias Brown (May 9, 1793 - July 7, 1857) was a **U.S.** Representative from **Maryland**. Born near **Baltimore**, **Maryland**, Brown attended the common schools ... He died near **Baltimore**, **Maryland**, and is interred in a private cemetery near **Eldersburg**, **Maryland** ...

Intra-sentence    (Maryland, country, U.S.)  
(Baltimore, located in, Maryland)  
(Eldersburg, located in, Maryland)  
Inter-sentence    (Baltimore, country, U.S.)  
(Eldersburg, country, U.S.)

- MultiModal



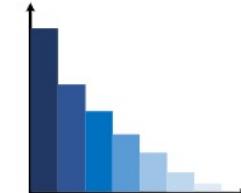
**DeepKE**

*Knowledge Extraction Tool*

Diverse Data



Low Resource



Off-the-shelf Usage    Flexible Training

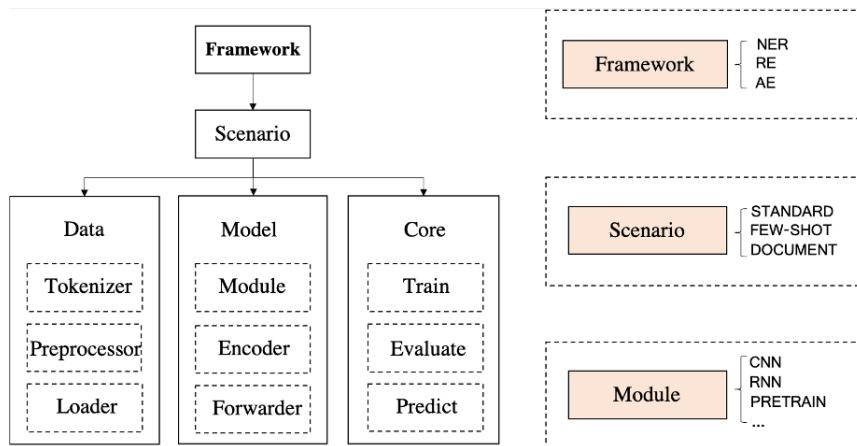
**cnSchema**  
OpenKG 开放的中文知识图谱



# DeepKE: with Traditional and Cutting-edge Methods



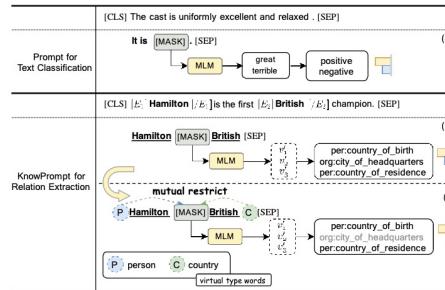
AAACL IJCNLP 2022



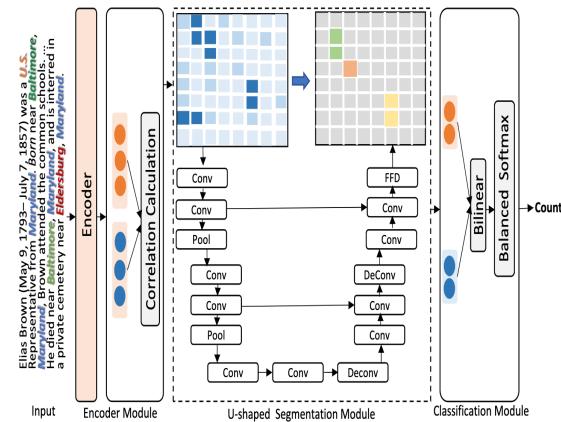
```

# Import packages
from torch import optim
import deepke.relation_extraction.standard.models as models
from deepke.relation_extraction.standard.tools import preprocess, train, FocalLoss
# Preprocess data
preprocess(config)
# Choose the model, users can choose from CNN, RNN, Transformer, GCN, Capsule and BERT
Model__ = {
    'cnn': models.PCNN,
    'rnn': models.BILSTM,
    'transformer': models.Transformer,
    'gcn': models.GCN,
    'capsule': models.Capsule,
    'lm': models.LM,
}
model = __Model__[config.model_name](config)
# Define the optimizer, scheduler and criterion
optimizer = optim.Adam(model.parameters(), lr=config.learning_rate, weight_decay=config.weight_decay)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, factor=config.lr_factor, patience=config.lr_patience)
criterion = FocalLoss()
# Train
train(epoch, model, train_dataloader, optimizer, criterion, device, writer, config)
# Save
model.save(epoch, config)
# Predict
model.load(config.saved_path, device=device)
model(preprocess(config.text))

```

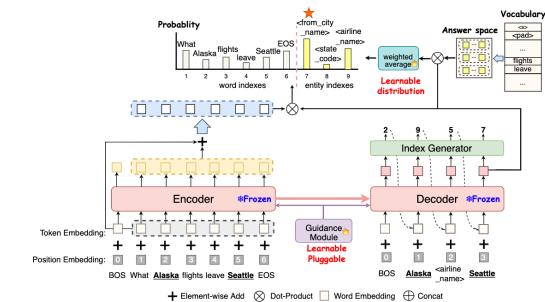


## KnowPrompt WWW'22



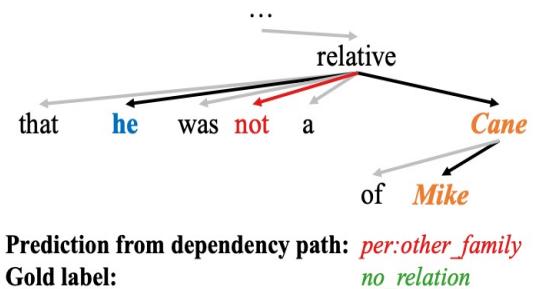
## DocuNet IJCAI'21

- KnowPrompt: Knowledge-aware Prompt-tuning with Synergistic Optimization for Relation Extraction WWW'22
- LightNER: A Lightweight Tuning Paradigm for Low-resource NER via Pluggable Prompting COLING'22
- Document-level Relation Extraction as Semantic Segmentation IJCAI'21
- Graph Convolution over Pruned Dependency Trees Improves Relation Extraction EMNLP'18



## LightNER COLING' 22

I had an e-mail exchange with Benjamin Cane of Popular Mechanics which showed that **he** was not a relative of **Mike Cane**.



Prediction from dependency path: **per:other\_family**  
Gold label: **no\_relation**

## C-GCN EMNLP'18

**Step1** : go to the task folder

```
cd example/ner/standard
```

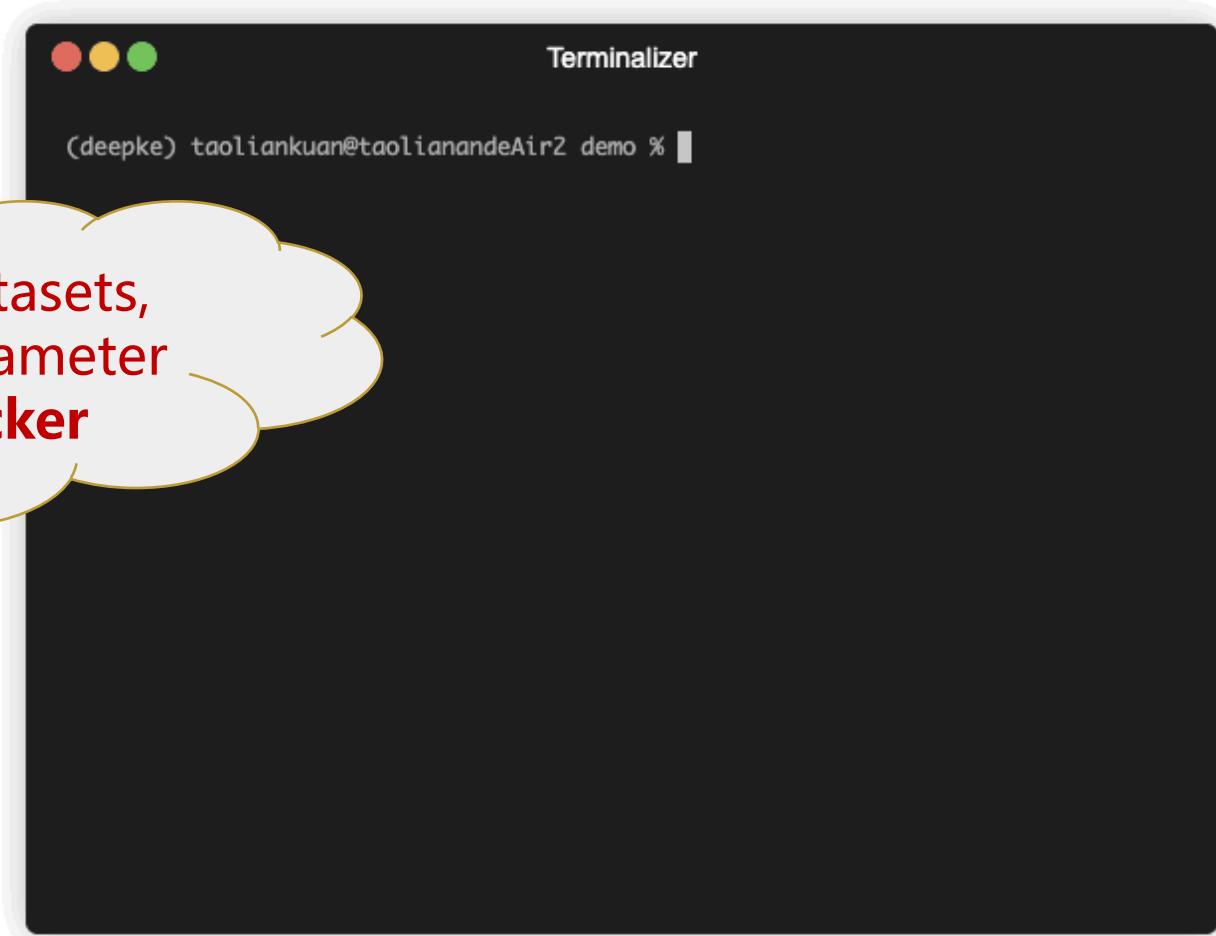
**Step2:** Train

```
python run.py
```

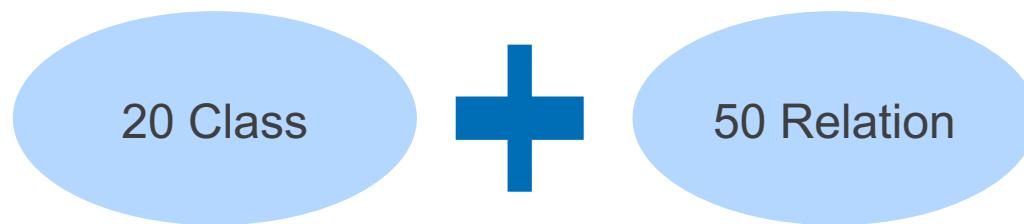
**Step3** : Evaluation

```
python predict.py
```

**Customize datasets,  
auto-hyperparameter  
tunning,docker**



DeepKE provides off-the-shelf extraction models based on cnSchema: supporting **28 entity types** and **50 relation classes**.



## Entity Recognition

DeepKE(NER), Roberta-wwm-ext, Chinese

[Google download](#)

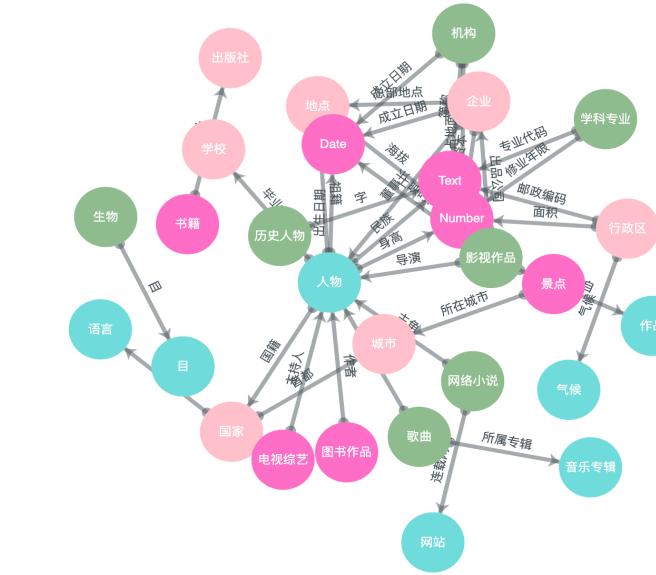
[Baidu Netdisk download](#)(password:u022)

## Relation Extraction

DeepKE(RE), Roberta-wwm-ext, Chinese

[Google download](#)

[Baidu Netdisk download](#)(password:2goe)



```
1 {  
2   "text": "查尔斯·阿兰基斯 (Charles Aránguiz)，1989年4月17日出生于智利圣地亚哥，智利职业足球运动员",  
3   "spo_list": [  
4     {"predicate": "出生地",  
5      "object_type": "地点",  
6      "subject_type": "人物",  
7      "object": "圣地亚哥",  
8      "subject": "查尔斯·阿兰基斯"  
9    }, {  
10       "predicate": "出生日期",  
11       "object_type": "Date",  
12       "subject_type": "人物",  
13       "object": "1989年4月17日",  
14       "subject": "查尔斯·阿兰基斯"  
15     }]  
16 }
```

# DeepKE: Documents&Tutorials

DeepKE

Search docs

GETTING STARTED

Start

Install

Example

FAQ

PACKAGE

DeepKE

Attribution Extraction

Name Entity Recognition

Few Shot

Models

Module

Utils

Standard

Relation Extraction

## Models

### deepke.name\_entity\_recognition.few\_shot.models.model module

```
class deepke.name_entity_recognition.few_shot.models.model.PromptBartEncoder(encoder) [source]
```

Bases: `torch.nn.modules.module.Module`

```
forward(src_tokens, attention_mask=None, past_key_values=None) [source]
```

Defines the computation performed at every call.

Should be overridden by all subclasses.



Although the recipe for forward pass needs to be defined within this function, one should call the `Module` instance afterwards instead of this since the former takes care of running the registered hooks while the latter silently ignores them.

```
class deepke.name_entity_recognition.few_shot.models.model.PromptBartDecoder(decoder, pad_token_id, label_ids, use_prompt=False, prompt_len=10, learn_weights=False) [source]
```

Bases: `torch.nn.modules.module.Module`

```
forward(tgt_tokens, prompt_state) [source]
```

Defines the computation performed at every call.

Should be overridden by all subclasses.



Although the recipe for forward pass needs to be defined within this function, one should call the `Module` instance afterwards instead of this since the former takes care of running the registered hooks while the latter silently ignores them.

```
decode(tokens, state) [source]
```

Comprehensive documents

The screenshot shows a Google Colab interface with a Jupyter notebook titled "ner\_few-shot.ipynb". The notebook has a single cell containing Python code for setting up a few-shot learning environment. The code imports various libraries like torch, tqdm, numpy, and transformers, and initializes a PromptBartModel, PromptGeneratorModel, and ConllNERProcessor. It also sets up a Trainer and Seq2SeqSpanMetric. A WandB configuration section follows, with commands to initialize the project and sync weights. Below the code cell is a terminal window showing the execution of wget and tar commands to download and extract dataset files.

```
import torch
from tqdm import tqdm
import numpy as np
from itertools import chain
from torch.utils.data import Dataset, DataLoader
from torch.nn.utils.rnn import pad_sequence
from transformers import BartTokenizer
import importlib
from deepke.name_entity_re.few_shot.models.model import PromptBartModel, PromptGeneratorModel
from deepke.name_entity_re.few_shot.module.datasets import ConllNERProcessor, ConllNERDataset
from deepke.name_entity_re.few_shot.module.train import Trainer
from deepke.name_entity_re.few_shot.module.metrics import Seq2SeqSpanMetric
from deepke.name_entity_re.few_shot.utils import *
from deepke.name_entity_re.few_shot.module.mapping_type import mit_movie_mapping, mit_restaurant_mapping, atis_mapping

import wandb
import logging
logger = logging.getLogger(__name__)
writer = wandb.init(project="DeepKE_NER_Few")

wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit: .....
wandb: [REDACTED] No API key specified.
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit: .....
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
wandb: wandb version 0.12.9 is available! To upgrade, please run:
wandb: $ pip install wandb --upgrade
Syncing run royal-bee-1 to Weights & Biases (docs).

!wget 120.27.214.45/Data/ner/few_shot/data.tar.gz
!tar -xzvf data.tar.gz

./data/
./data/atis/
./data/atis/test.txt
./data/atis/10-shot-train.txt
./data/atis/20-shot-train.txt
./data/atis/50-shot-train.txt
```

Google Colab Tutorials



# The End

Github: <https://github.com/NLP-Tutorials/AAACL-IJCNLP2022-KGC-Tutorial>

Questions: [zhangningyu@zju.edu.cn](mailto:zhangningyu@zju.edu.cn), [tgui@fudan.edu.cn](mailto:tgui@fudan.edu.cn), [nanguo2021@bupt.edu.cn](mailto:nanguo2021@bupt.edu.cn)